

Acknowledgements

[TODO: write at the end]

Contents

Acknowledgements	1
1 Introduction	4
2 Literature review	11
2.1 From text foundation models to single-cell foundation models	11
2.2 Cell sentences, and the model this thesis fine-tunes	11
2.3 Predicting chemical perturbation without a language model	13
2.4 The dispute: do these models beat uninformative baselines?	15
2.5 Six ways a drug-response benchmark can mislead	17
2.6 Connectivity mapping and signature transfer	18
2.7 Decodable, used, and the global workspace	18
2.8 Optimal transport for state matching	20
3 The investigation	22
3.1 The atlas, the model and the rules of measurement	23
3.2 The metric is the first result	29
3.3 The model is drug-blind	37
3.4 Is the drug represented at all?	43
3.5 Is the representation used?	45
3.6 Is the drug’s identity read?	47
3.7 Denoising the target: three nulls at one scoring	50
3.8 The residual frame: predicting only what makes this drug different	52
3.9 How much signal is there to encode	57
3.10 Re-encoding produces prompt sensitivity	60
3.11 Sensitivity is not fidelity	64
3.12 A lookup table wins	69
3.13 Channel availability for an unseen drug	71
3.14 Which part of the prompt is read remains unresolved	75
3.15 The investigation in one table	76
4 Limitations and Future Research Directions	88
4.1 Limitations of scope	88
4.2 Limitations of the unseen-drug evaluation	90
4.3 Directions	92
5 Conclusions	96

A	Appendix	107
A.1	Estimator validation against planted ground truth	107
A.2	Absent-gene convention	108
A.3	Swap-distance sweep	108
A.4	Training configuration	109
A.5	The identity test: failed designs and remaining guards	109
A.6	Calibration self-test	111

1 Introduction

A compound that enters a cell changes what that cell transcribes. Which genes move, by how much, and in which direction is the readout the field treats as its most direct account of what a compound actually does — not what it was designed to bind, but what the cell does in response. That is a working premise rather than an established fact, and how far it holds in any given experiment is itself a measurement: §3.1 finds a substantial share of this atlas’s treated conditions statistically inert or unidentifiable in principle, and nothing can be graded on those. Reading the response is routine. Reading it for every combination that might matter is not. A compound behaves differently at different doses, and differently again depending on the genetic background of the cell receiving it, so the space to be explored is the product of three large sets rather than any one of them. Even the largest screens sample a small corner of it, and every combination left untested is a question left open.

This is the gap prediction is meant to close. If the transcriptional response to a treatment could be computed rather than measured, the untested corners of that space would become inspectable, and the experiments actually run could be chosen because they are informative rather than because they happened to fit on the plate. The task, stated plainly, is this: given the transcriptome of an untreated cell and the identity of a drug, predict the transcriptome that cell would have shown had the drug been applied.

The task is harder than it sounds, for a reason that is easy to state and impossible to engineer around. The ground truth is a counterfactual. Sequencing destroys the cell it reads, so the treated and untreated profiles of the same cell can never both be observed, and a prediction can only ever be checked against a different cell that received the drug. Everything downstream inherits this. It is why paired training data must be assembled by matching cells rather than by tracking them, why evaluation is usually framed at the level of populations rather than individuals, and why distinguishing a real drug effect from ordinary cell-to-cell variation requires care rather than optimism.

Assembling data at the scale this problem demands became possible only recently. Two decades of falling sequencing costs, rising throughput and increasingly clever sample multiplexing turned single-cell profiling from a delicate experiment into something that can be run industrially, and Tahoe-100M [1] is what the end of that trajectory looks like: roughly a hundred million single-cell transcriptomes spanning some 1,100 small molecules across 50 cancer cell lines. Every compound and every cell line appears not as a summary statistic but as a cloud of individual cells.

That resolution changes which questions can be asked. The obvious one is conditional — what does this compound do to this cellular context, rather than what does it do on average across contexts that may have little in common. The less obvious question is whether the first has an answer. Because each treated condition is a population of

cells rather than a single measurement, it becomes possible to measure how much of an apparent drug effect reproduces across disjoint halves of the same treated well and how much is only variation between the cells that happened to be sampled. That second measurement turns out to matter as much as the first, and a good deal of what follows rests on the observation that standard evaluations quietly confuse the two.

The other development behind this thesis has nothing to do with cells. Over the same period, pretrained transformer language models became unreasonably good and unreasonably general, and the engineering poured into them — architectures, training recipes, scaling laws, the entire surrounding toolchain — came to exceed what any single scientific field could plausibly assemble for itself. That created a standing temptation: if a problem can be made to look like text, it inherits all of that work for free.

Cell2Sentence [2] takes the temptation about as literally as it can be taken. Rank the genes in a cell from most to least expressed, write their names out in that order, and the cell becomes a sentence — ordinary text that an ordinary language model can be fine-tuned on. The encoding is unsentimental about biology, discarding how much each gene is expressed and keeping only the ordering, on the argument that rank retains enough of the signal for the rest to be reconstructed. In exchange it requires no new architecture and no new training machinery, only cells written in a format that existing models already consume.

C2S-Scale [3] develops that idea into a family of models. Larger backbones are pretrained on over 50 million cells drawn from 825 single-cell datasets in CELLxGENE and the Human Cell Atlas, and then fine-tuned across a spread of tasks: predicting cell type and tissue, generating cells conditioned on a label, captioning clusters, answering questions about a dataset in natural language, and predicting perturbation responses. That last task is the one this thesis builds on, and it is worth being specific about its scope. Their perturbation work covers cytokine stimulation in peripheral blood cells, and, for small molecules, bulk L1000 profiles and the SciPlex3 panel. The single-cell results are scored as distributions — how closely the generated population for a condition resembles the observed one — rather than condition by condition. The checkpoint used throughout this thesis, **C2S-Scale-Pythia-1b-pt**, is the pretrained model from that family: a Pythia-1b language model [4], pretrained on natural language and then further pretrained on atlas cells, with no drug-treated cells among that atlas data.

What we do is extend that line of work to chemical perturbation at the scale Tahoe-100M now makes available, and evaluate it as a prediction problem in the strict sense — condition by condition, against controls chosen so that a model which ignores the drug cannot pass. The question is whether a cell-sentence model fine-tuned on this atlas learns what a particular compound does to a particular cell.

Two things are worth keeping straight about what the model has actually seen, since the scale of the atlas invites an easy overstatement. The fine-tuning does not consume

Tahoe entire — it draws a sample of a few hundred thousand treated cells from it, which is ample for the task but is not a hundred million. And nothing here is unseen by the whole pipeline: the backbone’s language pretraining is text in which compound names, targets and mechanisms occur, and every prompt names the compound outright and carries a **Mechanism** field — its annotated mechanism of action where one exists, the literal string **unclear** where it does not. What a held-out drug is held out of is the Tahoe fine-tuning, with its name and whatever mechanism annotation it carries still supplied at inference — not the model’s prior exposure to that name.

It does learn something, and the something is the difficulty. What the model reproduces is the part of a treated profile that is not specific to the treatment: the cell’s own state, and the programme that perturbed transcriptomes have in common whichever compound produced it. What it does not reproduce is the part that separates one compound from another — and that part is not a remainder. Decomposing the treated-minus-control shift gives 62.1% of its energy to the drug-specific residual against 39.0% to the generic programme, the two components essentially orthogonal, so the missing one is the larger of the two. Change the drug named in the prompt while holding the input cell fixed, and the prediction barely moves. Within plate on unseen drugs the model scores 0.498 against a within-well split-half precision reference of 0.576, and a baseline that copies the untreated cell and is told nothing about any drug scores 0.504. Restricting to the conditions where the task is provably winnable does not separate them: model 0.768, control-copy 0.766, a paired contrast of +0.002 [−0.026, +0.028] over 204 conditions whose own noise ceiling is 0.938. That interval is narrow relative to the range available on that stratum, so what it rules out is not only a large effect but a useful one.

The reply this invites is that it is a fact about one checkpoint and a larger one would do better. We cannot answer that directly: none was run, and Section 4 says so rather than argue around it. What can be said is which parts of the result a larger checkpoint would change, and the two measurements that bound the reply never touch the model at all. The first is the scoreboard: a predictor that places every gene at the middle rank, and therefore carries no information about drug, cell or biology, scores $DE-\Delta r = 0.9999$ [0.99991, 0.99992] on the field’s default metric, above both the two-cell noise ceiling and the model. The second is the target. Because a cell sentence ranks genes by expression, the string the model is trained to emit is nearly the same string for any two drugs in the same cell line — only 17% of its top 200 gene positions differ between them — and of the divergence that re-encoding recovers, 92% comes from subtracting the cell’s own baseline state rather than the response drugs share with one another. Most of what makes two targets look alike is the cell, not a common reaction to being drugged. Neither number involves a forward pass, and the within-well retrieval reference is unmoved by the re-encoding (0.742 against 0.743), so what changes is not how much the data distinguish the two drugs but how much of that distinction reaches the tokens. A larger backbone would

be fitted to the same targets and scored on the same metric. Scale is a live question for this task, and Section 4 says where we think it might bite; it is not a question about either of those two numbers.

The third measurement does involve the model, and it is the one that says what kind of failure this is. A per-drug average — the compound’s mean drug-specific residual, fitted from training conditions in *other* cell lines — beats the model, and not as a low bar. On the conditions where both are defined ($n = 1192$) the lookup scores 0.913 against the model’s 0.549; of the distance from chance to the within-well split-half precision reference of 0.857 it covers 115.7% where the model covers 13.9%. That the lookup passes the reference it is measured against is a property of the reference — one noisy half-sample against a many-line average, on a metric that compresses near its top — and is not an estimate of headroom. Nor is the ordering an artefact of averaging many lines: restricted to a *single* other cell line the same lookup still scores 0.822. What no lookup can do is condition on the cell in front of it, which is why this is the useful comparison: it separates the drug’s average effect from the part that depends on context. Section 3.9 measures that split with a transfer coefficient $\mathcal{T} = 0.553$ [0.514, 0.593], leaving 44.7% [40.7%, 48.6%] of the drug-specific residual not shared across the compared contexts. That figure is descriptive and should not be carried further than it goes: the estimator mixes cell line, dose, well and target estimator, so it is neither a drug-by-cell-line interaction share nor a bound on what a model could reach.

A negative result earns its place only if it can be made specific, and that is the organising principle of what follows. We establish first that the failure is real rather than an artifact of how this field measures success — a distinction less pedantic than it sounds, since the standard metric can be saturated by a predictor that knows nothing about drugs. We then ask where in the pipeline the failure lives, and the answer is not that the drug is missing. It is recoverable from the data wherever the task is winnable at all, and it is linearly decodable from the model’s own activations, more explicitly the deeper one reads; what the generation readout does with that representation is consult it at very low gain — a small effect, replicated across seeds in one training arm and absent in another, and surviving multiplicity correction only at one rung of the displacement ladder, offered as the modest reading rather than as a mechanism. At the token budget those arms were scored at, no repair that changed what the target *contains* separated from its scramble, which pointed at how the target is *written*, and re-encoding it as the drug-specific residual produces the first measurable drug effect in this project — first in sequence rather than uniquely, since a budget-matched re-run of the optimal-transport target clears zero as well. That is where the gap the thesis actually turns on opens: prompt sensitivity is not biological fidelity. The re-encoded model responds to the drug named in its prompt and still recovers a small fraction of that drug’s signature — a smaller fraction than a training-only lookup that never sees the cell it is predicting for. The thesis therefore ends not in

closure but in the one place a lookup cannot go by construction: the drug never seen in training. Of the three properties a compound could be generalised through, exactly one carries signal we are willing to call measurable, and only on one of the two axes it could be clustered over.

Four questions carry that argument, in order:

1. **Is the evaluation trustworthy?** Before asking whether a model is any good, one has to ask whether the metric could tell. Can a predictor carrying no drug information score well on the measures the field has standardised on, and if so, which measures survive?
2. **Does the model use the drug at all?** Holding the input cell fixed and changing only the compound it is told about, does the prediction move — and does it move in the direction the real response moved?
3. **If not, where does the failure live?** Is the drug missing from the data, missing from the model’s internal representation, or present in that representation but ignored by the process that writes the output? Each answer implies a different repair, and the cheapest one turns out not to be the one that is needed.
4. **Does it beat the thing it has to beat?** A table of per-drug averages, fitted on training conditions and blind to the cell it is predicting for, is the pre-registered win condition: a lookup can reproduce a memorised average, so only a model that tailors its answer to the control cell’s state can beat one. Does this model? And of the signal a lookup cannot express by construction — the part that does not transfer across contexts — how much does any predictor here recover?

The contributions follow the same order. We audit the calibration of the standard evaluation and show that it can be saturated by a predictor carrying no drug information whatsoever, which bears on any work built from the same family of metrics. Of the five metrics we test, only the normalised inverse rank has a positive dynamic range fraction — +0.635 across plates and +0.545 within plate, against −0.16 to −0.69 for the other four across plates and −0.10 to −0.52 within plate. The four negatives are the stronger half of that result rather than a footnote to it: a negative fraction means the metric ranks an uninformative predictor *above* a real replicate, so those four are inverted and not merely blunt. Both surviving figures are measured with the bootstrap clustered on cell lines rather than on the drug rows nested in them — [+0.604, +0.663] across plates over 25 lines, [+0.516, +0.569] within plate over 27, each with Holm $p = 0.002$. The audit then repeats in the residual frame, on a different target, a different comparison set and a different positive control, where the same metric returns +0.704 [+0.667, +0.736] over 42 cell lines against a zero-drug-information negative arm — a second calibration result, not a robustness check on the first.

We establish drug-blindness on that one surviving metric, using several independent instruments so that no single measurement carries the claim alone, and we separate out two things the failure is not. It is not that the drug is absent from the data: on the stratum where the task is provably winnable the signal is there and the model still does not use it. And it is not that the drug is absent from the model: drug identity is linearly decodable from the residual stream, more explicitly the deeper one reads, while the generation readout consults that representation at very low gain. That dissociation between what is encoded and what is read is a small effect rather than a mechanism — present in one training arm and absent in another, and surviving multiplicity correction only at one rung of the displacement ladder. It is the least well-evidenced link in the chain, which is why no argument downstream rests on it alone.

We identify the target encoding as the most probable binding constraint, and the exclusion behind that is narrower than an earlier framing of it claimed. What it excludes is the data and, at the token budget those arms were scored at, the content of the target: no repair that changed what the target holds separated from its scramble, which left how the target is written. It does not exclude the architecture. The one intervention this thesis recommends building — a learned per-drug modulation of the control-to-treated transformation — is architectural rather than a change of prompt, and it is recommended on the size of the readout deficit rather than on the activation evidence above, which is too weak to carry it. The re-encoding that tests the diagnosis also changes aggregation, control referencing and sequence length, so the attribution is to the encoding as a package rather than to token dilution alone. It produces the first measurable drug effect in this project. On training conditions the model beats its designated comparator (`scramble_orth`: the same prompt with the drug name replaced by a within-plate partner chosen so that its measured response is near-orthogonal to the target drug’s — chosen that way because a comparator handed a deliberately *dissimilar* drug scores below chance and inflates the gap it defines, which is a result of this work rather than a design detail, and one that any label-swap ablation scored by a discrimination metric inherits) by $+0.0898$, with an interval excluding zero whether the resampling is clustered on cell line $[+0.0382, +0.1415]$ or on drug $[+0.0291, +0.1506]$. Re-encoding is first in sequence rather than uniquely responsible: a budget-matched re-run of the optimal-transport target clears zero as well, at $+0.0292$ $[+0.0144, +0.0441]$.

We do *not* claim that this transfers. On held-out drug-and-cell-line pairings the same gap is $+0.0332$, which excludes zero clustered on cell line $[+0.0069, +0.0594]$ but not clustered on drug $[-0.0031, +0.0695]$ — and drug is the axis across which a claim about transferring a drug signature has to generalise. That split is in any case mostly recombination rather than novelty: 30 of its 34 drugs and 39 of its 42 cell lines also appear in the training split, and 82% of its conditions pair a training drug with a training cell line. On genuinely unseen drugs the gap is -0.0315 , with intervals spanning zero under

both clusterings ($[-0.0836, +0.0206]$ on cell line, $[-0.1049, +0.0419]$ on drug). Nor are the three gaps on a common scale: the within-well split-half precision reference differs across the splits (0.956, 0.824, 0.836), so their ordering cannot be read as a clean decay with novelty.

We then show that the sensitivity and the fidelity are different quantities, which is the result the rest of the thesis is organised around. The per-drug lookup quoted above — fitted on training conditions only, blind to the cell it is predicting for, 0.913 against the model’s 0.549 — holds up even when its averaging advantage is removed: drawn from a *single* other cell line it still scores 0.822. That the lookup passes the precision reference is a property of that reference — a single noisy half-sample against a many-line average, on a metric that compresses near its top — so the ordering is what carries the claim, not the excess. This is DrEval’s finding replicated in transcriptomic space, on a 946-gene response rather than a scalar potency, and it is what makes the negative a statement about the model rather than about the noise floor.

Finally we ask what remains reachable for a drug never seen in training, by grading the three channels through which such a drug could be predicted at all. Protein target is the only one that clears the pre-set relevance margin of 0.03, and it clears it clustered on cell line and not clustered on drug, which is the axis an unseen-drug claim rides on; it rests on 18 annotated drugs of which two are held out. Mechanism and chemical structure are inconclusive rather than closed, and mechanism is graded in a frame biased against it, since the plate-local generic removes part of a co-plated mechanism class’s own shared programme before the channel is scored. Alongside that we report how much of the reproducible drug response is not shared across the compared contexts — 0.447 [0.407, 0.486], mixing cell line, dose, well and estimator rather than isolating a drug-by-cell-line interaction. That last number is scope-setting rather than load-bearing: it says how much of the signal a lookup cannot express, not how much anything could recover.

The thesis proceeds as follows. Section 2 places the work against single-cell foundation models, the perturbation-prediction literature, and an unsettled dispute within it about whether deep models beat trivial baselines at all — a dispute this thesis adjudicates in one corner, using the instruments of the side it ultimately argues against. Section 3 is the investigation itself, told in the order it happened: because nearly every instrument in it was built in response to the measurement that came before, the data, the methods and the findings are presented as one continuous account, with each definition given where it is first needed. Section 4 sets out the scope honestly, separating what is closed by measurement from what is merely untested, and Section 5 concludes.

2 Literature review

2.1 From text foundation models to single-cell foundation models

The models this thesis builds on descend from a template established in natural language processing. A transformer is pretrained on a very large unlabelled text corpus with a single objective — predict the next token — and the resulting network is then adapted to specific tasks by fine-tuning or prompting. The GPT line of models made the case that this recipe scales: capability on tasks never explicitly trained for improves steadily with model size, data and compute, so a single pretrained backbone can serve as a general substrate rather than a task-specific tool. What made the recipe portable is that it assumes almost nothing about language. It requires only that the data be expressible as sequences of discrete tokens, and that some tokens be predictable from others.

That portability is what invited the transfer into biology. If a cell can be written as a sequence, the same machinery applies unchanged. scGPT [5] and Geneformer [6] both take this route, treating a cell’s expression profile as a sequence of gene tokens and pretraining on tens of millions of cells before transferring to classification, integration and perturbation tasks. The ambition is usually stated as some version of the *virtual cell* [7]: a model whose simulated responses to intervention are trustworthy enough to substitute for an experiment.

The interesting variation among these models is not architectural but representational. A cell is a vector of expression values over some twenty thousand genes; a transformer consumes a sequence of discrete tokens. Something must be discretised or discarded in between, and what gets discarded bounds what the model can learn. The approach this thesis uses makes that choice unusually explicit.

2.2 Cell sentences, and the model this thesis fine-tunes

Cell2Sentence [2] takes the most literal route available: rank a cell’s genes by expression and write their symbols out in descending order. The cell becomes a *cell sentence*, which is text in the ordinary sense — a string that a tokeniser will happily consume:

MALAT1 ACTB TMSB4X B2M EEF1A1 RPL13A FTL TPT1 RPS27 GAPDH ...

Nothing marks this as biological. The information is carried entirely by which symbols appear and in what order, and a language model fine-tuned on such strings is doing exactly what it does on prose: predicting the next token given the previous ones.

The encoding keeps rank and discards magnitude. C2S defends the loss with a reconstruction argument: expression can be recovered from rank through a fitted log-linear relationship $e_i = a_d \log(r_i) + b_d$, which achieves $R^2 = 0.815$ against the original values

on peripheral blood data. On that basis they argue rank encodes expression reversibly enough across datasets for the discarded magnitude not to matter for downstream use. It is worth noting what kind of criterion this is — reconstruction fidelity is a statement about *absolute* accuracy, how close the recovered profile is to the true one, which is not automatically the same as preserving whatever distinguishes one experimental condition from another.

Understanding what the model does with such a sentence requires a little of the underlying architecture, because several of its properties matter later. C2S-Scale builds on decoder-only transformers of the GPT family, and the checkpoint used here is built on Pythia [4], a suite of such models released with open weights and training data. A decoder-only transformer factorises the probability of a sequence into a product of next-token predictions,

$$p(x_1, \dots, x_T) = \prod_{t=1}^T p(x_t \mid x_{<t}), \quad (1)$$

and is trained by minimising the cross-entropy between each predicted distribution and the token that actually occurred. Each token is first mapped to a vector by an embedding matrix, and those vectors form what is usually called the *residual stream*: a running representation, one vector per position, which every layer reads from and writes back into additively. A layer comprises multi-head self-attention, which lets each position gather information from earlier positions and is *causally masked* so that position t attends only to positions $\leq t$, followed by a position-wise feed-forward network. Both are wrapped in normalisation and added back to the stream rather than replacing it. After the final layer an unembedding matrix projects the stream to one score per vocabulary entry, and a softmax turns those scores into the next-token distribution.

Three properties of that design are worth noting now, since later sections depend on them. The residual stream is additive and shared, so a representation can be read at a chosen depth, modified, and the remaining computation allowed to proceed on the modified value — which is what makes interpretability interventions of the kind used in Section 3.6 possible at all. The training loss is a sum of per-token cross-entropies that weights every position equally, so a target whose tokens are largely shared between conditions gives any individual condition only a small share of the gradient. And conditioning is mediated entirely by attention across a causal mask, so information supplied early in a prompt must be reached from every later position at which it is needed.

C2S-Scale [3] develops Cell2Sentence into a family of models on this architecture, spanning backbones from Pythia at 410 million and 1 billion parameters up to Gemma-2 at 27 billion. Pretraining uses over 50 million human and mouse cells drawn from 825 single-cell datasets in CELLxGENE and the Human Cell Atlas. The pretrained models are then fine-tuned across a range of tasks: predicting cell type and tissue, generating cells conditioned on a label, captioning clusters, answering natural-language questions about

a dataset, and predicting perturbation responses.

The perturbation task is the one this thesis extends, and its published scope is worth stating precisely. C2S-Scale evaluates it on three benchmarks: cytokine stimulation of peripheral blood cells, bulk L1000 profiles, and the SciPlex3 small-molecule panel. In each the model receives an untreated profile together with a description of the treatment and produces the treated profile, so the setup is conditional rather than label-only generation. Two features of it are relevant to the comparison this thesis draws. Pairing is by group rather than by cell — for the bulk benchmark, treated and untreated samples are matched by cell line name — which is unavoidable given that sequencing destroys the cell it reads. And the single-cell results are scored as distributions: for each condition the model generates as many cells as were observed, and the generated population is compared against the observed one, which asks whether the model produces a plausible population for a condition rather than whether it distinguishes that condition from another.

The checkpoint fine-tuned in every experiment below is **C2S-Scale-Pythia-1b-pt**, the pretrained member of that family, whose training corpus is atlas data alone and contains no drug-treated cells.

2.3 Predicting chemical perturbation without a language model

Cell sentences are an unusual way to attack this problem, and the usual way is worth setting out, because it is the comparison a reader reaches for first. The direct line of work predicts a treated cell as a displacement in a learned latent space. scGen [8] estimates the vector between control and perturbed populations in the latent space of a variational autoencoder and adds it to cells from a context where the perturbation was never applied: the perturbation is one offset, and the autoencoder supplies everything else. CPA [9] makes that arithmetic explicit and learnable — a basal cell state plus additively composed embeddings for drug, dose and covariates, decoded back to expression — so an unobserved combination is assembled from parts that were each observed separately. chemCPA replaces CPA’s one-hot drug embedding with a molecular-structure encoder, which is what allows a compound absent from training to be embedded at all.¹ GEARS [10] does the analogous thing for genetic perturbations, reaching unseen genes through a gene–gene knowledge graph rather than through chemistry. A second family replaces the additive displacement with a learned map: CellOT fits, for each perturbation, a neural optimal-transport map from the control population to the treated one, and its conditional extension amortises that map over a perturbation embedding so that a single network covers many treatments.²

¹Hetzel et al., *Predicting Cellular Responses to Novel Drug Perturbations at a Single-Cell Resolution*, NeurIPS 2022.

²Bunne et al., *Learning single-cell perturbation responses using neural optimal transport*, Nature Methods 2023; Bunne, Krause and Cuturi, *Supervised Training of Conditional Monge Maps*, NeurIPS 2022.

Two properties are common to all of them, and neither holds for the model fine-tuned here. The prediction is a vector in a continuous space, emitted in one pass by a decoder trained on a loss written directly on that vector; and the drug enters through a channel the architecture was built to receive it on — a learned embedding in CPA, chemCPA and GEARS, an estimated latent offset in scGen, and a separately fitted map per perturbation in CellOT, amortised over a perturbation embedding in its conditional extension. A cell-sentence model emits a token sequence sampled left to right, carries the drug as a few words of a natural-language prompt, and is trained by a per-token cross-entropy that weights every position equally. That difference is not decorative. In the standard cell-sentence target, 83% of the top-200 token positions are identical between any two drugs in the same context (Section 3.8) — the string to be emitted is nearly the same string whichever drug was named — and a latent-transition model has no target with that property. Whether that property is the binding constraint is not settled here: Section 3.10 reports what re-encoding the target does, and the chapter does not establish that tokenisation is the constraint that binds. What is reported below is therefore specific to conditional generation over an ordering, and it is not evidence about CPA, chemCPA or CellOT, none of which was run here.

The two families are also scored differently, and that difference is the subject of Section 3.2. The latent-transition literature is evaluated condition by condition against the observed population — distributional distances, R^2 on pseudobulk, overlap of differentially expressed genes — which asks whether a prediction resembles what its own condition produced. It does not ask whether the prediction resembles its own condition more than it resembles some other drug’s, and those two questions come apart precisely when most of the response is shared across drugs. Benchmark suites have consolidated the first: scPerturb [11] harmonises the datasets and standardises distances between perturbed populations, and PerturBench [12] assembles models and metrics for the same comparison across perturbation types. PerturBench also supplies the exception this thesis relies on, a rank-based score that grades a prediction against a set of candidate conditions rather than against one truth.

None of these methods is run as a baseline here, so the head-to-head a reader wants is missing; that is a real gap, and the results chapter concedes it as one where the comparison would fall due. What is offered in its place is narrower than a head-to-head: the thing these methods have in common is measured directly. In scGen and in CPA the drug-specific content carried into a new context is a single per-drug vector — dose-scaled, in CPA’s case — and everything else comes from the encoder and the decoder. Section 3.12 measures what such a vector is worth on this task by looking the drug up: the mean residual of the same compound in other cell lines, fitted from training conditions only. On the conditions where both are defined it scores 0.913 against the fine-tuned model’s 0.549 ($n = 1,192$), and 0.890 against 0.533 on the 893 held-out drug-and-cell-line combinations, where the

lookup cannot borrow the target’s own well; chance on that scale is 0.50 by construction. That is not a score for CPA or chemCPA — their decoders make the emitted effect context-dependent even when the drug channel is one vector — but it does say which question is the informative one. It is not which architecture wins. It is whether anything improves on a per-drug constant, and that is answerable, and answered below, without fitting a second architecture.

Two gaps bear on what follows, one in this literature and one in this thesis. Tahoe-100M [1] is recent enough that we are aware of no published evaluation of any of these methods on it, so the reference points used below are constructed rather than cited — which is why so much of the next chapter is spent building baselines before anything is scored against them. And the atlas’s own authors have since released Tahoe-x1 [13], a family of single-cell foundation models pretrained on a corpus that includes this atlas. It is not a cell-sentence model, so it does not isolate the representation variable this work turns on, but it is the obvious external check on whether a readout failure of this kind survives perturbation-specific pretraining; it was not evaluated in this thesis, and nothing measured here speaks to it.

2.4 The dispute: do these models beat uninformative baselines?

The field contains an unresolved and unusually sharp disagreement about whether any of this works, and this thesis is best read as an attempt to settle one corner of it.

On one side sits a series of benchmarks reporting that deep perturbation models do not beat trivial predictors. Ahlmann-Eltze, Huber and Anders [14] find that deep-learning prediction of genetic perturbation effects does not outperform simple linear baselines. Csendes et al. [15] reach the same conclusion for foundation cell models on post-perturbation RNA-seq, and Bendidi et al. [16] put it most bluntly in a title, reporting that a single PCA still rules them all. Kernfeld et al. [17] concur in a broad comparison of expression-forecasting methods. The pattern is not confined to transcriptomic readouts: on scalar drug sensitivity, Hauptmann and Kramer [18] hold the experimental protocol fixed across architectures and find that the measured differences between them largely dissolve, and DrEval [19] finds that deep models barely outperform a naive predictor built from mean drug and cell-line effects. The recurring result is that a per-perturbation mean is remarkably hard to beat.

Against this, Miller et al. [20] argue that these conclusions “largely stem from limitations of benchmarking metrics, not from the models themselves”, and the argument is methodological rather than rhetorical. They formalise the use of positive and negative controls in perturbation benchmarking and introduce a new positive control, the *interpolated duplicate*, which blends a technical duplicate with the mean baseline using weights drawn from per-gene differential-expression significance — giving a reference point of known and

tunable quality. They then define a quantitative measure of whether a metric can see that quality at all, the *Dynamic Range Fraction* (DRF), and apply it across 14 Perturb-seq datasets and 13 metrics. Widely used measures such as MSE and control-referenced Pearson correlation turn out to be poorly calibrated, while weighted and rank-based alternatives — weighted MSE, weighted $R^2\Delta$ [21], and Normalised Inverse Rank [12] — are consistently calibrated. Evaluated on the metrics that survive, they report, deep models *do* beat mean, control and linear baselines.

What is actually in dispute is narrower than it first appears, because both sides agree the standard metrics are defective. The disagreement is about what a corrected evaluation would reveal. One reading is that the models were always capturing real biology and the broken metrics were hiding it, so better measurement should expose genuine skill. The other is that the models were never capturing much, and better measurement will simply show that failure more clearly and more honestly. Put that way the question is empirical rather than interpretive — but it is settled only within the setting where it is tested, and neither camp has tested it in the setting of this thesis. Miller et al. and Ahlmann-Eltze et al. both study **genetic** perturbation, where an intervention is a designed, near-binary knockout with a named target gene. DrEval studies **drug sensitivity** as a scalar potency such as IC_{50} , predicted from cell-line omics. None of the works in the dispute takes as its object a full transcriptomic response to a small molecule — which acts at a dose and typically engages several proteins at once — at single-cell resolution, produced by an autoregressive generative model. The models of Section 2.3 do work on chemical perturbation at single-cell resolution, so the setting itself is not unoccupied. What we are aware of no one having done in that setting is running the calibration audit first, and letting it decide which metrics are admissible before any model is scored.

The position taken here is to adopt the proponents’ instrument rather than argue with it. Section 3.2 ports the DRF framework onto drug perturbation and uses it to decide which metrics are admissible in this setting before any model is scored. The framework and the interpolated-duplicate positive control it rests on are inherited; not every choice inside them is. The metric that survives is computed under a similarity function we picked — Euclidean distance in expression space rather than over rank vectors — on a criterion stated where the choice is made, and the two variants do not agree. The accurate statement is therefore that the design of the audit comes from the sceptics’ opponents and one parameter inside it is ours, disclosed at the point of use. What adopting the design buys is that the verdict cannot be dismissed by the objection Miller et al. themselves raise.

That verdict is where this thesis and its own instrument part company, and the divergence is a result rather than an embarrassment. The audit reproduces their central finding: metric choice dominates, the rank-based metric survives, and the correlation-based measures they flag as poorly calibrated are exactly the ones that fail here — four of

the five metrics audited are not merely uninformative but inverted, scoring a drug-agnostic baseline above the noise ceiling, each with an interval excluding zero. The agreement is not total, and the exception belongs in the same sentence: the weighted R^2 variant implemented here is inverted too, where their analysis rates that family as calibrated — its sign tracks cell count, and Section 3.2 therefore draws no directional conclusion from it. What it does not reproduce is their conclusion about models. They report that deep models beat uninformative baselines once the metric is fixed; on the metric their own analysis endorses, the checkpoint evaluated here does not (Section 3.3). Both results can stand: a repaired metric is a precondition for seeing skill, not a guarantee that there is skill to see. Which difference in setting accounts for the split — genetic against chemical perturbation, a discriminative regression model against an autoregressive generative one, or the strength of the per-drug baseline in this atlas — is taken up where the measurements are, in Sections 3.2 and 3.12.

2.5 Six ways a drug-response benchmark can mislead

DrEval [19] deserves separate treatment, because beyond its verdict it supplies a diagnostic vocabulary that this thesis adopts throughout. Auditing drug response prediction, it identifies six obstacles to clinical translation: non-reproducible models, data leakage producing weak generalisation, pseudoreplication, biased evaluation metrics, missing ablation studies, and inconsistent response data. Four of them shape the methodology of Section 3 directly.

Pseudoreplication, in Hurlbert’s original sense [22], is treating repeated measurements of one unit as independent observations. DrEval observes that GDSC2’s 395,025 measurements descend from only 969 cell lines and 287 drugs, so an interval computed over measurements is badly overconfident about how much independent evidence exists. Our data has the same shape — hundreds of thousands of cells resolving to a few thousand conditions across 50 cell lines — and every confidence interval in this thesis is consequently a clustered bootstrap over cell lines rather than over cells.

Biased evaluation they frame as a version of Simpson’s paradox. Because different drugs act at vastly different concentrations, most of the variance in a response dataset lies *between* drugs, so a model that learns nothing but each drug’s mean will explain most of it. Their illustration is stark: a global Pearson correlation of 0.90 alongside an average within-drug Pearson of 0.01. Section 3.2 explains why a discrimination metric evaluated *within* a comparison set is immune to this by construction.

Their treatment of splits is the part we borrow most directly. DrEval names four, each corresponding to a different application: Leave-Pairs-Out for missing-value imputation, Leave-Cell-Lines-Out for personalised medicine, Leave-Drugs-Out for drug design, and Leave-Tissue-Out for repurposing. Our holdout maps onto two of these — `unseen_combo`

is Leave-Pairs-Out and `unseen_drug` is Leave-Drugs-Out — and Section 3.11 uses their names. The mapping carries a warning we take seriously. Citing Partin et al. [23], DrEval reports that good Leave-Pairs-Out and Leave-Cell-Lines-Out performance is reachable simply by predicting each drug’s average response across the training cell lines, so a strong result on those splits demands that the per-drug average be measured explicitly rather than assumed away.

Their remedy for missing ablations is randomisation and robustness testing. The scramble ablation of Section 3.3 is such a test, specialised to the drug: hold every byte of the prompt fixed except the compound, and see whether anything downstream notices.

2.6 Connectivity mapping and signature transfer

The Connectivity Map programme [24, 25] built an entire drug-discovery methodology on a specific premise: that a compound induces a characteristic transcriptional signature stable enough across cellular contexts that querying a reference database by signature similarity is informative. The L1000 platform [25] operationalised this at scale, and its 978-gene landmark set defines the gene panel used here, intersected with Tahoe’s measured genes to give 946.

The premise — that a drug’s signature transfers across cellular contexts — is normally assumed rather than measured, and how much of it holds determines whether a conditional model has anything left to predict once a per-drug constant is accounted for. Section 3.9 measures it directly, at single-cell resolution and with the measurement noise divided out, and what comes back is a number rather than a yes or a no. On a disattenuated cross-context similarity scale the shared part is $\mathcal{T} = 0.553$ [0.514, 0.593]: a drug’s residual signature in one context recovers a little over half of itself in another, far above the matched negative controls (−0.017 and −0.008) and far below reproduction. The remaining 44.7% [40.7%, 48.6%] is not thereby shown to be biology. The compared pairs differ in cell line, in dose and in treatment well at once, so the figure is not a variance share and not a drug×cell-line interaction, and the section says so where it reports it. What the measurement delivers is a calibration of the transfer premise rather than a verdict on it, and we are aware of no prior estimate of that quantity for chemical perturbation at single-cell resolution.

2.7 Decodable, used, and the global workspace

Linear probing establishes that information is *present* in a representation, not that it is *used* by the computation, and the distinction is central to the analysis in Section 3.6. Elazar et al. [26] make the methodological case for keeping the two apart: probe accuracy licenses no conclusion about behaviour, and the way to ask whether a property is used is to remove it from the representation and measure what the model then does. Their

amnesic-probing design is the one adopted here, with the addition of a matched control, since an intervention that fails to move the output proves nothing unless a comparable intervention does.

A recent line of work gives that distinction a sharper form. Gurnee et al. [27] report that language models maintain a small privileged set of representations — which they term the *J-space*, identified by a Jacobian Lens that measures the average linearised effect of an activation on the model’s likelihood of producing a particular token — and argue it functions as a global workspace analogous to conscious access. The J-space is strikingly small: roughly ten to twenty-five concepts are active at any moment, typically accounting for under 10% of activation variance. Its defining property is causal rather than correlational, and they establish it by intervention: swapping one active workspace vector for another changes the model’s answer to match. Content in the J-space drives behaviour; content outside it may be richly represented and still inert with respect to what the model actually emits.

What this thesis takes from that account is an experimental design, and it is worth being exact about how much. The design is the swap: overwrite one candidate representation with another and ask whether the model’s behaviour follows. Section 3.6 constructs a domain analogue of exactly that test, substituting a drug’s learned representation for another drug’s and measuring whether the emitted profile moves toward the substituted drug. The account also sharpens the decodable/used distinction into something testable: a property can be recoverable from the residual stream at high accuracy and still sit outside whatever the readout consults, in which case making it more explicit need not change what the model emits — a prediction Section 3.7 tests directly and Section 3.10 bounds. What the thesis does not do is test the workspace claim itself. Nothing below identifies a J-space, measures its occupancy, or shows that a drug representation lies outside one; the substitution result is consistent with the workspace reading and equally consistent with a readout that reads the drug at low gain for some other reason, and Section 3.6 declines to choose between them. The framework is cited for the instrument it justifies and the distinction it sharpens, not as a hypothesis this thesis adjudicates.

The related account of why a model might allocate its representations this way comes from Elhage et al. [28], who show that networks reserve dedicated directions for features the computation needs on every forward pass and push the remainder into superposition, where a feature remains linearly recoverable without occupying an axis of its own. Read together, the two are a reason to measure decodability and use separately rather than to read the first as evidence for the second: a property required to emit any plausible output should sit in the causally active set, while one needed only for occasional, context-dependent composition need not. That asymmetry motivates building the instrument; it is not something the instrument returns.

2.8 Optimal transport for state matching

Predicting a treated cell from an untreated one runs into an obstacle that is not statistical but physical: sequencing destroys the cell it measures. No cell is ever observed both before and after treatment, so there is no ground-truth correspondence between the control population and the treated population — only two clouds of points in expression space, sampled from the same well at different times. Any construction of paired training data must therefore *infer* a correspondence, and optimal transport is the standard framework for inferring one.

The problem in its classical form asks for the cheapest way to move one distribution of mass onto another. Monge’s original formulation seeks a map $T : \mathcal{X} \rightarrow \mathcal{Y}$ pushing a source measure μ onto a target ν , so that $T_{\#}\mu = \nu$, while minimising

$$\int_{\mathcal{X}} c(x, T(x)) \, d\mu(x), \quad (2)$$

where c is a ground cost, typically squared Euclidean distance. This is intuitive — every source point is assigned one destination — but it is frequently ill-posed, because a deterministic map cannot split mass, and no such map exists when the target is more spread out than the source. Kantorovich’s relaxation resolves this by replacing the map with a *coupling*: a joint distribution $\pi \in \Pi(\mu, \nu)$ whose marginals are μ and ν , giving

$$\min_{\pi \in \Pi(\mu, \nu)} \int_{\mathcal{X} \times \mathcal{Y}} c(x, y) \, d\pi(x, y). \quad (3)$$

A coupling may send fractions of a source point to several destinations, which is what makes the relaxed problem always solvable and appropriate for noisy, overlapping empirical distributions. When c is the p -th power of a metric the optimal value defines the p -Wasserstein distance $W_p(\mu, \nu)$; the case $p = 2$ is the one used here, and its geodesics give a principled way to interpolate between distributions.

Solved exactly, the discrete problem is a linear program with roughly $O(n^3 \log n)$ cost, which is prohibitive at single-cell scale. Entropic regularisation makes it tractable. Adding a negative entropy term to the objective,

$$\min_{\pi \in \Pi(\mu, \nu)} \langle C, \pi \rangle - \varepsilon H(\pi), \quad (4)$$

renders the problem strictly convex, so the optimum is unique and takes the scaling form $\pi = \text{diag}(u) K \text{diag}(v)$ with Gibbs kernel $K_{ij} = \exp(-c(x_i, y_j)/\varepsilon)$. The Sinkhorn–Knopp algorithm recovers u and v by alternating matrix–vector products [29], which are cheap and parallelise well on a GPU.

The regularisation parameter ε is not merely a computational convenience, and it is the reason this framework appears in this thesis at all. It controls how sharply mass is

assigned. As $\varepsilon \rightarrow 0$ the coupling concentrates and approaches the unregularised solution, in which each control cell is matched to essentially one treated cell. As ε grows the plan blurs, and in the limit every control cell is coupled equally to every treated cell — the product of the marginals. Given a coupling, the *barycentric projection* maps each control cell x_i to the π -weighted average of its partners,

$$\hat{y}(x_i) = \frac{\sum_j \pi_{ij} y_j}{\sum_j \pi_{ij}}, \quad (5)$$

and the two extremes of ε therefore produce two familiar objects: at small ε a specific observed treated cell, and at large ε the group mean. A sweep over ε is thus a continuous interpolation between two target constructions that would otherwise have to be treated as unrelated alternatives, which is how Section 3.7 uses it.

Applied to biology, this machinery has largely been put to two uses, and this thesis makes neither of them exactly. The first is trajectory inference. Waddington-OT [30] treats the cell population at each time point as a distribution in expression space and recovers a probabilistic ancestor–descendant link between snapshots, on the assumption that cells follow paths minimising transport cost. Subsequent work relaxes the assumption that mass is conserved — cells divide and die, so the two populations need not have equal total mass — using unbalanced formulations built on the Wasserstein–Fisher–Rao metric. A separate line of work exploits the fact that when both measures are Gaussian the transport cost has a closed form, the Bures–Wasserstein distance, which splits cleanly into a shift in means and a change in covariance structure and so distinguishes a population that moved from one that merely spread out. The second use is prediction itself: a further line of work fits a transport map from control cells to treated cells and takes that map *as* the predictor. The use made here is a third one, upstream of the model. Transport supplies the correspondence from which a training target is built, and the object that then has to generalise is the language model trained on that target, not the coupling — which is why what matters below is ε , the parameter deciding which target the coupling defines, rather than the fidelity of the solver.

Two practical points govern how couplings are computed here. The ground cost must be evaluated in a space where Euclidean distance is meaningful, and raw counts are not such a space: they are sparse, overdispersed and dominated by library size. Couplings are therefore computed in the released Tahoe scVI latent [31], a variational autoencoder embedding that models counts with a negative-binomial likelihood and yields a compact continuous representation in which distances behave. And because the cost of a full coupling grows quadratically, transport is solved within matched groups — a single cell line on a single plate — rather than across the atlas, which is both computationally necessary and biologically correct, since a control cell from one cell line is not a plausible counterfactual partner for a treated cell from another.

3 The investigation, in the order it happened

A model that has learned something about a drug should answer differently when it is told a different drug. That is the whole question, and it is cheap to ask — swap the compound named in the prompt, change nothing else, and see whether the prediction moves. Asking it honestly turned out to be expensive, because the first thing the investigation had to measure was its own instrument, and the instrument was reporting near-perfect scores for a predictor that had never been told which drug it was predicting.

What follows is the investigation in the order it happened. Every instrument here was built in response to the measurement before it, and that order is kept rather than tidied away: a question, the tool that can answer it, what the tool showed, and what that showing forced next. Design choices appear as the failures they avoid rather than as best practice — nearly every one was forced by a specific way the obvious alternative breaks on this data. Definitions appear where they are first needed. Dead ends stay in, including three substitution designs that failed and two printed claims that had to be withdrawn, because the confound that killed each one is the reason the surviving design looks the way it does. Every section that reaches a verdict closes with its central claim, set off from the text, in language that could be shown wrong. One section closes without one, and deliberately: §3.14 asks which span of the prompt the model reads, cannot answer it, and ends in the open question rather than in a claim block that would give an unresolved measurement the typography of a finding. Two conventions hold throughout: every number is quoted with its interval where one exists, and “not established” and “shown to be zero” are different statements — the difference is load-bearing in several places.

The findings, stated before the apparatus that produced them. First, the field’s standard perturbation metric is saturated by a predictor that carries no information at all, so the instrument was repaired before the model was measured. Second, on the repaired instrument the model is drug-blind in the expression frame: it returns the same answer whether told the true drug or a different one, and a baseline that copies the untreated cell matches it exactly. Third, the blindness is not an encoding failure. Drug identity is written into the model’s activations, increasingly explicitly with depth, and the read-out barely consults it; consistently with that, a sweep of denoised training targets — which changed what the model had to produce but not how the drug was encoded — produced no separation from its scrambles under the test available at the time. Fourth, re-encoding the training target as the drug-specific residual makes the model demonstrably sensitive to the drug in its prompt. That is the project’s first non-null drug effect, though a budget-matched re-run later showed it to be first in sequence rather than the only target construction that produces one: one rung of that earlier sweep had been read as inert on a truncated generation budget. Fifth, the sensitivity does not become reliable

prediction on held-out treatment wells. A lookup table of training averages scores above the within-well split-half precision reference itself, covering 115.7% of the distance from chance to that bar — a property of the bar, not a headroom estimate — where the model covers 13.9%, and the chapter’s argument is that dissociation, between prompt sensitivity and biological fidelity. Sixth, for a drug the model has never seen, the prior question is whether any shared property carries a usable signal at all, and the answer is a lead rather than a closure: of the three channels Tahoe’s own metadata offers, only protein target clears the relevance margin set in advance, and only with the cell line as the clustering unit rather than the drug, on 18 annotated drugs of which two are held out of training. Mechanism and chemical structure are inconclusive, not dead. That is the one place the chapter points forward. Figure 1 is the map.




Figure 1: The investigation in nine beats, each told with the instrument it required. The chapter’s remaining two subsections are not beats: §3.1 is the apparatus and §3.15 is the summary table. The ninth beat is the one the chapter does not close.

3.1 The atlas, the model and the rules of measurement

This section is apparatus, and it comes first because on this atlas the apparatus decides what may be claimed. Four things have to be settled before any score means anything — what is being measured, on what population, against what reference, and with what uncertainty — and Tahoe answers none of them by convention. Each is settled below by a measurement rather than a stipulation, and each of them narrows a claim later. The pitfalls recorded here are kept for the same reason the dead ends are kept elsewhere: each is why the surviving construction looks the way it does.

The corpus is Tahoe-100M [1]: drug-treated single-cell transcriptomes across cancer cell lines, with plate-matched DMSO controls. Metadata supplies cell line (Cellosaurus ID), drug, dose (`drugname.drugconc`), plate, and a fine-grained mechanism-of-action annotation (`moa-fine`). All analysis is restricted to the $G = 946$ genes in the intersection of the L1000 landmark set [25] with Tahoe’s measured genes, library-normalised to counts

per 10,000 and transformed as $\log_{10}(1 + x)$. The panel is a budget choice rather than a biological one: a cell is represented as a token sequence, the model’s context is 8192 tokens, and a prompt must carry an instruction line and two cells, so the full transcriptome does not fit. Neither step changes the rank sentence, and it is worth being exact about that because it is easy to assume otherwise: library normalisation multiplies every gene in a cell by one constant and the logarithm is monotone, so within a cell the ordering — and therefore the sentence — is invariant to both. What they change is everything downstream of a single cell. Counts scaled by depth are not comparable *across* cells, so the pseudobulk means, the drug-specific residual that is a difference of two such means, and every cosine and Euclidean quantity computed on them all depend on the normalisation even though the sentences do not. Sequencing depth does still shape a sentence, through which genes are detected at all rather than through their order, and no rescaling recovers a transcript that was not observed.

One representational pitfall is worth recording because it is silent. Tahoe’s per-cell **genes** field holds vocabulary token identifiers, not positional indices into the gene metadata table. Mapping them positionally scrambles gene identity quietly: the values exceed the metadata row count, so the error surfaces only as an out-of-bounds index at the tail. All expression construction here maps `gene_metadata['token_id']` to gene symbol to panel index.

Validating that mapping needs more care than it first appears. An earlier draft offered a PCA cell-line probe on the resulting panel, which reaches 0.864 against a chance rate of 0.02. That probe cannot establish gene identity, and the reason is worth stating because the same mistake is easy to make elsewhere: permuting the gene columns is an orthogonal transform, so it preserves every pairwise distance between cells and leaves PCA and classification accuracy exactly unchanged. A panel with every gene label scrambled scores 0.864 too. The probe shows the panel retains cell-line-discriminative structure, which is worth knowing, but it is silent on which column is which gene.

What does discriminate is abundance, because abundance is a property of the gene rather than of the panel. On the measured profiles, **GAPDH** is the single most abundant of the 946 genes and **RPS6** the second; across the canonical high-expressors carried by this panel the median abundance rank is 19.5, against 473.5 expected if the index were permuted ($p = 0.00015$, 20,000 draws; `gene_identity_check.py`). A permutation would erase exactly this signal, which is why it is the test that was needed.

The physical layout of the atlas drives every inferential choice in this chapter, so it is stated before anything that depends on it. Tahoe applies one drug at one dose to one *well* containing roughly 49 cell lines simultaneously: across the three plates this project caches (`plate4`, `plate5`, `plate6`) there are 289 treatment samples, a mean of 48.6 cell lines per well (median 49, range 43–50), with 286 of 287 distinct (drug, dose) treatments occurring in exactly one well. Those figures are properties of this cache and not of the atlas: Tahoe

includes further plates, among them a replication plate, and where the text below says a treatment occurs in one well it means one well *here*. No cross-plate replicate of a (drug, dose) treatment exists within this cache, which is what the inference rules below depend on; conditions that do appear on more than one of these plates differ in dose. Each drug is assayed on its own plate or plates, with plate-matched controls. Three consequences run through everything below. Conditions from one well share a physical assignment, so they are not independent — this sets the inference rules below. Drug and plate are partially confounded, so discrimination must be measured within plate (§3.3). And a (drug, cell line) pair cannot be held out cleanly, because the pair’s well remains in training through its other ≈ 48 lines (§3.11).

Three units of analysis run through everything that follows, and confusing them is the easiest way to misread a number. A *condition* is a (drug, cell line, plate, dose) tuple, the level at which responses are scored. The physical treatment assignment is the *treatment sample*, one treated well, so a condition is an observation nested inside an assignment rather than an assignment. Plates are numbered within cell line rather than globally, so the batch unit is the pair (cell line, plate) and never the plate label alone. A *pseudobulk* is the mean expression profile of a group of cells; a *control pseudobulk* is the same object over plate-matched DMSO cells. Dose is carried as a molar concentration resolved from `drugname_drugconc`, not as a unitless float: an earlier pipeline kept only the numeric part, which collides 1 μ M with 1 nM. A direct audit of all 289 treatment samples found 289 usable molar doses with zero unit collisions: every drug in this atlas states its doses in a single unit, so the defect was a demonstrated possibility that did not materialise. Combination samples are dropped rather than analysed as their first component.

Table 1 states the quantities every n in this chapter should be located against. Two populations are used and they are not interchangeable. The broad characterisation of the task (§3.3) streams a wide sample of the atlas, over more plates and more drugs than the three cached above, so the drug and condition counts in its block below are counts of that sample and not of the cache; the residual arm (§3.10 onward) works from a smaller cached subset built once and reused, so that every arm sees identical cells.

The upper block of that table describes Tahoe-100M rather than this project, and it is worth reading as a result and not only as scaffolding, because it holds for anyone else scoring predictions on this atlas. Across the 5,014 of 6,628 conditions whose (cell line, plate) group carries plate-matched DMSO cells — the only ones a control-referenced test can reach — the median has no differentially expressed gene at BH $q < 0.05$ and a signal-to-noise ratio of 0.75 against its own within-well split-half noise, and 72.7% are active on the whole profile: the typical condition carries a real response measured with too few cells, not no response, and §3.3 is where that distinction is settled rather than asserted. Over all 6,628, under a per-condition identifiability bar — a condition’s own split-half NIR against its plate-mates clearing 0.8 — 46.2% qualify, and 78.7% have a plate-mate

lying closer to them than their own second half does. Identifiability is a property of the pairing rather than of the compound: over the 227 drugs assayed in at least three cell lines, the median drug’s own split-half reference spans 0.83 between its most and least identifiable context (`drug_biology_atlas.json`). A benchmark built on this atlas that does not report those figures is quoting accuracies over a population of which barely half is identifiable in principle, and a per-drug leaderboard on it is ranking pairings.

Three evaluation tiers, defined by what the training set contained, are used for the broad characterisation. *Tier 1* holds conditions whose (drug, cell line) pairing was seen in training, so it measures fit rather than generalisation. *Tier 2* holds conditions whose drug never appeared, which is Leave-Drugs-Out in the vocabulary of Section 2.5. *Tier 3* holds pairings whose drug and cell line were each seen but never together, which is Leave-Pairs-Out. Tier 3 is small (99 conditions) and is reported only where it is adequately powered; the residual arm builds its own, larger holdout, and the unit that holdout is keyed at turns out to be one of the most consequential choices in the chapter (§3.11).

The model side is equally compact. A cell is written as a *cell sentence*: its expressed panel genes in descending order of expression, terminated by `[END_CELL]`. The prompt supplies cell line, drug, dose, mechanism, and the untreated control cell’s own sentence; the response is the treated cell’s sentence. Two tokens, `[END_CELL]` and — for the residual arm — `[DOWN]`, are registered as atomic special tokens; left unregistered, `[DOWN]` splits into sub-words and the up/down boundary ceases to be a clean symbol. The backbone is `C2S-Scale-Pythia-1b-pt` [3], a C2S-Scale checkpoint built on the Pythia-1b decoder [4], cold-started for every arm from the same base checkpoint with identical hyperparameters (1 epoch, learning rate 10^{-5} , gradient accumulation 16, maximum length 8192, bf16), so that in each experiment the training target or objective is the only variable. What that checkpoint has already seen bears on how the held-out splits should be read. It is pretrained on over 50 million human and mouse cells from 825 single-cell datasets in CELLxGENE and the Human Cell Atlas — an atlas corpus, carrying no chemical or genetic perturbation data — and its model card lists cell and tissue prediction, conditioned generation and gene-set conversion among its capabilities, but not perturbation prediction. The base model therefore reaches our fine-tuning having never seen a drug-treated cell. Two channels nevertheless stay open. The decoder underneath is Pythia-1b, pretrained on natural-language text in which compound names, targets and mechanisms occur; and every prompt supplies the drug’s `moa-fine` mechanism annotation next to its name. “Held out” therefore means held out from Tahoe fine-tuning, with mechanism supplied, and every unseen-drug result below must be read under that narrower scope.

Vocabulary. Six terms carry the argument and are collected here before first use. Two *measurement frames* run through the chapter and are never interchangeable. In the *expression frame* the prediction is a cell profile, scored against other drugs on the same plate.

In the *residual frame* (§3.8) the prediction is a drug-specific signature, scored against other drugs in the same cell line. The *comparison set* J — the conditions a prediction is ranked against — is what fixes chance at 0.50 in both frames, and it is declared explicitly at every score below. The *generic* is the drug-agnostic mean response subtracted to form the residual: operationally, the mean shift over the other drugs at the stated scope, and conventionally read as the stress and cell-cycle programme common to all compounds — a reading used here for intuition and nowhere tested. It is the most important object in the chapter, for a measured reason: it carries 39.0% of the perturbation shift’s energy (§3.9), so subtracting it does not clean the target up but replaces it, and every number quoted in the residual frame is a score on what is left. The *residual target* is what remains once the generic is removed. It is built *split-before-fit* (the holdout is fixed from metadata before any statistic is estimated) and *leave-one-drug-out* (no condition’s generic contains its own drug). The *scope* of the generic, `cell_line` or `plate`, is the set of drugs the mean is taken over, and §3.8 shows the choice is consequential. *Control-copy* is the baseline that predicts the untreated cell unchanged; it carries zero drug information by construction. The residual arm’s split has three parts — `train`, `unseen_combo` (an unseen *well* of a seen drug) and `unseen_drug` — and samples conditions under an explicit per-split *quota*, because a uniform draw reproduces the pool’s proportions and starves the held-out splits (§3.11). Five training *arms* (model variants differing only in target or objective) are used; they are enumerated in Table 20 of the Appendix. Where a claim below is multiplicity corrected, the family is those five arms by up to seven depths.

Ceilings. Almost every negative result in this chapter is stated as a distance from a *ceiling*, and the word names one estimator, not nine: a *within-well split-half precision reference*, the score obtained when a disjoint cell subset of the same treated well is substituted for the prediction. It is not a replicate ceiling. What varies is the frame, the comparison set and the cell budget, and Table 2 says which is which. Two properties matter for how the ceilings can be read. They are *split-half* estimates rather than biological replicates, because this cache places 286 of 287 treatments in exactly one well — for those, no biological replicate exists in this cache. And a split-half ceiling is a within-well sampling-precision estimate, which is optimistic relative to true biological reproducibility; the two coverage figures quoted against the residual-frame ceiling in §3.12 inherit that optimism.

Rules of measurement. Two rules apply to every result below, plus one statement of provisionality. The first concerns uncertainty. Conditions are not independent, and they are not independent in two crossed ways at once: two conditions from the same treated well share a physical assignment and a batch, and two conditions from the same cell line share the biology. Treating conditions as independent gives intervals that are too narrow

— here, about 1.3 to 1.5 times too narrow — and neither grouping alone captures the dependence. The tempting shortcut is actively misleading: clustering on the assignment unit alone uses the larger cluster count (289 wells against ≈ 50 lines) and would tighten every interval while sounding more rigorous. Every interval on a *mean* in this chapter is therefore two-way cluster-robust (Cameron–Gelbach–Miller),

$$V = V_{\text{line}} + V_{\text{well}} - V_{\text{intersection}}, \quad (6)$$

where the intersection of the two groupings is the single condition, so $V_{\text{intersection}}$ is the ordinary independent-sampling variance. When the estimator goes non-positive in a finite sample, the wider one-way variance is used and the branch recorded. A cluster-robust variance is asymptotic in the number of clusters, not observations, so the reference distribution is Student’s t with degrees of freedom one less than the smaller cluster count. This matters where an arm spans only 28 wells — the `unseen_drug` arm, whose reference is t_{27} rather than the normal, widening its interval by 4.7%. For statistics computed over pairs of conditions, such as the transfer coefficient, neither grouping partitions the observations, a pair belongs to two nodes at once, so those intervals use the Fafchamps–Gubert dyadic estimator, under which two pairs are dependent whenever they share either member. Throughout, an estimate is quoted as value [lower, upper], a two-way cluster-robust 95% interval unless stated otherwise. This is the structure DrEval [19] identifies in this field more broadly, citing Hurlbert’s pseudoreplication [22]: hundreds of thousands of cells resolving to a few hundred conditions over 50 cell lines.

The second rule concerns generation validity, and it exists because the generation budget turned out to be a claim-level variable rather than a runtime setting. A 600-token cap truncated 26% of generations in an early run, and it did so silently — not because the truncated outputs were intact, but because nothing measured them: a generation cut at the cap never reaches `[DOWN]` and is scored with its whole down-block missing, and no completion rate was quoted beside the score to say so. What it cost was a conclusion. The optimal-transport arm of §3.7, scored at that cap, gave a gap against its neutral comparator of +0.0144 [−0.0051, +0.0338], spanned zero, and was printed as inert; re-run at the 1400 tokens the residual arm had, it gives +0.0292 [+0.0144, +0.0441], and that printed negative had to be withdrawn (§3.10). The same cap had already halved the residual arm’s own effect. The general form is not specific to cell sentences: wherever a paired comparison is scored on variable-length generations, the token budget is part of the claim, an under-budgeted arm can be reported as a null, and a null quoted without its budget cannot be read. Hence the rule — every generation-based evaluation reports the `[END_CELL]`/`[DOWN]` completion rate, the fraction of emitted symbols that are valid panel genes, and the duplicate rate, before any score is quoted, together with the budget it was generated under. The generation configuration itself is fixed and is recorded here

because the intervals do not cover it: sampling at temperature 0.8 with top- p 0.9, four samples per condition in the residual-frame evaluations of §3.10 and their control re-runs, and eight in the expression-frame evaluations that feed §3.3 and Table 8. Every contrast is additionally generated with a fixed per-(condition, arm) seed, so the two arms of a paired comparison are not drawn from different points of one global stream.

The provisionality is stated once. Every residual-frame number in this chapter is taken from one generation run and its companion audits (`re_v3` with the scramble-stratum audit, the matched variance decomposition, the two calibration runs and the channel gate). Numbers from the two earlier runs are not reconciled with these and are not quoted except where a superseded estimate is named as such and the conclusion drawn from it withdrawn, because between runs the sign of several estimates changed and not only their magnitude. Two components carried defects found after that run and have since been repaired and re-run: the DRF calibration bootstrap (§3.2, both arms) and the channel-gate partner pool (§3.13). Both components are reported here on the corrected runs, and neither leaves an arm awaiting one.

3.2 The metric is the first result

The investigation begins with the scoreboard, because every later claim depends on it. That the metric audited here is the field’s default is a premise the rest of the chapter leans on, so it is worth pinning down. A correlation between predicted and true change from control, restricted to the genes that actually move, is among the scores the transcriptomic benchmarks in the dispute of §2.4 report; control-referenced correlation is among the widely used measures Miller et al. [20] find poorly calibrated; and DrEval [19] lists biased evaluation metrics among six standing obstacles in this field. What follows is the top- K member of that family. Writing x for the control profile, y for the treated truth and $\hat{y}$ for the prediction, let

$$S_K = \arg \text{top-} K_g |y_g - x_g| \quad (7)$$

be the K genes with the largest true change. Then the differential-expression delta-correlation is

$$\text{DE-}\Delta r = \text{corr}((\hat{y} - x)|_{S_K}, (y - x)|_{S_K}). \quad (8)$$

We report $K = 50$ with Pearson correlation throughout; $K \in \{20, 100, 200\}$ and the Spearman variant were computed and move the numbers by less than 0.02 except at $K = 200$, where the top- K set starts to include genes that do not move.

What does this metric reward? The way to find out is to score a predictor that contains nothing: every gene placed at the middle rank, carrying no information about drug, cell or biology. It scores $\text{DE-}\Delta r(K=50) = 0.9999$ [0.99991, 0.99992] — above the two-real-cell noise ceiling (0.76) and above the model (0.73). The ordering is inverted:

less information yields a higher score, which is the signature of an artefact rather than of skill. The same predictor’s partial-DE- Δr and panel- τ are undefined rather than merely poor: its shift from control is identically zero, so there is nothing left to correlate once the control is regressed out.

What travels beyond this atlas is the audit rather than the verdict. The predictor is defined from the metric alone: it needs no access to a model, a training set or an atlas, and it exploits a single property of the data it is run on — that treated cells regress toward a shared response. Whether that property holds strongly enough elsewhere to invert the score is an empirical question this cache cannot answer, but any result reported on this metric family can be checked the same way, at the cost of a few lines of code.

The cause is visible in the definition, and three properties of it, each innocent in isolation, compose into the failure. First, S_K is selected using the truth, so every predictor is scored on the genes the condition actually moved: the exam questions are chosen from the answer key. Second, treated cells in this atlas regress toward a shared response, so a prediction that reverts toward the population centre correlates with $y - x$ by construction. No drug knowledge is required to score well, only knowledge of where the centre is. Third, the fix is available in the same breath: partial-DE- Δr regresses the control profile out of both arguments before correlating, which removes the reversion route and is therefore the honest version of the same quantity. Removing the control collapses the raw 0.95 to ≈ 0.26 for the two baselines that were scored, and that residue is matched by a baseline that performs no fitting at all. The model’s own partial-DE- Δr was never computed,³ so the comparison is anchored instead to panel- τ , Kendall’s τ between predicted and true rank over the full 946-gene panel, which selects no gene subset and is measured for every arm. The anchor is used only to rank arms against one another on a fixed scale; no absolute claim rests on it, because the calibration audit below rules panel- τ inadmissible for absolute claims. It tells the same story: **revert_mean** 0.268, ridge 0.265, model 0.257 (tier 2, **worst** convention). The response decomposes into control reversion (an artefact), a generic drug programme (≈ 0.26 , drug-agnostic, trivially matched), and a drug-specific component that no predictor captures.

A cell sentence lists only expressed genes, so a predicted profile leaves some panel genes unranked, and any rank-based metric needs a convention for the absent tail. Two were used throughout and reported side by side: **worst** assigns every absent gene the last rank G , and **francesca** assigns them a fixed middle rank $G/2$. The choice is immaterial — spike-in DE- Δr reads 0.952 against 0.954 under the two, and no baseline or ceiling in this chapter moves by more than 0.01 — so all figures are quoted under **worst** and the convention is not mentioned again. That comparison was a one-off diagnostic and was not written to an artifact, so it is quoted as a check that was run rather than as a

³It requires a GPU pass that was not run, so the cell is marked rather than filled with the expected value.

reproducible result.

The metric this investigation relies on instead is the Normalised Inverse Rank (NIR), introduced by Wu et al. in PerturbBench [12], a benchmark suite for cellular perturbation models (there called the “transposed-rank” metric), and identified by Miller et al. [20] as one of the few consistently calibrated metrics across 14 Perturb-seq datasets. Figure 2 shows the construction. For a prediction $\hat{y}$ of condition i , with truths $\{y_j\}$ over the comparison set J ,

$$\text{NIR}(\hat{y}, i) = \frac{1}{|J|} \sum_{j \in J} \left[\mathbf{1}\{\sigma(\hat{y}, y_i) > \sigma(\hat{y}, y_j)\} + \frac{1}{2} \mathbf{1}\{\sigma(\hat{y}, y_i) = \sigma(\hat{y}, y_j)\} \right], \quad (9)$$

the fraction of the other conditions in J whose truth the prediction resembles less than its own, with ties counted at one half. Here σ is a *similarity*: higher means closer. Two of the three frames below score with a distance rather than a similarity, and there σ is its negation — $\sigma(\hat{y}, y) = -\|\hat{y} - y\|_2$ for `nir_expr` — which is what reverses the inequality in the implementation. The construction is otherwise identical across frames; only σ and J change, and both are declared wherever a figure is quoted. The comparison set is declared at every score: conditions are grouped by (cell line, plate) in the expression frame and by cell line in the residual frame, and the two are not interchangeable: widening J across plates moves a drug-agnostic baseline from 0.161 to 0.399 (§3.3). Two properties do work that no other candidate does. Chance is 0.50 by construction, not empirically: it is a statement about J and nothing else, so “this model is at chance” is a statement rather than a comparison against a fitted null, and that is what makes the negative results in this chapter quotable without a control arm. The condition it rests on should be named, because it is an assumption and not a theorem: the target’s own truth y_i must be exchangeable with the comparison truths $\{y_j\}$, so that an uninformative prediction ranks y_i uniformly among them. That holds within every arm as scored here, since each arm is compared against the same truth set it is scored on. One asymmetry survives *between* arms and is stated where the ceiling is used: the split-half reference is scored against half-sample truths while every baseline is scored against full-sample truths (`residual_eval.py`), which makes the reference the harder of the two and the comparison conservative rather than flattering. And because the generic response is present in every condition’s truth in equal measure, it cancels in the comparison: NIR is a discrimination metric, and only the drug-specific difference decides it, which is precisely the component the reversion artefact exploits. The tie term is not cosmetic. In residual space (§3.8) a predictor of “the average drug response” is exactly the zero vector, every similarity against it is zero, and a strict inequality scores it 0.000 rather than the 0.500 it deserves.

Our instruments compute NIR under three similarity functions and store all three: `nir_expr`, Euclidean distance in expression space; `nir_rank`, the same construction over rank vectors; and, in the residual frame of §3.8, cosine similarity between signed rank-




Figure 2: **NIR asks whether a prediction resembles its own condition's truth more than it resembles the truths of the other conditions in the comparison set J .** Own truth ranked first gives NIR $\rightarrow$ 1; ranked at random gives 0.5; and a predictor whose similarities are all identical — the zero vector, which is what “predict the average drug response” becomes in residual space — scores 0.5 by the half-tie rule, not 0.

discounted residual vectors. Every *expression-frame* NIR figure in this chapter is `nir_expr`. That choice is load-bearing, and it was not made by the DRF audit below: `nir_rank` never entered the graded family. It is excluded by a cheaper admissibility screen that makes grading it pointless — a held-out disjoint sample of a condition's own cells must be retrievable above chance. Under `nir_rank`, on the same within-plate tier-2 run that later delivers the drug-blindness verdict ($n = 606$, `nir_sameplate.json`), every arm sits below the 0.50 chance line and the noise ceiling is among them: ceiling 0.380, control-copy 0.337, model 0.313, drug-agnostic linear map 0.285, mean baseline 0.070. A metric on which a disjoint sample of a condition's own cells cannot be retrieved above chance is not measuring what its name suggests, and no DRF was needed to see it. Nor does the exclusion contradict the calibration literature's preference for rank-based metrics: the rank in NIR's construction is a ranking of *conditions* by similarity, not of genes, and the similarity itself is computed in expression space. Two measurements therefore stand behind the metric this chapter uses, and they ask different questions — whether the instrument registers a true positive at all, and how much of the range between a drug-agnostic baseline and a noise ceiling it spans. All three variants are stored; the graded one is named here so the artifacts cannot be misread.

Choosing NIR over its rivals is itself a measurement, made with the framework of Miller et al. [20], which grades each metric rather than each model by asking how well it separates a positive from a negative control. The Dynamic Range Fraction is

$$\text{DRF}(m) = \frac{m(\text{pos}) - m(\text{neg})}{m(\text{perfect}) - m(\text{neg})}, \quad (10)$$

on a scale where 0 is the drug-agnostic baseline and 1 a perfect predictor. Here pos is an interpolated-duplicate noise ceiling — a technical duplicate interpolated toward the mean baseline with per-gene weights given by differential-expression significance — and neg is a drug-agnostic leave-one-out mean. A negative DRF means the metric scores the uninformative baseline above the noise ceiling: it is inverted, and rewards the generic

gene ordering rather than the drug. Three hardening choices matter for how the result can be read. Tied expression values receive average ranks, because breaking ties by array index makes a tied gene’s rank depend on its position in the panel. The DRF ratio is recomputed whole inside every bootstrap resample, because its denominator is estimated from the same data as its numerator and holding it fixed understates the interval. And the family of five metrics is controlled simultaneously with Holm’s procedure, because the claim “NIR alone points forward” is a claim about all five at once. Those three choices govern how the point estimate is formed. A fourth — what the bootstrap resamples — governs whether the interval around it means anything, and in the stored runs it was not made: each of the all-plates run’s 1,820 rows is one drug within one cell line, and the resampler treated every row as its own cell line, against the 25 that exist (the same-plate run carried the identical defect over its own rows). The one-sided p was read off the observed resampling distribution rather than one with the null imposed, so it degenerates to the floor $1/2001$ wherever $\text{DRF} > 0$ and to 1 wherever $\text{DRF} < 0$, and no metric in the family was in fact tested. Both are repaired in the analysis code and both arms have been rerun under the repair. The same-plate rerun also fixed a mismatch the first run exposed: its cluster was the (cell line, plate) pair the comparison is grouped within, but plates nest inside cell lines and share their biology, so resampling pairs independently repeats the original error one level up. The cluster is now the cell line in both arms, and the grouping unit and the cluster unit are reported separately because they are not the same thing.

Under the hardened protocol, NIR is the only one of the five whose DRF is positive. In the all-plates run, $\text{DRF}(\text{NIR}) = +0.635$, 95% CI $[+0.604, +0.663]$, clustered on the 25 cell lines that the 1,820 drug rows are nested in rather than on the rows themselves. The four rivals are all negative and all exclude zero on the same clustering: -0.163 (`weighted_r2`), -0.319 (`panel- $\tau$`), -0.650 (`spearman_expr`) and -0.686 (`DE- $\Delta r$`). Against a null imposed by centring the resampling distribution at $\text{DRF} = 0$, no draw of 2,000 reaches the observed value for NIR, so the one-sided p is at the resampling floor and reads 0.002 after Holm correction across the five metrics. That the p is at the floor is now a statement about the effect rather than about the resampler: the point estimate sits some forty cluster-robust standard errors from zero. Restricting swap partners to the same plate (586 drugs over 44 cell-line $\times$ plate groups spanning 27 cell lines) pulls every value toward zero and changes no sign: NIR $+0.545$ $[+0.516, +0.569]$, against -0.104 , -0.251 , -0.500 and -0.515 in the same order, each excluding zero. The two arms are matched on the thing the comparison is about — 27 cell lines within plate against 25 across plates — which an earlier configuration was not: capping the run at 25 *groups* admitted only 18 distinct cell lines. That narrower run gives NIR $+0.519$ $[+0.479, +0.546]$ over 335 drugs, so the figure is not an artifact of how many plate-groups are admitted. What the argument uses is the sign pattern, which is identical in every run; no ranking among the four negative metrics is claimed. The inversion is a property of the data, not a leak: a leave-one-out correction barely moves

the numbers, since real cell lines carry 20–204 drugs, and it survives holding the plate constant.

This replicates the methodological core of Miller et al. and then sharpens it. Metric choice dominates, and the control-referenced correlation they flag is the worst performer in both arms ($\text{DE-}\Delta r$, -0.686 across plates and -0.515 within). But their calibrated set does not transfer intact to drug perturbation. Two of the four rivals audited here are themselves rank-based — $\text{panel-}\tau$ over the gene panel and `spearman_expr` — and both come out inverted in every calibration run computed for this chapter; `weighted_r2`, the counterpart of the weighted $R^2\Delta$ [21] in their calibrated set, is negative in both audited arms, though it sits closest to zero and its sign tracks cell count, so it carries no directional claim of its own. What survives the port is not the rank-based family but the single *discrimination* metric within it: NIR is the only one of the five that scores a prediction against a set of other conditions rather than against its own truth, and it is the only one that points forward. Being rank-based is not what protects a metric here; refusing to score anything the generic response can supply is. That is the statement worth carrying elsewhere, and it is not the one a clean replication would have produced.

Miller et al. then conclude that under calibrated metrics deep models outperform uninformative baselines; §3.3 evaluates on the calibrated metric and finds the model at chance. Two differences in setting are candidates for the discrepancy — genetic versus drug perturbation, and a discriminative regression model versus an autoregressive generative one — and §3.12 supplies a third and more structural possibility. Adopting their framework is also deliberate in a second sense: evaluating our model on the metric their analysis endorses means the drug-blindness result cannot be dismissed on the grounds their own paper raises.

Everything above grades `nir_expr` only. The residual chapter scores a different σ over a different comparison set, and a metric is not calibrated by having a relative calibrated, so the residual frame carries its own DRF — a second calibration result, on a different frame and against different controls, rather than a formality discharged in passing. Computed the same way, with a split-half replicate of the condition as the positive control and the drug-agnostic arms already in the evaluation as the negatives, it reads $\text{DRF} = +0.704$ [$+0.667, +0.736$] against `control_copy`, clustered on the 42 cell lines. It does not depend on which negative is chosen: a random vector and the neutral scrambled partner both give $+0.707$, and all three negatives sit at chance as a matter of measurement ($0.506, 0.501, 0.501$) rather than of stipulation. It is also reached under the more conservative of the two positive controls — a raw split-half replicate rather than the denoised interpolated duplicate the expression-frame protocol uses — so the residual frame is the better separated of the two ($+0.704$ against $+0.635$) despite being graded against the harder reference, which is the direction the re-encoding was intended to move. One arm is reported for completeness and carries no weight: `generic` scores exactly 0.500 with zero variance




Figure 3: **Under the standard metric, less information yields a higher score.** (a) A predictor carrying nothing about drug, cell or biology scores $\text{DE-}\Delta r = 0.9999$, above the two-real-cell noise ceiling and above the model. Tier 2, $n = 400$ conditions. (b) Grading the metrics rather than the models: only NIR points forward. The panel shows the same-plate run (586 drugs over 44 cell-line $\times$ plate groups), the confound-free version, in which NIR reads +0.545 against -0.104 (`weighted_r2`), -0.251 (`panel- $\tau$`), -0.500 (`spearman_expr`) and -0.515 (`DE- $\Delta r$`); the all-plates run (1,820 drugs over 25 cell lines) gives +0.635 for NIR, 95% CI [+0.604, +0.663] clustered on cell line, with the four rivals between -0.163 and -0.686 and all excluding zero. Both arms are clustered on the cell line — 25 across plates, 27 within — and every interval shown is a percentile cluster bootstrap with no small-sample correction (see text). `weighted_r2` carries no directional claim: its sign tracks cell count.

because the residual is defined by subtracting it, so using it as the negative would make DRF a rescaling of the ceiling and test nothing.

One exception to the rule of §3.1 is worth flagging where the numbers are, not in a footnote. DRF is a ratio of means rather than a mean, and the two-way estimator above is written for a mean, so both DRF intervals — expression frame and residual frame — are *one-way* percentile cluster bootstraps that resample whole cell lines and recompute the ratio inside each draw. That respects the cell-line dependence and the ratio structure, and it ignores the crossed well dimension. Wells are crossed with lines here, so the interval is robust on the cell-line axis alone and is narrower than a fully crossed one would be. The NIR separation it grades is far too large for that to change the sign, but the caveat belongs beside the number.

Two questions about the instrument remain, and they are not the same question. The first is whether the discrimination task is doable at all at pseudobulk resolution: if a disjoint sample of a condition’s own cells cannot be told from another drug’s cells, no repair to the metric would help. A spike-in forced choice answers it. Within one plate, a reference pseudobulk drawn from drug A’s cells must be scored more similar to a disjoint sample of A’s own cells than to a sample of drug B’s, over many resamples, with chance at 0.50; contaminating B’s population with a growing fraction of A’s cells titrates the difficulty, so the curve of accuracy against contamination characterises each metric’s sensitivity. At zero contamination three of the six metrics scored separate the two populations essentially perfectly (panel- $\tau = 1.000$, DE- $\Delta r = 0.995$, top- $n \tau = 0.987$, plate held constant), and at full contamination those same three return 0.501–0.502, so the chance value arrives here by measurement and not by stipulation. The whole-profile similarity scores on the same instrument are much weaker even at zero contamination — cosine 0.740, Spearman 0.680, cosine on the shift 0.620 — which is worth registering, because a whole-profile similarity is what NIR uses in the expression frame. The task is not the bottleneck; how much of it a given σ recovers is a separate matter.

The second question — whether NIR itself registers a true positive — the spike-in does not answer, since NIR was not among the metrics run on it. The evidence is the positive control inside the DRF audit, on the same two runs that produce the figures above: under NIR the interpolated duplicate scores 0.802 against the drug-agnostic mean’s 0.457 across plates, and 0.668 against 0.272 within plate. In the residual frame a split-half replicate reaches 0.854 against negatives at 0.501–0.506. None of these is the reference the drug-blindness verdict is measured against, and it is worth saying which is: on the tier-2 within-plate benchmark ($n = 606$) the within-well split-half precision reference reaches only 0.576, above chance but not far above it, and that — rather than 1.0 — is the bar §3.3 states the model against.

THE CLAIM. The field’s standard score is saturated by a predictor carrying no information ($\text{DE-}\Delta r = 0.9999$). Of the five candidates audited, NIR is the only one with a positive DRF ($+0.635$, $[+0.604, +0.663]$ clustered on the 25 cell lines, Holm $p = 0.002$; $+0.545$ $[+0.516, +0.569]$ over 27 cell lines with same-plate partners, Holm $p = 0.002$), and the other four are inverted rather than merely uninformative: each ranks the drug-agnostic baseline above the noise ceiling, with intervals excluding zero in both audited arms. Under NIR the positive control is retrieved well clear of the negative (0.802 against 0.457 across plates, 0.668 against 0.272 within), so the instrument registers a true positive. Every drug-use claim in the rest of this chapter is therefore made on the calibrated metric.

3.3 The model is drug-blind

With a ruler that cannot be gamed in hand, the investigation can ask the question it exists to answer: does the model use the drug at all? The answer is no, and it is stated before the instruments because the instruments matter only as the objections they close. On the calibrated metric, within plate on tier 2 ($n = 606$), the model scores 0.498 against a within-well split-half precision reference of 0.576 . It is matched by predictors that know nothing about the drug: a drug-agnostic linear map scores 0.500 , a control-copy 0.504 , the global mean 0.180 . And on the identifiable stratum, where the task is provably winnable, a zero-drug-information control-copy reaches 0.766 against the model’s 0.768 .

Four instruments produce this verdict, and each exists because a specific objection would otherwise be fatal. They are the benchmark itself (the model against drug-agnostic baselines on the calibrated metric, within plate); the scramble ablation (the only manipulation that isolates drug use); forced-choice grading (a truth-based check whose chance value is fixed); and output invariance (a truth-free check that no scoring choice can confound). The benchmark is the verdict above; the other three are built and reported below.

One measurement rule comes first, because the verdict depends on it. Drug and plate are partially confounded by the atlas’s design (§3.1), so discrimination grouped by cell line alone lets the batch signature stand in for drug identity. Measured directly, by scoring the same tier-2 conditions with and without the plate held constant:

The mean baseline is the clearest demonstration: it contains no drug information whatsoever, and it scores 0.399 when the comparison set may span plates against 0.161 when it may not. Every discrimination instrument therefore gained a `--same_plate_only` mode and all discrimination results in this chapter were re-run within plate. Conclusions are unchanged; some magnitudes shrink.

The central manipulation is the *scramble ablation*: the prompt is left byte-identical except that the drug name and mechanism are replaced by another drug’s, while the con-

trol cell and the scoring truth are unchanged. Because both arms share the same control cell, control- and batch-derived leakage cancels in the scramble gap $\Delta_{\text{scr}} = \text{NIR}_{\text{model}} - \text{NIR}_{\text{scramble}}$. This is the only manipulation that isolates drug use, and every drug-use claim in this chapter is anchored to it rather than to a comparison against a drug-agnostic predictor.

Two constraints on the swap partner are required for a gap measured this way to mean what it says, and both exist because the unconstrained pool fails in a specific way. The partner must be a different drug: allowing the same drug at another dose or plate substitutes the drug for itself, an unchanged output is then correct rather than blind, and the contamination concentrates in the most similar stratum because a drug’s own other doses are usually its nearest neighbours. And the partner must be on the same plate: residuals retain plate structure, so a cross-plate partner differs from the target in batch as well as in drug, and the gap absorbs both. A scramble that swaps in a drug with a similar true response is also a weak test, and in this atlas that risk is structural rather than incidental. In full-profile space the mechanism annotation is close to uninformative about transcriptional response: in the characterisation behind Table 1, two drugs sharing a **moa-fine** annotation sit a mean distance 2.322 apart against 2.377 for two drugs that do not — 1,244 same-mechanism against 17,563 different-mechanism pairs, a ratio of 0.977. Same-mechanism drugs in Tahoe are just over two per cent closer to each other than arbitrary pairs are. That is a fact about the data rather than about the scramble, and it recurs: it is why conditioning on mechanism fails to move $\text{DE-}\Delta r$ in the grouping ladder below, and it is why the scramble’s own preference for a partner of different mechanism cannot on its own be relied on to deliver a different response — “different mechanism” does not imply “different response”. (Mechanism does carry signal in the residual frame, §3.13; the ladder below takes up why the two are not in conflict.) Two remedies are used. A swap-distance sweep — two hundred (drug, swap) pairs over twenty cell lines on tier 2, binned by quantile of the response distance $\|y_A - y_B\|$ between the real and swapped-in drug — turns the single test into a curve of Δ_{scr} against that distance. And a stratified scramble defines three partners per condition, chosen by residual cosine within the cell line: **near** (most similar), **orth** ($\cos \approx 0$) and **opposite** (most anti-correlated). The strata make visible an assumption that a single comparator hides: the gap has two parts — how much naming the truth helps, and how much naming a lie hurts — and only the first is wanted. A usable null is therefore a stratum whose own score sits at chance. §3.11 shows this is true of exactly one of the three, and not the one an evaluation would naturally default to.

The third instrument is forced-choice grading: a prediction is scored against its own condition’s truth and against another drug’s truth from the same comparison set, recording how often the own truth wins. Chance is 0.50 by construction, and the precision reference — the same forced choice with a disjoint sample of the condition’s own cells




Figure 4: **The scramble replaces the drug name and its mechanism and nothing else.** Both arms are conditioned on the same control cell and scored against the same truth, so any leakage through the control cell or the plate enters both arms in equal measure and cancels in the difference Δ_{scr} . This is why Δ_{scr} , rather than a comparison against a drug-agnostic predictor, is the quantity every drug-use claim in this thesis is anchored to: the drug-agnostic comparison is not leak-immune, and a drug-agnostic mean baseline scores $\text{NIR} = 0.399$ when the comparison set may span plates against 0.161 when it may not, precisely because it inherits that leakage. A refinement matters here and is developed in Section 3.3: if the swapped-in drug happens to have a *similar* true response then an unchanged output is correct rather than blind, so the swap partner is chosen from three defined similarity strata and the evidence is the *gradient* across them rather than any single value.

substituted for the prediction — lands between 0.67 and 0.83 depending on cells per condition. The fourth is output invariance, which asks whether the drug token changes the generated text at all, without reference to any truth. For a fixed control cell, the model generates twice from the same prompt and once from each of several prompts differing only in the drug, and the instrument compares

$$\text{invariance gap} = \underbrace{\overline{\text{sim}}(\text{same prompt, resampled})}_{\text{sampling noise alone}} - \underbrace{\overline{\text{sim}}(\text{prompts differing in the drug})}_{\text{noise plus any drug effect}}, \quad (11)$$

with similarity measured as top-100 Kendall τ and as Jaccard overlap. A positive value means two generations for the same drug resemble each other more than generations for different drugs do; zero means swapping the drug perturbs the output no more than resampling at the same temperature does. Reported over 8 contexts $\times$ 12 drugs at temperature 0.8. This instrument is the weakest of the four — it can only detect whether the output moves, not whether it moves correctly — but it is also the least able to be confounded, since it needs no ground truth at all.

Two of the four return intervals, and an interval that covers zero earns its keep by what it excludes, not by the word *null*. The output-invariance gap is 0.000 $[-0.019, +0.016]$ in top-100 Kendall τ ; two generations from a byte-identical prompt agree at $\tau = 0.198$, so the interval caps any drug-driven divergence at roughly a twelfth of the agreement two identical prompts retain, and the Jaccard version of the same comparison is tighter still ($+0.000 [-0.008, +0.008]$). On the identifiable stratum the scramble gap is $+0.014 [-0.016, +0.042]$, against the 0.170 that separates the model’s 0.768 from that stratum’s 0.938 ceiling: at the very top of the interval, naming the right drug rather than a wrong one would buy a quarter of what is there to be bought, and the point estimate puts it at 8% of that. Neither figure shows the effect is zero and no experiment here could; both put it far below the headroom these conditions leave. Read that way the result is stronger than the word null, not weaker.

Four objections to the verdict close on measurement rather than argument. First, *is there anything to know?* A ladder of mean-shift baselines — each predicting the average response of a progressively narrower group — answers that directly. On tier 2, predicting the global mean scores DE- Δr 0.056; conditioning on cell line raises it about two-and-a-half-fold to 0.142; conditioning on mechanism does not move it (0.052); and conditioning on mechanism and cell line is no better than cell line alone (0.120 against 0.142). Tiers 1 and 3 tell the same story. The informative conditioning variable in this data is therefore the cellular context, not the compound class. It is worth holding onto, because §3.6 finds precisely that asymmetry inside the model’s activations, and §3.13 finds that mechanism does carry signal once the generic programme is subtracted. The apparent tension is a frame difference rather than a contradiction: in full-profile space the generic response

dominates and swamps a mechanism effect that becomes visible only in the residual frame.

Second, *is the task winnable at all?* Not for most of the atlas, and an honest account needs two tests rather than one. Only 46.2% of (drug, cell line) conditions clear the identifiability bar of Table 1 — their own within-plate split-half NIR against their plate-mates reaching 0.8 — and 78.7% have a plate-mate closer to them than their own split-half replicate is. How inert the rest are depends on what is asked. A whole-profile permutation test against DMSO calls 72.7% of the 5,014 testable conditions active, leaving 27.3% statistically inert; a per-gene test is far harsher, and the median condition has no gene at all surviving BH correction at $q < 0.05$ (Table 1). The two do not contradict each other — a whole-profile test pools evidence that no single gene carries — but 27.3% quoted alone would make this atlas sound better resolved than it is, and any gene-level reading of these conditions answers to the harsher number.

Restricting to the identifiable subset does not rescue the model. It appears to beat the linear baseline there (0.768 against 0.531), yet a zero-drug-information control-copy reaches 0.766 on the same conditions — the paired model-minus-control contrast is +0.002 $[-0.026, +0.028]$ — and the leak-immune Δ_{scr} stays bounded as above. The model is drug-blind even where the task is provably winnable.

The control arm carries a second result that is easy to walk past, and it constrains how a ceiling of this kind should be read. Across the three strata the control-copy’s score tracks the ceiling almost step for step: 0.310 against a ceiling of 0.287 on unwinnable conditions, 0.540 against 0.689 on marginal ones, 0.766 against 0.938 on identifiable ones. A predictor that reproduces the untreated cell and knows nothing whatever about any drug does best exactly where the ranking problem is easiest. The reading this invites is that what makes a condition identifiable on this criterion is largely what makes its untreated cells distinctive among their plate-mates, rather than what makes its drug’s effect distinctive — so the distance between a predictor and a high ceiling is not all of it room that better drug modelling would fill. The ceiling itself rises with aggregation (0.61 at 10 cells, 0.76 at 40, 0.85 at 120), so the drug signal is real and noise-limited; what the strata add is that most of the ranking it is scored on can be done without it.

Third, *does the representation throw the drug away?* Cell sentences keep only the rank order of genes and discard expression magnitude, and if the drug lived in the magnitudes the diagnosis of the following sections would be attacking the wrong thing. The objection is testable without any model, by asking whether true expression discriminates drugs any better than rank does. It does not: at single-cell resolution, cosine, Pearson and Spearman similarities on true normalised expression all sit within 0.04 of chance (0.529–0.531 $[0.515, 0.543]$), the same near-chance regime as rank. What recovers the signal is aggregation, not representation: at fifteen cells the control-referenced shift cosine reaches 0.793 $[0.775, 0.809]$. The drug is not hiding in the magnitudes, which matters twice over. It rules out the simplest alternative explanation for everything that follows, and it means the repair




Figure 5: **The task is winnable where the noise ceiling is high, and precisely there, a predictor containing no drug information matches the model.** (a) The ceiling rises steadily with the number of cells per condition, cross-plate and within plate, so identifiability is noise-limited rather than absent: the drug signal is real and emerges under aggregation. (b) Conditions binned by their own ceiling, with five arms in every stratum. On the identifiable stratum ($n = 204$) the model reaches 0.768 — and a control-copy baseline that simply reproduces the untreated cell, and therefore knows nothing whatever about the drug, reaches 0.766. The paired contrast beneath each stratum is $+0.002$ $[-0.026, +0.028]$ there, clustered over cell lines. Showing the model alone across these strata would make the visual claim that it does better on winnable drugs; the control arm is what shows that it does not.

of §3.10 cannot be read as “restoring the magnitudes”: that repair keeps rank throughout and changes only what is being ranked. The same measurement bounds the task from below — at single-cell resolution the same-drug versus different-drug discrimination gap is only $+0.006$ $[+0.003, +0.008]$, so whatever a model can learn here, it cannot be learned from a single cell.

Fourth, *does the scramble swap in similar drugs?* The swap-distance sweep resolves this: Δ_{scr} is null in every distance bin including the far tail (-0.019), across a $2.3\times$ spread in swap distance, with a trend correlation of $+0.051$. Telling the model it is a drug whose true response is maximally different still does not change the output.

THE CLAIM. On the calibrated metric the model is drug-blind in the expression frame, and even where the task is provably winnable: at chance overall (0.498 against a 0.576 within-well split-half precision reference), matched on the identifiable stratum by a zero-drug-information control-copy (0.768 against 0.766 , paired contrast $+0.002$ $[-0.026, +0.028]$), and unmoved even when told the most dissimilar drug on offer.

3.4 Is the drug represented at all?

A null result invites two readings, and they call for opposite repairs. Either the drug never enters the model’s computation — an encoding failure, which target- and input-side fixes might reach — or it enters and is ignored, a readout failure, which only an objective-side fix could reach. Every behavioural measurement so far is consistent with both, because both predict the same output. The previous section has already closed a third possibility, that the task itself cannot be done: the ceiling rises with aggregation, nearly half of all conditions (46.2%) are identifiable in principle, and on those the model is matched exactly by a baseline that copies the untreated cell (0.768 against 0.766). So the failure is the model’s, and the useful question becomes where in the model it lives. That question cannot be answered behaviourally; it needs the weights. It decomposes into three sub-questions, each with its own instrument, because each instrument’s blind spot is exactly what the next must close. Is drug identity present in the activations? Does the network use the region where it is written? And when it uses that region, does it use which drug is written there, or only that something is? This section answers the first; the next two take the others in turn.

The first instrument is a linear decode probe, and its logic is deliberately weak in a way that has to be said out loud, because the obvious reading of a probe result is the wrong one here. The prompt names the drug in plain text (§3.1), so the drug is in the input by construction, and a reader that recovers it from the activations has established only that

the model did not discard what it was handed. A positive was close to guaranteed. The informative outcome would have been a negative — the drug field already gone — and what the probe buys, read as that kind of floor test, is the depth at which such a negative would appear. No claim is made yet about whether anything downstream reads what it finds.

The construction uses forward passes only, no generation and no training of the model itself. Real tier-2 evaluation prompts are grouped by drug; for each, the residual-stream activation is captured at the last prompt position, the position from which response generation begins, at every layer. A cross-validated multinomial logistic-regression classifier then learns to name which of the twelve drugs the prompt named, and its accuracy is compared against the floor measured on shuffled labels ($1/12 \approx 8\%$). The twelve drugs are a random draw from the tier-2 drugs with at least forty scored cells each; the count sets the 8% floor and, in §3.6, bounds the drug subspace at rank 11. One limitation is stated rather than repaired: the cross-validation folds are stratified over prompts and not grouped by well, so well-mates can straddle a fold boundary. The accuracy is therefore read as a presence test, not a generalisation estimate, and the argument below rests on the dissociation between decodability and the functional weight of the same region rather than on the absolute figure.

The negative never arrives. The probe reads 82% at layer 9 and 76% at layer 16, the deepest measured, against the $\approx 8\%$ floor; the between-drug to within-drug variance ratio collapses about three-fold across the final layers, so the drug is present throughout but geometrically de-emphasised as the output approaches. The survival is what carries weight here, not the peak: the identity is still linearly available at the deepest layer measured, at the position the first response token is emitted from, so nothing between the prompt and that layer has discarded it. That is the end of what a probe can say. It is an external reader with access to information the downstream layers may never consult, and whether the network itself reads this region is a causal question — the next section’s.

THE CLAIM. The drug is not missing from the computation: its identity is linearly readable from the residual stream at every layer measured, 82% at the peak and 76% at the deepest against an 8% chance floor. The prompt names the drug, so this is a floor test and its positive is close to free; what it rules out is that the drug field is lost before the readout runs — an encoding failure in the activations — and not that the readout ignores it. The drug is in the model; whether anything inside the model reads it is the next question. (§3.8 later uses *target encoding* for a different thing — how the target string is written — and nothing here bears on that.)

3.5 Is the representation used?

Asking whether the network uses the drug region means intervening on it, and intervening requires localising it first. The second instrument is built in four steps, and each step exists because the unguarded version of the same measurement misled us once already.

Step one: where the drug lives. On 480 tier-2 prompts — twelve drugs, forty real cells each — at seven layers, the twelve per-drug mean activation vectors are centred on the global mean and stacked, and their singular value decomposition contributes its top ten orthonormal directions: V_{drug} , the slab of the residual stream in which drug identity is written. Twelve centred class means span at most eleven directions, so the slab keeps essentially the whole span, noisiest directions included; for a removal claim that choice is conservative. The layer grid is even (2, 4, 6, 8, 12, 16) plus layer 9, added after the raw probe peaked there. A cell-line subspace is built identically and serves throughout as the positive control, because the model demonstrably does use cell line (§3.3’s grouping ladder).

Step two: purification. The raw drug slab is partly a batch slab. Drugs are confounded with plate and dose by the atlas’s design, and probing dose from inside the drug subspace recovers it at 0.90–0.95 at every layer: read naively, the “drug” axis would partly be measuring the batch. A context subspace is therefore built from the union of cell line, plate and dose directions, and the drug axis is orthogonalised against all of it, giving V_{pure} — the part of the drug axis not explained by batch. Every causal claim below is made on V_{pure} , never on the raw slab.

Step three: the removal gate. Before any “ablating the drug changes nothing” result can be believed, the instrument must prove that the slab actually contains the drug. The decode probe’s accuracy is measured before and after projecting the drug subspace out, and the fraction of above-chance signal removed must exceed 0.8 — otherwise the null is unfalsifiable, a measurement of nothing. The gate passes at every layer: accuracy falls from 0.41–0.82 to a constant 0.121 floor, removing 88–95% of the above-chance signal. The nulls below are real properties of the model, not broken hooks.

Step four: the causal measurement, in logit space. The true response is teacher-forced and its per-token log-probabilities are read off clean; a forward hook then projects the slab out of the residual stream at every position ($h \leftarrow h - (hV)V^\top$) and the log-probabilities are read again; the measure is the mean KL divergence from clean to ablated over the response tokens, on 60 prompts. Scoring in logit space rather than by generating is a deliberate repair of the project’s first causal attempt, which injected a steering vector and generated: the steering vector was small against the residual-stream norm, its effect on the output sat under the sampling-noise floor, and its “no effect” was uninterpretable. The ablation is run for V_{pure} (headline), for a random subspace of matched dimension (the null), and for the cell-line slab (the positive control). The prompt count is small and no

interval is quoted: the ablation is a sign-and-magnitude instrument, and what intervals the probe arc has come from the identity test of §3.6 — which is discounted there in turn, so no conclusion in this chapter rests on the probe arc alone.




Figure 6: **The model computes an increasingly explicit drug representation and barely uses it.** (a) Context-removed decodability by layer, with the raw probe drawn for contrast. (b) Ablation of the purified drug subspace against a matched-dimension random subspace, on a log axis because late-layer divergences are small for everything. This is a functional-variance comparison; whether drug *identity* is read requires the substitution test reported in the text.

What the ablation shows is a dissociation, and Table 6 quantifies it. With cell line, plate and dose projected out, drug decodability rises with depth, from 0.354 at layer 2 to 0.883 at layer 16: the representation becomes more explicit, not less. Yet the purified drug subspace’s variance share stays flat at ≈ 0.03 , against ≈ 0.6 for cell line, and ablating it costs $3.6\text{--}18.6\times$ the matched-dimension random null across layers. That effect is real. It is also small beside the cell-line positive control, which the same ablation reads at $1.0\text{--}11.2\times$ the drug’s: removing the drug subspace costs 9–37% of what removing cell line costs at layers 2–9, rising to 98% at layer 12 and 25% at layer 16. In variance-share terms the gap is wider still, cell line carrying 15–22 $\times$ the purified drug slab. The instrument is not the limit; the drug is a bit player.

This is also where ablation reaches its limit as an instrument. It proves the slab carries functional variance that downstream layers consume; it cannot prove that the slab’s drug-specific content is what is read: the purified directions could carry generic strong-programme energy that merely co-varies with drug across twelve compounds. Whether the network reads which drug the slab encodes requires replacing what the slab says rather than removing it, and that substitution test is the next section’s.

THE CLAIM. The drug representation is used, but at a bit player’s weight: decod-

ability rises monotonically with depth while the slab’s variance share stays flat at ≈ 0.03 , roughly 5% of the cell line’s ≈ 0.6 , and ablating it costs 9–37% of what ablating cell line costs at layers 2–9. Ablation cannot say whether the slab’s content is read — that needs a substitution.

3.6 Is the drug’s identity read?

The verdict first, because the instrument that produced it runs for pages and the verdict is small. In one of the five training arms of Table 20, at one depth of seven, injecting half the class-mean displacement that points from drug A to drug B inside the purified drug subspace ($\alpha = 0.5$, the rung the multiplicity correction below is applied to) moves the model’s preference towards drug B’s own held-out response — relative to an identically constructed injection that names a third drug — by $+0.0197$ $[+0.0063, +0.0335]$ of the preference gap the model already has. That is one to two per cent, against a 10% margin pre-registered for calling an effect material. Two further seeds reproduce it, and of twenty-six headline tests it is the only survivor of multiplicity correction. Read at face value that says the generation readout does consult which drug the region names, at a gain far too small to account for anything in §3.3. Read with the multiplicity discipline set out below it says less, because which result survives correction depends on which rung of the magnitude ladder the correction is applied to. This is the weakest load-bearing measurement in the chapter and it is carried as such: no argument downstream rests on it alone, and a reader who discards it entirely loses nothing that the drug-blindness result depends on.

The instrument is nonetheless given at length, because an effect of one to two per cent is worth reporting only if the design that produced it can be shown not to manufacture one — and three earlier versions of this design could not. The substitution test is the only instrument that can answer the question at all, and the design that survived audit is stated first; the three designs that failed are then the constraints it satisfies, one sentence each, with their algebra deferred to Appendix A.5. The test makes two structural choices so that its arms differ in nothing but which drug they name, and it ablates nothing. First, a contrast score,

$$\log P(g_B \mid \text{prompt}_A) - \log P(g_A \mid \text{prompt}_A), \quad (12)$$

where g_A and g_B are real response strings emitted under drugs A and B by held-out cells, and each log-probability is a per-token mean. Being a mean is what makes the contrast robust: any shift that is uniform across positions — an inflated logsumexp normaliser included — enters both terms equally and cancels exactly, whatever the two responses are and however many tokens each has. It does *not* cancel algebraically in the stronger sense one might assume. The normaliser is a function of the hidden state, and once g_A and g_B diverge the two teacher-forced sequences visit different states, so the state-dependent

part of the normaliser differs between them and survives the subtraction. That residual is not bounded here; what controls for it is the **permuted** arm, which shares the anchor, the construction and the norm and differs only in naming the wrong drug, so anything the contrast fails to cancel appears in that arm too. What survives the comparison of the two arms is the model’s preference between drug B’s own held-out response and drug A’s. Second, every arm is anchored at drug A as an on-manifold class-mean displacement of identical norm. Writing Π for the orthogonal projector onto the purified slab, the **swap** arm injects $\Pi(\mu_B - \mu_A)$, pointing at the drug whose response is scored; the **permuted** arm injects $\Pi(\mu_C - \mu_A)$, rescaled to identical norm — the headline control: same anchor, same construction, wrong drug; and an isotropic random direction in the slab is kept only as an off-manifold damage diagnostic, never as the comparison, because a random direction disrupts the model in ways an on-manifold class-mean displacement does not. The subtraction $\mu_B - \mu_A$ cancels the global mean exactly, so the injected vector is pure drug difference, and no ablation anywhere means there is no removed term for the control to trip over.

The three failed designs taught the constraints. A control drawn in the full 2048-dimensional stream puts $\approx 99.5\%$ of its energy where the intervention never acts, so the control must live in the same subspace as the treatment. Ablate-then-add delivers (added – removed), so an uncentred class mean turns the substitution arm into a restoration — hence no ablation, and an exactly centred displacement. A norm-matched random direction is off the class-mean manifold and inflates the logsumexp normaliser (measured +2.75 nats), depressing every target whether or not it carries identity — which is what the contrast score and the **permuted** arm above are for.

What the planted-world selftest does and does not cover follows from the same algebra. It plants a readout that genuinely reads the drug subspace and checks that the swap path recovers it, which validates the directional logic end to end. It cannot speak to the normaliser, because in the planted world each drug’s “response” is a single token: there is no prefix, so nothing diverges, and the one term the argument above had to be careful about is the one term the selftest never exercises.

Everything is out of sample, and everything is on a comparable scale. The subspaces and class means are fit on one slice of cells; g_B is drawn from a third disjoint slice — never from the rows defining μ_B , else the injected vector and the scored response would share cells and the test would be partly circular; all scoring happens on the held-out slice. Because the five training arms emit different target formats, raw log-probability differences live on different scales, so every effect is divided by that model’s own clean preference gap between A’s and B’s responses, measured with no intervention: the reported unit is the fraction of the existing preference gap the injection closes, comparable across arms.

Four guards carry the inference; the remainder — a drug-proxy screen, an injected-norm diagnostic and a three-world planted selftest — are engineering and are deferred

to Appendix A.5. Uncertainty is a cluster bootstrap over ordered (A,B) pairs, because rows sharing a pair share both the injected vector and the scored response. Equivalence is by TOST: the interval must lie inside a pre-registered margin and contain zero, with the margin set at 10% of the model’s own clean preference gap. A per-layer headroom gate excludes any layer where the cell-line positive control shows less than $5\times$ dynamic range, because a saturated instrument can support neither verdict. And a magnitude ladder repeats the test at $\alpha \in \{0.5, 1, 2, 5, 10\}$ times the natural displacement, because a real identity effect should show a monotone dose–response.

This design was run across the five training arms of Table 20 and, for the two residual arms, replicated at three further seeds. Of 26 headline tests — five arms by up to seven depths — exactly one survives Benjamini–Hochberg correction: the residual arm at layer 12 ($p = 0.001$, $q = 0.026$). The consensus arm at layer 2 ($p = 0.012$) does not survive ($q = 0.156$).

Which rung that correction is applied to matters, and the answer is not the one a reader would assume. The headline is the $\alpha = 0.5$ rung — half the natural displacement — and the ladder’s first entry is therefore the headline restated in raw nats (+0.0013) rather than an independent point beside it. At the full displacement $\alpha = 1$ the effect is *larger* and noisier: +0.0265 [+0.0021, +0.0524] against +0.0197 [+0.0063, +0.0335], at $p = 0.033$ against $p = 0.001$. Each of $\alpha = 0.5$, 1 and 2 excludes zero on its own, so the dose–response is real (in raw nats: +0.0013, +0.0017, +0.0044), and the two highest rungs do not ($\alpha = 5$ and 10 span zero, as an on-manifold displacement pushed far enough must). But applying the 26-test correction to the $\alpha = 1$ slice instead leaves *nothing* surviving, and applying it to $\alpha = 2$ promotes a different arm at a different depth (**residual_holdout** at layer 4, $q = 0.007$) while demoting this one. The slice decides not merely whether a survivor exists but which one it is, which is the signature of a family dominated by noise rather than one containing a single robust effect. The multiplicity-corrected claim is therefore a statement about the $\alpha = 0.5$ slice and is not robust to that choice — which is why this section opened by discounting its own result rather than closing on it.

The value of this is that it is a positive control on real data, and the argument for it borrows a measurement this chapter does not report until several sections later. The residual arm is independently known to respond to the drug its prompt names: on **train**, swapping the named drug and its mechanism for an orthogonal partner costs it +0.0898 in NIR, excluding zero clustered on cell line ([+0.0382, +0.1415]) and on drug ([+0.0291, +0.1506]). That is a behavioural result, measured without touching an activation, and it is established in §3.11; used here it is a forward reference, and a reader unwilling to take one on credit should read that section before accepting this control. Two things bound what it licenses. It is a training-split fact — the same establishment does not extend to held-out combinations — so the independent knowledge this control rests on is knowledge about that arm on conditions it was trained on. And it is behavioural, so it says the arm responds

to the drug its prompt names, not that the response is routed through the subspace this test injects into. What the pairing buys is narrower than that, and it is enough: the probe detects identity reading, replicably, in the arm behaviour calls drug-sensitive, and detects none at any depth it could measure in the arm that is behaviourally drug-blind. A null from an instrument that has never been shown to register a positive is uninterpretable; this one has.

THE CLAIM. The generation readout consults the drug representation at very low gain: $+0.0197$ [$+0.0063, +0.0335$] of the model’s own preference gap in the anchor seed and $+0.0133$ averaged over the three seeds that measured it — one to two per cent, a fifth or less of the 10% margin pre-registered for calling an effect material — in one training arm and at one rung of the displacement ladder. The multiplicity-corrected version of that statement holds for the $\alpha = 0.5$ rung and not for its neighbours: at $\alpha = 1$ nothing survives correction, and at $\alpha = 2$ a different arm does. The claim is not that the readout cannot read drug identity, and it is not the stronger mechanistic story that earlier drafts of this thesis told.

Three limits belong with it. The effect appears at one depth of seven, and layer 4 is positive in a single seed and null in the other three, which is what multiplicity discipline is for. The single-cell arm has no layer-12 measurement at all — its cell-line positive control reached $3.88\times$ there, under the $5\times$ headroom gate, so the layer was excluded — meaning the cleanest comparison, the same depth across two arms, does not exist; the dissociation rests on that arm being null at all six depths it could measure. And the sibling residual-holdout arm does not reproduce the effect at the headline rung ($+0.0028$ and $+0.0049$ in the two seeds where it was measured) — it is instead the arm that surfaces at $\alpha = 2$, which is the rung-dependence stated above rather than a second replication. Read together with the decodability series, the picture is a representation that becomes more explicit with depth while the process that writes the output barely consults it. That is suggestive of the workspace account of Section 2.7, and it is offered as an interpretation of the pattern rather than as a demonstration of the mechanism.

3.7 Denoising the target: three nulls at one scoring

That interpretation is not established, and this section does not proceed as though it were. What it does is test the cheapest thing the interpretation predicts, which is worth testing on its own account because it is the family of repairs a reader reaches for first. The prediction: if the defect is a readout that barely consults the direction of the drug region, whatever is written there, then improving what sits in that region cannot help — and improving the training target is exactly that kind of intervention. The inference runs one

way only. A positive here would tell against the low-gain reading; a negative is evidence about target-side fixes and not confirmation of the mechanism. The obvious objection is that one failed target construction proves nothing, so rather than try another we swept the family. Two constructions anchor the sweep. *Consensus* groups cells by (drug, cell line, dose) and averages them into one pseudobulk sentence, with every original prompt retained so only the response changes. *Optimal transport* computes, for each (cell line, plate) group, a coupling π between control and treated cells minimising $\sum_{ij} \pi_{ij} \|x_i - y_j\|^2$ under uniform marginals, solved by Sinkhorn iteration [29] on the kernel $\Gamma = \exp(-C/\varepsilon)$ with C the pairwise squared-distance matrix; couplings are computed in the released Tahoe scVI latent [31] ($n_{\text{latent}} = 10$), joined per cell by barcode rather than re-encoded — a re-encode produced a degraded latent (cell-line probe 0.11 versus 0.975 for the shipped latent) — and targets are built decoder-free, as convex combinations of real treated cells in expression space, so they always lie on the data manifold. The regularisation parameter then defines a ladder whose endpoints are experiments we had already run: $\varepsilon \rightarrow \infty$ gives a uniform coupling and hence the treated mean (consensus); $\varepsilon \rightarrow 0$ gives near-hard assignment, approximately a single state-matched treated cell. That is what makes this a sweep over a family rather than a third anecdote. It is a sweep at one generation budget and against one comparator, though, and the caption records the endpoint that a differently scored evaluation does not find null. Every rung was scored on the identifiable stratum — the conditions §3.3 showed to be winnable in principle — under the random-partner scramble of §3.3.

Training was not inert, and what it did instead is a result rather than an aside. The consensus model stops echoing the control: on the identifiable stratum it scores 0.653 against a zero-drug-information control-copy at 0.766, where the single-cell model sits on the control at 0.768. Its predictions are not mode-collapsed either — there is real between-condition variation. But that variation is uncorrelated with real drug identity: the Mantel correlation between the model’s between-drug distance structure and the true one averages +0.03 over cell-line and plate groups. The failure this literature names first for a perturbation model that underperforms on well-calibrated metrics is mode collapse — the prediction ignores the input and reverts to a mean, and the repair is to restore output diversity [21]. This arm is not that. It moved off the control, it kept the variation, and it spent it on the wrong axis. Producing no structure and producing structure that does not track the drug are different failures with different repairs, and diversity restoration addresses only the first.

THE CLAIM. No rung of the ladder separates from its scramble on the identifiable stratum under the random-partner test. Two things bound that negative: the consensus arm was stopped mid-epoch, so within the ladder it rests on the single-

cell and optimal-transport rungs, and the optimal-transport arm is not null in every frame — scored against the neutral comparator over all its conditions at the residual arm’s generation budget, its gap excludes zero (§3.10). By the inference rule stated above that positive tells against the low-gain reading rather than for it; it is measured in a different frame and so does not settle the matter, but it is one more reason the interpretation of §3.6 is carried here as an interpretation and not as a result. What the ladder establishes is that denoising the target changes what is in the target without changing what the readout does with it under this scoring — not that no target-side fix can work. The next two sections turn to a different property of the target — how much of it its tokens expose to the drug, and how much drug-specific signal it carries. Measuring those is not the same as showing that the target encoding is the defect, because the residual construction changes aggregation, control referencing, sign coding and target length alongside token divergence.

3.8 The residual frame: predicting only what makes this drug different

If the *target* encoding is the constraint — how the target string is written, a different sense of the word from the activation-side encoding ruled out above — the target must be redefined so that its tokens are the drug-specific part. How little of the standard target the drug occupies is measurable directly, and the measurement motivates everything that follows. Because a cell sentence ranks genes by expression, the target strings for two drugs in the same context are nearly identical, and the loss weights every token position equally. The measurement counts, for pairs of drugs sharing a context, how many of the 200 target positions carry a different gene in the two drugs’ target strings, averaged over pairs; the precision-reference column scores within-condition split-half drug retrieval through the identical representation, so the two columns are measured in the same units.

The natural inference is that only a small fraction of the gradient can be carrying drug identity. That inference is a hypothesis about the optimisation: what is measured here is the overlap of the target gene sets, not a token-level loss or gradient attribution, and it should be read at that strength. Its warrant is predictive — it explains why every change to target content failed identically in §3.7, and it predicts that changing what the tokens encode would not. That warrant does not identify tokenisation as the unique binding constraint. Gene positions are not loss-token positions under byte-pair encoding, and the residual target defined below alters aggregation, control referencing, generic subtraction, sign coding and sequence length at the same time as it alters token divergence. Separating those factors needs a matched-budget factorial that this thesis does not run.

The residual target is the repair. For each condition (drug, cell line, plate, dose) with

at least 40 treated and 20 plate-matched control cells:

$$r(d, c) = \underbrace{(\bar{y}_{\text{treated}} - \bar{y}_{\text{control}})}_{s(d, c)} - \underbrace{\frac{1}{|D_c|} \sum_{d' \in D_c} s(d', c)}_{\text{generic}(c)}. \quad (13)$$

The control subtraction removes the cell’s own state; the mean-over-drugs subtraction removes the generic drug programme shared across compounds. What remains is not what makes this drug different from all drugs: it is what makes it different *from the other training drugs that happen to share its plate and cell line*. The comparison set is fixed by the plating design rather than chosen, so the same compound plated among different neighbours would carry a different residual and a different target string. That is the price of estimating the generic locally.

The generic is a plate-local mean, so whatever a compound shares with its plate-mates leaves along with the stress and cell-cycle programme. Mechanism is the one channel where that matters: §3.13 finds mechanism-related drugs co-plated 0.404 of the time against 0.362 for a random draw, the only channel where the difference is material. For a co-plated mechanism class, then, part of the class’s own shared programme sits inside the generic and is subtracted out of every member’s residual by construction — so the mechanism channel is graded in a frame that has already removed some of what it is asked to find, and the bias runs against it.

The direction is testable rather than arguable. Joining Tahoe’s `moa-fine` labels to the plate rosters the gate itself scored, same-mechanism plate-mates supply a mean 4.4% of a condition’s generic, small per condition but present for 83% of the 917 labelled mechanism-arm conditions. The 155 conditions with *no* same-mechanism plate-mate — whose residuals therefore retain the mechanism programme — give a gap of +0.1145 [+0.0276, +0.2014], against +0.0737 [+0.0150, +0.1324] for the 762 that have one, both against the same plate-matched null under the same two-way clustering. That is the direction the construction predicts. It is not a corrected estimate: the two intervals overlap across almost their whole range, the split is post hoc, and the cleaner stratum rests on 43 wells. What it supports is a statement about sign rather than size — the frame understates this channel rather than flattering it, which makes §3.13’s inconclusive verdict on mechanism the conservative reading rather than a generous one.

The two subtractions are not equal partners, and the table above says which does the work. Removing the control moves the between-drug token difference from 34.6 to 111.4; removing the generic adds a further 6.8, so 92% of the recoverable difference is bought by the control subtraction alone. Almost everything the standard target hides is hidden by the cell’s own baseline. That is a property of the target encoding rather than a bad choice of target on our part: a list ranked by absolute expression puts the same highly expressed genes at the top whatever the compound, so what distinguishes two treatments

has to fit into the positions the baseline ordering leaves free. The measurement behind that reading is an overlap of gene sets, not a variance decomposition, and it is one atlas at $K = 200$ (6,602 conditions, `target_divergence.json`); how much of that ordering a perturbation displaces is an empirical quantity, so no margin is claimed for other atlases or other K . It is also the correct reading of §3.2: what saturates $DE-\Delta r$ there is reversion toward a central point, with `revert_center` scoring 0.9999 without fitting anything at all, rather than the generic drug programme specifically. The generic is nonetheless a real and separable component — it carries 39% of the shift’s energy (§3.9) — and subtracting it is what isolates the drug-specific remainder once the baseline is already gone. The inclusion thresholds are set where the identifiability curve of §3.3 flattens: the within-well split-half precision reference is 0.61 at 10 cells, 0.76 at 40 and 0.85 at 120, so below 40 treated cells a condition is still paying steeply for cells it does not have. Four properties of how the generic is estimated decide whether the remainder can be trusted, and each was forced by a measured failure.

First, the split precedes the fit. An earlier build estimated the generic and the reliability filter over every condition and assigned the holdout afterwards, so each held-out target was defined partly by other held-out conditions. The pipeline is therefore staged — inventory (metadata only), split (metadata only, before any pseudobulk exists), fit (the generic is handed the training keys explicitly), transform (held-out conditions are projected through the fitted generic and never enter it) — and a poison test stands behind the staging: corrupting every held-out expression vector must leave the training targets byte-identical, and corrupting a training condition must change them. What the split is, and why it is keyed at whole wells, is the subject of §3.11; here it matters only that it exists and is metadata-only.

Second, the scope is the plate, not the cell line. The original argument for cell-line scope was reproducibility: 61.7% of conditions reproduce across a half-split at cell-line scope against 19.3% at plate scope (`target_divergence.json`, cosine threshold 0.2, $n = 6602$), on the reasoning that a plate holds too few drugs to estimate a mean and subtracting a noisy one injects noise. That comparison does not hold up: both figures were computed with a shared control pseudobulk subtracted from both halves, which makes the halves correlate through their common control error, and the inflation is not neutral between scopes — it favours cell-line scope precisely because the cell-line generic is itself less noisy. That alone retracts the retention argument. A disjoint-control recomputation exists, but it was never written to an artifact in the canonical build and the two surviving records of it disagree on the plate figure, so no corrected percentages are quoted here; what both records agree on is that plate scope retains at least as many conditions as cell-line scope, so retention is not a reason to prefer cell-line scope. Two further diagnostics settle it in plate’s favour and are reported in §3.9: at cell-line scope the dose-transfer ordering inverts and the shared-minus-split control bias reads +0.242, against −0.000 at plate

scope. A plate group with fewer than three other training drugs yields no generic, and the condition is dropped and counted rather than silently falling back to the contaminated frame.

Third, the mean is over drugs, not conditions, and leaves the condition’s own drug out. A drug measured at five doses would otherwise contribute five times the weight of a single-dose drug. And once the generic is train-only, an asymmetry appears that had not existed before: a training condition’s generic would contain its own drug while a held-out condition’s would not, so the two splits would be scored against differently defined targets and the generalisation gap would partly measure that difference. The generic for every condition is therefore “the mean over training drugs other than this one, within scope” — identical in definition on both sides.

Fourth, the reliability line is resolved from the data, not inherited. A condition is kept for training only if its residual reproduces across a half-split, $\cos(r_A, r_B)$ above a threshold, with each half measured against its own control half so the two share nothing. The threshold itself was the single largest determinant of the training set, and it had been inherited (0.2) from an era when reliability was measured with shared controls — a different quantity on a different scale, under which that same 0.2 retained 62% of conditions at the cell-line scope then in use, against the 19.0% (948 of 4,999) it retains once each half is measured against its own control half. A moderate quality bar had silently become a stringent one. Rather than pick a new number, the builder reports the whole retention curve against a null that pairs half A of one condition with half B of a different condition — same encoding, same dimensionality, same noise, no shared biology — and the threshold is set at the null’s 95th percentile: a condition is kept when its two halves agree more than two unrelated conditions’ halves do. At the measured null (95th percentile +0.109, resolving to +0.1086) this retains 2,523 of 4,999 training conditions (50.5%). The inherited 0.2 survives as an argparse default in the downstream analysis configs — `truth_repro_thr` in `re_v3.json`, `repro_thr` in `vardecomp_matched.json` and `channel_gate_v3.json` — but it did not act as the training filter, and the canonical artifact shows it: of the 300 scored training conditions, 181 sit below 0.2, the minimum is +0.109, and none falls at or below the resolved +0.1086. The held-out sets are not filtered by this criterion: dropping held-out conditions for failing their own reproducibility bar would select the evaluation set on its own outcome, so noise in retained training conditions costs learning efficiency, not validity.

What this fraction means is stated carefully, because it is easy to over-read. The half-split agreement is a within-well sampling-precision criterion: it asks whether two disjoint cell subsets of the same well give the same residual. It is not biological reproducibility — 286 of 287 distinct (drug, dose) treatments in this cache occur in exactly one well, so no independent-well replicate exists here to measure that with — and the retained fraction is a statement about how much drug-specific signal survives pseudobulk noise at this cell

budget, not about how much biology is real.

The target string is `<up genes> [DOWN] <down genes> [END_CELL]`, 100 genes per block, ordered by signed residual magnitude; sign is biology, so it is given a symbol rather than being folded into a magnitude ranking. Predicted and true residuals are both mapped to a signed rank vector (top- k up positive, top- k down negative, weight $1/\log_2(\text{rank} + 1)$) so a generated sentence and a continuous truth are compared like for like; similarity is cosine. The score is NIR over other drugs in the same cell line — a comparison set that spans plates, in apparent violation of the same-plate rule of §3.3. It stands because the plate-scoped generic subtracts each plate’s mean response from every condition on it, so plate structure enters neither the truths nor, in expectation, the predictions; the rule’s purpose was to keep plate variance out of the comparison, and here the target construction has already removed it.




Figure 7: **The information was always present; the standard target’s tokens barely expose it.** Left: the drug-specific residual is built by two subtractions. The control removes the cell’s own state, leaving the *shift*; the mean over the other drugs in the same plate group removes the *generic* programme that every compound triggers, leaving what makes this drug different. Right: that same measurement, written two ways. Ranking genes by expression (a) puts housekeeping genes first, so only 34.6 of the top 200 tokens differ between two drugs in the same context; ranking by signed residual (b) raises that to 118.2. (Divergence counts quoted at cell-line scope, as in Table 9; plate scope is similar.) The bracket is the point of the figure — the within-well split-half retrieval reference is essentially identical under both encodings, so no information has been added. What changed is how much of it the token sequence exposes to the loss. That is a fact about the target’s tokens; it does not on its own identify tokenisation as the binding constraint (Table 9).

THE CLAIM. Ranking the target by signed residual exposes 59% of its top-200 gene positions to the drug against 17% for the standard target (Table 9; the residual row at cell-line scope), at an unchanged within-well split-half retrieval reference; and the residual is estimable without leakage — split before fit, generic plate-scoped and drug-measured, reliability resolved from the data at the matched null’s 95th

percentile (+0.109). What that establishes is that the standard target’s tokens hide a difference the data contain — and that this is a fact about ranking genes on absolute expression rather than a fact about these drugs, since 92% of the recoverable between-drug difference returns as soon as the cell’s own control is subtracted. It does not establish that tokenisation is the binding constraint: the residual package changes aggregation, control referencing, sign coding and length at the same time.

3.9 How much signal is there to encode

The residual defines what the model should predict. Before reporting what happened when a model was trained on it (§3.10), this section measures the target itself — how much drug-specific signal it carries, and of what kind — because the same training result reads differently against a large reachable signal and against a small one.

The shift’s energy decomposes exactly into residual 0.621, generic 0.390 and a cross term of -0.011 : 62.1% of the perturbation response’s energy is drug-specific, 39.0% is the generic programme, and the two components are essentially orthogonal, so the shares are genuine (`vardecomp_matched.json`). Put that number beside Table 9 and the encoding problem states itself in one line: the drug-specific part is the majority of the perturbation response, 62.1% of its energy, and two drugs in the same context still share 83% of the top-200 genes in the standard target. The drug is not a small effect in the response. The response is not what a cell sentence ranks. Every claim that follows is scoped by the first of those numbers, and it has not previously been computed for this task.

Within the drug-specific part, two components call for different responses. Modelling

$$r(d, c) = \beta(d) + \kappa(d, c) + \epsilon, \quad \mathcal{T} = \frac{\text{var } \beta}{\text{var } \beta + \text{var } \kappa}, \quad (14)$$

the per-drug constant $\beta(d)$, the drug *main effect*, is retrievable without any model: it is what a lookup table reproduces. The remainder $\kappa(d, c)$ is the only component a conditional model could win on. It is written here as a drug×cell-line interaction because that is what the decomposition assumes; the estimator actually used does not identify such a term, for the reasons set out below, so κ is called the *context-specific remainder* and $\mathcal{T}$ a disattenuated cross-context residual-similarity coefficient rather than a variance share. The decomposition’s additive part is what remains identified. This is DrEval’s naive predictor, in transcriptomic space: their `NaiveMeanEffectsPredictor` [19] predicts $\hat{y}_{ij} = \mu_i^c + \mu_j^d - \mu$, an additive model of cell-line and drug main effects with no interaction term, and our `generic(c) +  $\beta(d)$` is precisely that quantity lifted from a scalar potency to a 946-dimensional profile. It follows that κ is the term their naive model omits. It does not follow that $1 - \mathcal{T}$ measures the headroom available above it: the cross-line pairs

entering $\mathcal{T}$ may also differ in dose and in treatment well, and each condition’s residual is taken against a generic its own drug helped define, so $1 - \mathcal{T}$ bounds that headroom from above at best. DrEval finds that deep models barely beat this additive baseline on drug sensitivity; the transfer coefficient asks whether, in the transcriptomic setting, there is anything above it to win in the first place.

$\mathcal{T}$ is estimated by disattenuated cross-line correlation rather than by fitting the variance-components model directly, for a reason this cache’s layout forces: with 286 of 287 treatments in exactly one well there is no replicate from which to separate $\text{var } \kappa$ from $\text{var } \epsilon$, so the decomposition is unidentified on its own. The cosine route buys identification through split-half reliability:

$$\hat{\mathcal{T}} = \frac{\cos(r(d, c_1), r(d, c_2))}{\sqrt{\rho(d, c_1) \rho(d, c_2)}}, \quad \rho(d, c) = \frac{2 \cos(r_A, r_B)}{1 + \cos(r_A, r_B)}, \quad (15)$$

where r_A and r_B are disjoint halves of condition (d, c) and $\rho(d, c)$ is the Spearman–Brown reliability of its full-sample residual. Two cautions attach, and both are stated where the numbers live. Dividing out the reliability, the disattenuation, is essential: measurement noise alone depresses the raw cross-line cosine and would otherwise be read as interaction. And the correction is a Spearman–Brown formula applied to a cosine, though it is derived for correlations between parallel halves; the planted-truth validation of Appendix A.1 is what licenses it empirically. The reliability entering it is a within-well split-half precision, optimistic relative to true biological reproducibility. That optimism sits in the denominator: an over-large ρ makes $\hat{\mathcal{T}}$ too small, so the unshared fraction $1 - \hat{\mathcal{T}}$ is, if anything, *overstated*. An earlier draft of this section had the direction the wrong way round.

The measurement itself carries a control-design lesson, stated because it changed the protocol. The first run compared $\mathcal{T}$ against a negative control of different drugs on the same plate, which returned +0.485 against $\mathcal{T} = 0.511$, and was correctly declared void: that control differs from the estimand in two ways at once, different drug and same cell line, and plate structure surviving a cell-line-scoped generic inflates same-plate comparisons (the +0.485 just quoted, against –0.017 at plate scope) while leaving cross-line comparisons untouched. A structure-matched control — different drug, different cell line, breaking only the drug match — is the one that bounds the estimand, and a negative control that differs in more than one way cannot do so. Five checks in all accompany the estimate, and Table 10 states what each bounds.

$\mathcal{T}$ is a descriptive sensitivity analysis, not a load-bearing estimate, and the reason is in its own configuration. The scoped calculation retains conditions at a reproducibility threshold of 0.2; it permits cross-line pairs that differ in *dose* (`same_dose_only=false`); it uses a *self-including* generic (`loo_generic=false`), so each condition’s own drug contributes to the programme subtracted from it; and a cross-line pair may also sit in a

different treatment well. What the number is, precisely, is a disattenuated cross-context residual-similarity coefficient over pairs that differ in cell line, dose, well and target estimator at once: $\mathcal{T} = 0.553$ [0.514, 0.593], so 44.7% [40.7%, 48.6%] of that similarity scale is not shared across the compared contexts. It is not a share of variance and it is not a drug×cell-line interaction; reading it as either would require the other three axes to be inert, and nothing here establishes that. Its one useful asymmetry is directional: for the reason given above, an optimistic reliability in the denominator pushes $\hat{\mathcal{T}}$ down, so $1 - \mathcal{T}$ errs high and bounds from above what a per-drug constant leaves on the table rather than measuring it.

One contrast in this analysis comes close to varying a single factor, and it deserves stating as a result rather than spending as a check. Holding drug and cell line fixed and changing only the molar concentration, transfer is 0.707 [0.632, 0.782] on 363 pairs, against 0.553 [0.514, 0.593] for the pooled cross-context comparison; the two intervals do not overlap. Dose is the milder perturbation of a drug’s residual. The ordering is between one clean axis and a mixture rather than between dose and cell line alone — the cross-context arm differs in well and in dose as well as in cell line — but the dose arm is the closest this section comes to a single-axis contrast: drug and cell line are fixed, and only the well necessarily moves with the dose.

That ordering also did duty as a diagnostic when the generic’s scope was chosen, and it is the weaker of the two that were used. At cell-line scope it inverts: the dose arm falls to 0.485 while the cross-line coefficient holds at 0.511 (`vardecomp_baseline.json`), which is the wrong way round for what one expects of dose against cell line. But that expectation is a prior rather than anything measured here, so the argument is a plausibility check and is used as one. The second diagnostic needs no prior. At plate scope the shared-control and split-control builds of $\mathcal{T}$ agree to -0.000 ; at cell-line scope they differ by $+0.242$, which is plate structure leaking through a cell-line-scoped generic and inflating same-plate comparisons. That is what the scope decision rests on, and it is also why the cell-line-scoped build cannot be offered as corroboration when it returns a similar $\mathcal{T} = 0.511$: that build is the one declared void above, its own same-plate negative control reading $+0.485$.

The per-drug constant $\beta(d)$, which the lookup of §3.12 reproduces almost exactly, is therefore not the whole signal. What it cannot express is the part that does not transfer across contexts — a quantity that, for the reasons given above, does not isolate a drug×cell-line interaction, and no size for one is estimated here.

Whether anything could learn that interaction is a separate question. If a cell line modulated drug responses in a consistent direction, that direction would be estimable from the line’s own cells and a conditional model would have a concrete target. Writing $\kappa(d, c) = r(d, c) - \hat{\beta}(d)$, the test compares $\cos(\kappa(d, c), \kappa(d', c))$ for different drugs in the same cell line against the same quantity across different cell lines, with $\hat{\beta}$ estimated leave-

one-cell-line-out: estimated over all lines it would contain $r(d, c)$ itself, and subtracting a mean containing its own term induces a $-1/(m-1)$ correlation that biases the signal arm downward by construction, which would make a structured interaction look idiosyncratic — the exact error the test exists to avoid. Both arms are disattenuated by κ ’s own split-half reliability, since κ is a difference of two noisy quantities and is therefore noisier than the residual it comes from.

We looked for structure in κ and did not find it: no cell line modulates different drugs in a consistent direction, and a direct test of whether mechanism-matched drugs deviate from their per-drug constants in the same way within a given cell line returned nothing detectable. The temptation is to read two negative results as irreducibility, and they do not support it. The second test bounds its own power severely: κ ’s split-half reliability in this atlas is 0.195 at the 0.05 reliability floor (0.324 at the stricter 0.20 floor), so the quantity being correlated is barely estimable, and same-mechanism compounds are so heavily co-plated that only 3–7% of matched pairs survive a different-plate requirement (1% at the stricter floor), which leaves the plate confound in place. The defensible reading is that κ ’s structure is unresolved in this dataset. Resolving it would take more cells per condition, or a plating design that crosses mechanism with plate — neither of which is an analysis choice.

THE CLAIM. There is a conditional target that a per-drug constant cannot express, and no learnable structure in it is found. Its size is quoted as what it is: $\mathcal{T} = 0.553$ [0.514, 0.593] is a disattenuated cross-context similarity coefficient, so 44.7% of *that similarity scale* remains unshared across the compared contexts. It is not a variance share and not a drug $\times$ cell-line interaction — the estimator pools pairs differing in cell line, dose, well and target estimator — so it bounds what a per-drug constant leaves on the table without establishing that a model could reach any of it.

3.10 Re-encoding produces prompt sensitivity

Re-encoding the target as the drug-specific residual raises the differing fraction to 59% at an unchanged within-well split-half retrieval reference, and the model trained on it responds to the drug it is told about. On the 300 trained conditions of the residual-frame evaluation it beats the neutral scrambled comparator by +0.0898 [+0.0291, +0.1506], clustered on the drug and the treatment well, the drug being the axis a claim about drugs has to survive (`re_v3.json`; Table 13 sets this beside the held-out splits, and the two instruments that need no comparator arm at all are read below). That is the first non-null drug effect in this project — first in the order the experiments were run, and, as the budget-matched re-run below shows, not the only target construction that produces


one.

The gradient that first exposed the effect is worth stating in its own terms. On the earlier target build, the stratified scramble of §3.3 gave a gap that grew with swap dissimilarity — +0.0351 [+0.0061, +0.0656], +0.0716 [+0.0371, +0.1059] and +0.1429 [+0.1112, +0.1787] across the **near**, **orth** and **opposite** strata. The monotone ordering is itself the evidence: a metric artefact would not order itself by swap dissimilarity. Two qualifications ride with the magnitudes. They were read against the **opposite** comparator, which §3.11 shows is not a null, so they are quoted here as a gradient and the neutral-comparator magnitudes are adjudicated there. And they are the residual arm’s 1400-token run, while the two control arms plotted beside them were scored at the 600-token cap of §3.1, where their stratum gaps run from −0.021 to +0.026 (**fig-repair.json**). Those controls were read at the time as flat and null at every stratum. That is true of the single-cell arm, at 600 tokens and again at the residual arm’s budget. It is not true of the optimal-transport arm, whose 600-token gaps are themselves monotone in swap dissimilarity, and which does not stay flat once it is given the residual arm’s budget, as the re-run below establishes.

Three further diagnostics on that build agreed, and only the residual model passed all three. The correlation between condition reproducibility (**repro_cos**, the condition’s split-half residual cosine) and model NIR was +0.144 — the model does better where the drug effect is real, against −0.117 and −0.216 for the controls. Prediction diversity was highest for the residual model. And while $\cos(\text{prediction}, \text{own truth})$ exceeded $\cos(\text{prediction}, \text{other drugs})$ for all three arms, the residual model’s margin (+0.084) was roughly twice the optimal-transport arm’s (+0.049) and three times the single-cell arm’s (+0.026).

One measurement lesson and one withdrawal belong with this result. The effect was first measured at the 600-token generation cap of §3.1, which roughly halved it at every stratum where it was positive there: on **opposite** the gap reads +0.0716 at 600 tokens against +0.1429 at 1400, and on **orth** +0.0349 against +0.0716. The value +0.0716 therefore appears twice in this section — once as the 600-token **opposite** gap and once as the 1400-token **orth** gap quoted above — and the two are not the same quantity. Both controls have since been re-run at the residual arm’s 1400-token budget, and the result narrows this section’s claim rather than confirming it. The single-cell control stays flat, as it was at 600: −0.0128 [−0.0458, +0.0202] against its own neutral comparator, spanning zero. The optimal-transport control does *not*. At 600 tokens its gap was +0.0144 [−0.0051, +0.0338] and spanned zero; at 1400 it is +0.0292 [+0.0144, +0.0441] and does not. The token budget was suppressing a real effect in that arm, which is the same lesson the residual arm taught, applied to a control that was previously read as inert.

The comparison that decides the matter is between the two arms’ own gaps, and it costs this section its strongest phrasing. The residual arm’s pooled gap against the neutral comparator is +0.0361 [+0.0159, +0.0564] over 1,394 conditions; the optimal-transport arm’s,



residual frame. Intervals are 95% bootstrap, clustered over cell lines; $n = 250$ conditions per point.

Figure 8: On the earlier target build, the residual model’s advantage over its own scramble grew with how dissimilar the swapped-in drug’s response was; the single-cell control scored through the identical pipeline stayed flat and null, while the optimal-transport control did not — its 600-token gaps are themselves monotone, and at the residual arm’s budget against a neutral partner they clear zero (text). The horizontal axis is the measured mean residual cosine between the real and swapped-in drug, which is byte-identical across all four runs, so the positions are exact. The three solid series are at a matched 600-token generation budget; the residual model’s 1400-token run is drawn separately as the dotted overlay, since the budget alone roughly doubles the gap on the two strata where it is positive at 600 (+0.0349 $\rightarrow$ +0.0716 on *orth*, +0.0716 $\rightarrow$ +0.1429 on *opposite*). Intervals are 95% bootstrap clustered over cell lines, $n = 250$ conditions per point. The *near* stratum is the weakest test by construction: where the swapped-in drug’s true response is similar, an unchanged output is correct rather than blind.

at the same budget, is $+0.0292$ $[+0.0144, +0.0441]$ over 248 (`ctrl1400_summary.json`). Those intervals overlap almost entirely. Re-encoding the target therefore produces prompt sensitivity, but it is *not* the only target change that does, and the claim that it produced the first non-null effect in the project holds in the order the experiments were run, not as a statement that the residual encoding is uniquely responsible. Three limits bound the comparison in the other direction: the arms are scored in different frames, so their absolute NIRs are not comparable and only each arm’s gap against its own comparator is; they are separate evaluations over different condition sets, so this is an arm-level comparison and not a paired one; and those two condition sets differ in split composition by more than the arms differ from each other. The residual figure pools 300 trained, 895 unseen-combination and 199 unseen-drug conditions, whose own gaps are $+0.0898$, $+0.0332$ and -0.0315 ; the optimal-transport evaluation carries no split labels at all, so its 248 conditions cannot be sorted into trained and held-out. A spread of that size across splits, against a 0.0069 difference between the two pooled numbers, means the near-equality does not order the arms in either direction.

“OT achieved nothing” is therefore withdrawn, and not weakly. The optimal-transport arm’s 1400-token gap is measured against a partner whose true response is on average orthogonal to the target’s ($\cos = +0.035$), so it carries none of the comparator inflation §3.11 identifies. The 600-token evidence pointed the same way and was read too cautiously: the arm’s stratum gaps there were $+0.0073$, $+0.0144$ and $+0.0263$, monotone in swap dissimilarity like the residual arm’s, with only the **opposite** stratum excluding zero — and **opposite** is the one stratum this chapter shows is not a null. What the matched budget supplies is the neutral-partner reading, and it clears zero. The consequence reaches backwards: the optimal-transport target is a rung of the ε -ladder of §3.7, and that ladder’s negative holds only at the budget and comparator it was measured with. At a matched budget and a neutral partner the optimal-transport rung is not null; the single-cell rung is the one whose interval still spans zero.

The comparator-free evidence is already in hand: regressing each generation’s NIR on the cosine between the named drug’s truth and the target’s truth gives a slope of $+0.3728$ $[+0.2797, +0.4660]$ on trained conditions. Zero slope is the natural null, so no comparator defect can reach it — though the regression’s points are the three scrambled partner prompts and nothing else, and the target-drug prompt is never one of them, which bounds what it says (§3.11). Re-encoding the target made the model respond to the combined drug-and-mechanism prompt. What that response is worth is the question the rest of the chapter answers.

THE CLAIM. Re-encoding the target produced the project’s first non-null drug effect. On trained conditions the model beats the neutral comparator by $+0.0898$

[+0.0291, +0.1506] and beats chance outright — with no comparator arm at all — by +0.0981 [+0.0407, +0.1554], both clustered on the drug, the axis the claim rides on. The comparator-free slope of retrieval against the named drug’s truth is +0.3728 [+0.2797, +0.4660] — zero is its natural null, so no comparator defect can reach it, though it is fitted over the scrambled partner prompts alone, with the target-drug prompt never a point on the line. It is first in sequence and not uniquely responsible: at a matched 1400-token budget the optimal-transport target clears zero against its own neutral comparator too (+0.0292 [+0.0144, +0.0441] over 248 conditions that carry no split labels, against the residual arm’s +0.0361 [+0.0159, +0.0564] pooled over 1,394 that mix trained and held-out splits), on a separate evaluation in a different frame — so the two are not ordered against each other. Sensitivity is established; fidelity is the next question.

3.11 Sensitivity is not fidelity

A model that responds to the drug token has cleared a low bar, because reproducing a training average requires responding to it too. Everything scored above came from conditions the model trained on, so the result so far shows that the token is consulted — not that anything transferable was learned. Separating those two requires withholding data, and the physical layout of §3.1 forces which withholding is meaningful. A (drug, cell line) holdout — the field’s standard Leave-Pairs-Out split, and the split an earlier version of this work used — leaves the held-out condition’s treated well sitting in the training set through its other ≈ 48 cell lines. Measured against the holdout this investigation previously used, 124 of the 268 treated wells in this cache placed a held-out condition and a training condition in the same physical sample. What that estimand measures is within-atlas matrix completion, not independent generalisation, and no processing of the targets changes it, because it is a property of what was randomised.

Holding out whole wells is the only assignment-respecting choice the design offers, and it answers a different and narrower question: generalisation to an unseen treatment well of a seen drug — a new physical preparation, not a new cellular context. Fixed-dose cross-context transfer, present in one line and absent in another, is structurally unidentifiable in this atlas; this is a property of Tahoe’s design, not of our pipeline, and no rebuild fixes it.

The split used throughout the residual arm is therefore keyed at the well. `train` holds every condition of the remaining wells (4,999 conditions before the reliability line, 2,523 after it). `unseen_combo` holds 36 treated wells out whole (904 conditions, 900 after inventory filters), and the holdout never takes a drug’s last training well — without that guard a two-well drug could silently become an unseen drug while sitting in the unseen-combination arm, making the two arms measure the same thing; zero wells contribute

conditions to both sides of this split. `unseen_drug` holds every condition of 10 held-out drugs (714 conditions, 711 after inventory filters): Tahoe fine-tuning never sees the token, so this is the control — though the Pythia-1b backbone was pretrained on natural language and the prompt still supplies the `Mechanism:` field, so “unseen” means held out from fine-tuning with mechanism supplied, not unknown to the pipeline. A large effect there would indict the comparator rather than report a biological finding. We do not build a Leave-Cell-Lines-Out split here; Section 4 argues it is the natural next axis and explains why its interest is contingent on the transfer coefficient. Two warnings ride with the design. DrEval reports, citing [23], that apparently good performance on pair-holdout splits can be reached simply by predicting each drug’s average response across the training cell lines; our held-out-well estimand sits exactly in the regime where that warning bites, which is why the baseline ladder of §3.12 is not optional — the drug-average predictor is not a formality, it is the hypothesis under test. And evaluation samples conditions under an explicit per-split quota, because a uniform draw reproduces the pool’s proportions and starves precisely the split the experiment exists to test: in a first run, 250 constructed held-out combinations yielded only 33 scored conditions, purely as an artefact of the sampler. The residual-frame evaluation therefore scores 300 train, 895 `unseen_combo` and 199 `unseen_drug` conditions, 1,394 in all — the `unseen_combo` arm now takes almost the whole pool of 900, while `unseen_drug` remains a quota of the 711 available, so that verdict is powered on a subset of the conditions the split affords and its detection limit should be read with that in mind.

The first result from this design is about the comparator, and it is a finding rather than an erratum: it reframes every gap quoted so far. The comparator’s own score gives it away. Across the three swap strata the scrambled model’s NIR is 0.550, 0.501 and 0.423 as the named partner’s true response moves from aligned ($\cos = +0.22$) through orthogonal (0.00) to anti-aligned (-0.20) with the target (`re_v3.json`). A comparator whose own score is a monotone function of which drug it was handed is not a null: it is an active treatment that tracks the partner drug. The swap partner is itself a drug the model was mostly trained on, so telling the model a lie costs it precisely because it knows the partner. A gap measured against such an arm decomposes as

$$\Delta_{\text{scr}} = \underbrace{(\text{NIR}_{\text{model}} - 0.5)}_{\text{how much naming the truth helps}} + \underbrace{(0.5 - \text{NIR}_{\text{scramble}})}_{\text{how much naming the lie hurts}}, \quad (16)$$

and only the first term is wanted. The second term is not a rounding error. Pooled over all 1,394 scored conditions the model’s own advantage over chance is +0.0367 and the **opposite** comparator’s deficit is +0.0772: 68% of the +0.1139 gap that arm reports is the comparator failing, not the model knowing. Against `orth`, whose own score sits on chance (0.501), the second term is zero. Every gap in Table 13 is therefore measured

against the neutral **orth** partner, the only comparator that can carry a magnitude claim; the **opposite** figures are quoted in prose below, solely as the diagnostic that would have misled us.

The diagnostic is free and it is not specific to this experiment: score the comparator arm against chance, and if its score moves with which label it was handed, the gap it defines is not a measurement of the model. That condition is easy to meet by accident. A swap drawn to be maximally dissimilar — the obvious way to build a strong negative control — is exactly the swap that maximises the second term, so the harder the control is made to look, the more of the reported gap belongs to it rather than to the model.

The control behaves exactly as the diagnosis predicts. On **unseen_drug** — drugs held out of Tahoe fine-tuning, though not of the base model’s pretraining, and with the mechanism still supplied in the prompt — the gap against **opposite** is +0.0379 [−0.0083, +0.0841], positive where the neutral estimate is negative. The same conditions, the same generations and the same model scored against the neutral stratum give −0.0315 [−0.0836, +0.0206], an interval that covers zero as a control should — though the point estimate is negative and of the same size as the **unseen_combo** gap with the sign reversed, so this row is read as uninformative rather than as a confirmed null. The interval on the **opposite** arm is itself a lesson in the inference regime: a cluster-robust variance is asymptotic in the number of clusters rather than observations, and this arm spans 40 cell lines but only 28 treatment wells, so the two-way variance is carried on 27 degrees of freedom. A cell-line-only cluster bootstrap puts the gap clear of zero at [+0.0044, +0.0727]; adding wells as the second clustering axis, as every interval in Table 13 does, returns it to zero, and it does so at the normal quantile as well as at $t_{27} = 2.05$. What decides this row is the unit of clustering, not the quantile.

Read against the neutral partner, **train** clears zero under either clustering (+0.0898 [+0.0382, +0.1415] on cell line and well, [+0.0291, +0.1506] on drug and well) and **unseen_combo** clears it under only one of the two the claim needs: +0.0332 [+0.0069, +0.0594] clustered on cell line and well, and [−0.0031, +0.0695] clustered on drug and well. Drug is the axis a claim about transferring a *drug* signature has to generalise across, so generalisation to an unseen treatment well of a seen drug is not established. The precision matters: not established is not the same as shown to be zero. The drug-clustered interval reaches +0.0695, so an effect of up to roughly +0.07 remains compatible with these data at $n = 895$ spread over 34 drugs — a limit of power at the sample size the holdout affords, reported as one. (On **unseen_drug**, $n = 199$ leaves a half-width of 0.055 around the model’s own NIR and an 80%-power detection limit of 0.079 against chance.) The difficulty of the three splits is *not* comparable: their within-well split-half precision references are 0.956, 0.824 and 0.836, a spread of 0.132. The earlier conclusion that the splits are equally hard, and with it the inference that the held-out shortfall cannot be a noise-floor difference, is withdrawn — the two held-out arms sit against a markedly lower precision reference than **train**, and

this evaluation cannot separate that from a genuine loss of transfer.

The difference between the trained and held-out estimates is $+0.0567 [-0.0014, +0.1148]$ with cluster-robust $p = 0.0556$ (an unclustered permutation returns $p = 0.0142$, quoted only for contrast, because ignoring the clustering this chapter mandates everywhere else flatters the result). It does not clear the conventional threshold, and the two arms differ in three ways at once in any case: whether the condition was trained on, which drugs they contain, and how much headroom they have — their precision references are 0.956 and 0.824. Restricting the comparison to the 30 drugs present in both arms reverses the sign: the trained arm sits *behind* the held-out one, $-0.0606 [-0.1506, +0.0295]$, $p = 0.1797$, on 79 trained and 750 held-out conditions. The pooled difference is a composition effect: 83.8% of held-out-combination observations sit on those shared drugs against 26.3% of trained ones. The claim of a training-exposure or memorisation advantage is therefore withdrawn: none is established, the drug-matched estimate points the other way, and the unequal precision references confound what is left.

Two instruments that need no comparator at all cannot be touched by any of this, and they carry the positive content of the section. First, the slope test: regressing each generation’s NIR on the cosine between the named drug’s truth and the target’s truth gives $+0.373 [+0.280, +0.466]$ on `train`, $+0.336 [+0.223, +0.450]$ on `unseen_combo` and $+0.342 [+0.215, +0.469]$ on `unseen_drug` — all three exclude zero. What that regression is built from bounds what it can say. Its points are the three *scrambled* partner prompts and nothing else — exactly three per condition, so 900, 2,685 and 597 points on the three splits — and the real target-drug prompt is never a point on the line. On `unseen_drug` the swapped-in partners are overwhelmingly drugs the model *was* trained on. The slope therefore shows that the model responds to a named trained drug even in a scoring context whose own drug was held out; it is not evidence that the model knows the held-out compound. Second, scored against chance with no comparator, the model’s own NIR clears 0.5 on `train` (0.598 [0.549, 0.647]) and is not distinguishable from it on either held-out split (0.533 [0.492, 0.575] and 0.459 [0.404, 0.514]). A positive slope alongside a chance-level NIR is not a contradiction: it says the trained-drug-to-signature map exists but is coarse. The model has a coarse sense of what a named trained drug does and not of what it does in this well — and NIR asks the second question, ranking against the ≈ 40 other drugs measured in the same cell line. That is the drug main effect without the part of the residual that is *not shared across contexts*, which §3.9 measures at 44.7% ($1 - T$, with $T = 0.553 [0.514, 0.593]$) — a figure that mixes cell line, dose, well and estimator and is not a drug $\times$ cell-line interaction share. It is also why the lookup wins in the next section: `drug_lookup` is that main effect, estimated straight from data far more precisely than the model learned it.

The directional-agreement measurement says the same thing at condition level. The cosine between the model’s prediction and its own condition’s truth on trained condi-

tions averages +0.0389 (sd 0.1119, median +0.0220), with individual conditions spanning $[-0.3046, +0.3976]$ and 33.7% clearing $|0.1|$: weak on average and substantially heterogeneous, so a single summary number would misrepresent it in either direction.

One asymmetry attaches to every held-out number in this chapter, and it is stated here where those numbers live. The evaluation’s truths are not refitted: the generic is reconstructed from the build’s own fit inventory and verified against the build’s `fit_digest`, and a mismatch aborts the run rather than being absorbed, so the evaluation frame is the frame the checkpoint was trained against. What the reliability line does is select the *training* split: it is applied at build time to the conditions the generic is fitted on, so `train` is by construction the reliability-surviving population. The two held-out splits are *not* filtered: `unseen_combo` carries all 895 of its conditions, 499 of them below the +0.1086 reliability line and 74 with a negative split-half cosine, and `unseen_drug` all 199, with 88 below the line and 18 negative. This is deliberate — filtering the evaluation set on its own outcome is the selection this chapter removed — but it means the two sides are not matched on reliability, and the comparison is conservative in the direction that matters: the held-out arms are scored on their full, noisier populations while the trained arm is not.

THE CLAIM. Re-encoding the target made the model respond to the combined drug-and-mechanism prompt. That much is established and survives every stress: on trained conditions it holds clustered on the cell line and on the drug, it holds against chance with no comparator at all, and the comparator-free slope excludes zero on all three splits under both clusterings (+0.373 [+0.280, +0.466], +0.336 [+0.223, +0.450], +0.342 [+0.215, +0.469]).

Transfer to held-out combinations is not established. The gap there is +0.0332, which excludes zero clustered on cell line and well [+0.0069, +0.0594] and does not clustered on drug and well $[-0.0031, +0.0695]$, and it is the drug axis that the claim has to survive. The split is also weaker than its name: 30 of its 34 drugs and 39 of its 42 cell lines appear in training, and 82% of its conditions recombine a seen drug with a seen cell line. The effect is a majority tilt rather than a uniform one, positive for 20 of 34 drugs (sign test $p = 0.39$).

What the slope on the *unseen-drug* split does and does not add has to be stated exactly. At +0.342 [+0.215, +0.469] it shows that the output still tracks the drug *named in the prompt* when the condition’s own drug was held out. But the drugs named in that regression are the three scrambled partners, which on this split are overwhelmingly training drugs, and the real target-drug prompt contributes no point to it. The slope is evidence that the model responds to named trained drugs in an unseen-drug scoring context, and not evidence that it knows the held-out

compounds.

3.12 A lookup table wins

Sensitivity without fidelity is a real property, but it says nothing about quality: a model can respond to the drug in its prompt and still predict badly. Quality is what remains to be measured. The pre-registered bar was therefore never “beats chance” — residual targets are plausibly memorisable, and a predictor that simply looks up what this drug did elsewhere needs no model at all. That predictor is the comparison that matters, and it needs a fair construction. Each baseline in Table 14 is a predictor of the drug-specific residual, scored through an identical pipeline — same truths, same signed-rank encoding, same comparison set — and two disciplines make the ladder fair. The lookups are fitted from training conditions only: fitted from the full cache, a lookup scoring a held-out condition reads residuals the model was never shown (an oracle, not a baseline), so the oracle construction is excluded from the ladder rather than reported alongside it — every lookup row below is the training-only version. And `drug_lookup_1` selects a cell line and averages that line’s conditions rather than selecting a single condition, because 62% of drugs are multi-dose and a single condition would confound “different cell line” with “different dose”. `moa_lookup` is not an oracle in any split either: its partner pool is restricted to *training* conditions in the same cell line, so a held-out condition never enters the prediction made for it, and no unrestricted variant of it is reported here. The unrestricted-pool defect audited in §3.13 was in the channel gate, which generalises this construction; the restriction was already in force for the lookup arms here and the gate’s repair mirrored them, so the $n = 218$ row is on the restricted side of that repair.

The pre-registered interpretation, written before the result, anticipated three branches; the middle one fired. The model reads the drug and recovers only a small fraction of its signature; transfer to held-out combinations is not established (§3.11). Whether its prompt response is reducible to the **Mechanism:** field is not settled here: on the 218 conditions where that comparison exists, model – `moa_lookup` = +0.0436 [-0.0480, +0.1352], clustered on cell line and well, an interval that spans zero, so the two are not separated and the comparison is uninformative rather than favourable to either.

This is DrEval’s finding, in transcriptomic space, and it is the part of this chapter that travels furthest, because it needs no checkpoint, no prompt and no cell-sentence encoding — only a residual frame and a table of per-drug means. DrEval reports that deep learning models barely outperform a naive predictor of mean drug and cell-line effects [19], and our `drug_lookup` is the $\beta(d)$ term of that additive model, measured in a frame where the generic(*c*) term has already been subtracted. The finding replicates and sharpens: on a 946-gene transcriptomic response rather than a scalar potency, and against a within-well split-half precision reference, the additive drug main effect reaches 0.913 — above the

0.857 reference itself, on the common support — while a 1B-parameter model reaches 0.549. The warning DrEval inherits from [23], that good pair-holdout performance is reachable by predicting each drug’s training average, is treated here as the hypothesis rather than the caveat, and measured. It holds.

The first objection to this chapter is that its model is the wrong one: a conditional model built for perturbation response — CPA, scGen, GEARS and the chemical-perturbation line that follows them [9, 8, 10] — was never run on this split, so the head-to-head a reader wants is missing. It is missing, and that is a real gap. But the number that decides such a comparison is already in the table, and it is not the model’s. The ladder is architecture-free by construction — same truths, same signed-rank encoding, same comparison set — so a competitor can be dropped into it and scored, and what it then has to clear for a seen drug in a new context is `drug_lookup`: 0.913 pooled, or 0.890 on the held-out split where the well advantage is absent. Not 0.500. For a drug that never appeared in training no lookup exists at all, so that regime needs a different bar and gets one in §3.13; the two must not be allowed to share a number. What this section contributes is the bar, and the bar is reusable whether or not the checkpoint on the row beneath it is interesting.

On the common support the lookup scores *above* the reference, `drug_lookup`–reference = +0.056, and it is tempting to read that as “the modelling target has vanished”. That inference is not supported, for three reasons that all favour the lookup: the ceiling is scored against a half-sample truth while every baseline is scored against the full-sample truth; `drug_lookup` averages ≈ 38 conditions against the ceiling’s single noisy half; and NIR near 0.91 is compressive, collapsing most of the cosine range into a few thousandths. The two coverage figures are also distances to a split-half ceiling, which §3.1 notes is optimistic relative to biological reproducibility — both are inflated by an unknown amount, and the ordering is what carries the claim.

Those three reasons deflate the comparison with the *reference*. Only one of them — the averaging — could also deflate the comparison with the model, and the arm that tests it is `drug_lookup_1`, which draws on one other cell line and so carries no cross-line averaging advantage. It scores 0.822, and paired on the same common support the model trails it by -0.2723 $[-0.3074, -0.2372]$, clustered two-way on cell line and well. This drug’s mean response in a single other cell line beats a 1B-parameter model that was handed this condition’s own control cells. Averaging is why the lookup reaches the reference; it is not why it beats the model.

THE CLAIM. The model reads the drug and recovers a small fraction of its signature. The honest summary is the paired gap on common support ($n = 1192$), -0.3631 $[-0.3982, -0.3280]$: the model covers 13.9% of the distance from chance

to the within-well bar where a training-only lookup covers all of it and more. That the lookup exceeds the bar it is measured against is a property of the bar, not a headroom estimate — the bar is a single noisy half-sample while the lookup averages many cell lines, and NIR compresses near its top. The averaging is also not what beats the model: a lookup restricted to a *single* other cell line scores 0.822 and the model still trails it by -0.2723 $[-0.3074, -0.2372]$ on the same support. This is DrEval’s finding replicated in transcriptomic space: the additive drug main effect reaches 0.913 where a 1B-parameter model reaches 0.549. Being a bar rather than a result is what makes it reusable — it needs no checkpoint and no prompt, and any method proposed for seen-drug transfer on this atlas has to clear it rather than chance.

3.13 Channel availability for an unseen drug

An unseen drug can only be predicted through a property it shares with drugs seen in training, so the question is answerable by enumerating the channels and measuring each in the same frame. For a channel X , a drug’s residual is predicted as the mean residual of the other drugs in the same cell line that share property X with it — the construction of `moa_lookup`, generalised. It was intended to draw partners only from drugs that were in training, which would make it a fair contest in every split, unlike a same-drug lookup, which on an unseen drug is an oracle. An audit of the gate’s own implementation — not of the lookup arms of §3.12, which were already restricted — found that its partner pool was *not* restricted to training drugs but spanned every retained condition in the cell line. The restriction is now in the code and the gate has been re-run under it: of the 4,131 retained conditions only 2,523 are training, so nearly two in five of the conditions the old pool could draw on were ineligible. Every verdict in this subsection is from the corrected run, and the effect of the repair was not uniform — it moved the target gap down, the mechanism gap up, and chemistry’s down by half. The gate’s estimand requires the channel arm and its null to differ in exactly one respect, and two disciplines enforce it. Every channel is paired with a plate-matched random null: the same number of partner drugs, drawn at random from the same cell line, and the same split of partners across the target’s own plate and other plates, because mechanism- and target-related drugs are co-plated more often than chance — a null matched on count alone would differ from the channel arm in plate membership as well as in the channel, which is two differences where the estimand allows one. A co-plating diagnostic is printed before any verdict, so the size of the correction is a number rather than an argument, and where a plate-matched null spans too few cell lines the channel is reported UNTESTABLE rather than closed. “We could not build the control” must never be written as “closed”. Coverage is likewise reported before any verdict, because a channel annotated on few drugs cannot be closed

by a null result: it was never tested.

Before the clustering axis, a plainer question about population. The repair of the partner pool restricted where the *predictors* come from; it did not change which conditions are *scored*. The gate’s targets are all 4,131 retained conditions, spanning 94 drugs, of which 14 are held out of training entirely — and the target channel’s arm covers 18 drugs, of which **two** are held out. Restricting the targets to held-out drugs, which is the population the unseen-drug question is actually about, and clustered two-way on cell line and well as everything else in this subsection is, the target gap is +0.1197 [+0.0499, +0.1894] on those two drugs; mechanism falls from +0.0805 to +0.0439 [−0.0846, +0.1723] on four; and chemistry reads +0.0378 [−0.0029, +0.0785] on fourteen. Both of the latter span zero. Only the target channel survives the restriction at all, and it survives on two drugs.

The 18 are almonertinib, amsacrine, cepharanthine, clonidine, elagolix, gemfibrozil, goserelin, idarubicin, lumateperone, mitoxantrone, neratinib, norepinephrine, olanzapine, rapamycin, relugolix, temsirolimus, tofacitinib and tofacitinib citrate. The composition of that roster does as much work as its size. The target arm’s mean partner count is 1.19: 579 of its 712 predictions are made from a *single* partner drug and the remaining 133 from two, against 2.46 partners for mechanism and 5.00 for chemical similarity. And in some cases that single partner is an unusually close relative rather than an independent member of a target class — **Tofacitinib** and **Tofacitinib (citrate)**, the free base and the citrate salt of one molecule, are scored in exactly the same 20 cell lines and have exactly one partner each; so are rapamycin and its ester prodrug temsirolimus, in 22 lines. Neratinib and almonertinib pair the same way in 20 lines, and those two are distinct molecules, so the pattern is not only one of duplicate chemistry. Four in five of the target arm’s predictions therefore rest on one partner drug, and in a documented subset of them that partner is a salt or prodrug form of the same molecule. What the channel establishes is that a same-target neighbour predicts the drug. That a target *annotation* generalises across a class is the stronger statement, and testing it would mean removing the nearest analogue from the partner pool. That was not run.

What this section measures is therefore *channel availability* — whether the information a channel carries is present in the atlas at all — and not unseen-drug predictive performance, which two drugs cannot establish for the target channel and which mechanism does not survive at all. Every verdict below should be read in those terms.

One more thing decides how far the target verdict can be pushed, and it is the clustering axis. The gate clusters on cell line and treatment well, which is right for a statement about this experiment. But the claim the channel is being asked to support — that an *unseen drug* is reachable through its target annotation — generalises over drugs, so the drug is the unit, and only 18 annotated drugs carry the target arm. Recomputed on that axis the interval widens and the verdict changes: against the plate-matched null the target gap is +0.1351 [+0.0287, +0.2416] clustered on drug and well (df = 17), whose

lower bound falls 0.0013 short of the margin, and against the different-plate null it is +0.1322 [+0.0199, +0.2445]. Both still exclude zero; neither clears 0.03. Mechanism, which excludes zero on the cell-line axis, spans it on the drug axis in all three constructions. The honest summary is therefore narrower than *live*: protein target is the strongest channel of the three by a clear margin and carries signal on every construction and both clusterings, but its clearance of the relevance margin holds on the cell-line axis and not on the drug axis, and rests on 18 drugs. This is the same axis question §3.9 settles against a cross-line claim, and it is answered the same way: the axis is chosen by the claim, not by which interval is narrower.

The size of the co-plating correction is itself informative, so it is reported rather than argued away. Mechanism partners are co-plated with the target 0.404 of the time against 0.362 for a random draw (protein target 0.368 against 0.366; chemical similarity 0.370 against 0.367 — negligible for both, as expected for channels uncorrelated with plating; mechanism is the one channel where co-plating is material). Under a third construction that draws all partners from other plates, the target advantage is +0.1322 [+0.0648, +0.1996] ($n = 599$), mechanism’s is +0.0746 [+0.0243, +0.1249] ($n = 781$), and chemistry’s interval spans zero (+0.0118 [−0.0097, +0.0333], $n = 4090$). The channel arms’ absolute NIR in that construction sits at 0.638, 0.575 and 0.499 against nulls of 0.506, 0.500 and 0.487, so removing the shared plate does not push the arms below chance. Holding the plate apart barely moves target or mechanism — two and seven per cent off their plate-matched gaps — while chemistry’s roughly halves, which is where the co-plating correction actually bites. The bound this establishes is on similarity transfer through each channel, not on any model: a model could in principle combine channels or learn a non-smooth structure-to-response map that a nearest-neighbour average cannot express, though at a few hundred drugs the lookup is effectively the bound.

The ordering the three channels fall into is itself a result, and it is stable. Protein target sits above mechanism and mechanism above chemical structure under the plate-matched null, the count-matched null and the different-plate null alike, and again when the targets are restricted to held-out drugs. Clustered on cell line and well, the target interval and the chemistry interval do not overlap under any of the three nulls; each adjacent pair’s intervals do, and on the drug axis, under the plate-matched and different-plate nulls, even the two ends overlap. So this is an ordering of point estimates whose ends are separated and whose middle is not, and what it says about the atlas is that a compound’s transcriptional response tracks the protein it hits more closely than the chemical class it belongs to, with the mechanism label in between. One correction to it runs in a known direction: the residual frame subtracts a plate-local generic, and §3.8 shows that for a co-plated mechanism class part of that class’s own shared programme is subtracted along with it, so mechanism’s middle position is a floor rather than an estimate.

That ordering overturns a pre-registration made in this project. Two design documents

concluded that unseen-drug generalisation was out of reach, and both inferred it from the chemical-structure channel — Morgan fingerprints and a molecular language model — which the gate now places last of the three. The drug-knowledge specification written earlier in the same project had ranked raw structure the lowest-value channel and protein targets “the key feature”, on the reasoning that transcriptional response is driven by which protein a compound hits, something structure encodes only indirectly. The pre-registration generalised from the weakest channel to the strongest, and it was wrong: the gate reproduces the specification’s ordering, not the pre-registration’s conclusion. Tahoe’s own metadata carries protein-target annotations for 264 drugs (280 distinct symbols, 550 edges) which no analysis in this project had read.

Three caveats bound the target verdict. It is measured on the annotated subset: 264 of Tahoe’s drugs carry a protein-target annotation and 180 carry a mechanism label, and within this cache only 18 and 26 of the 94 scored drugs received the corresponding arm. What a shared property buys is a fraction of what the drug itself buys: on this metric the target arm sits +0.14 above chance where a same-drug lookup for a seen drug sits +0.41 (§3.12, a different condition set, so this is a comparison of scale and not a paired one). And target and MoA overlap, so whether target adds anything over MoA is a further question this measurement does not settle: settling it would need a leave-one-mechanism-out split rather than the leave-one-drug-out split used here, since a held-out drug’s mechanism class is generally still represented in training.

THE CLAIM. Protein target is the strongest of the three channels by a clear margin. It excludes zero on every construction under both clusterings, and on the population the unseen-drug question is actually about — targets restricted to drugs held out of training — the gap is +0.1197 [+0.0499, +0.1894], on two drugs. It clears the pre-set relevance margin of 0.03 clustered on cell line and well and does not clear it clustered on the drug, which is the axis a statement about unseen drugs rides on, and the arm rests on 18 annotated drugs of which two are held out of training. Mechanism excludes zero on the cell-line axis, falls 0.0021 short of the margin there, spans zero on the drug axis, and spans zero again when the targets are restricted to held-out drugs (+0.0439 [−0.0846, +0.1723]); it cannot carry an unseen-drug claim in any conditioning, and the frame understates it rather than flattering it. Chemistry’s interval is the narrowest of the three and the margin falls inside it. None of the three may be written as closed, and none may be written as live without naming the axis. Similarity transfer through the target channel is measurable and is the only channel worth pursuing; whether it is large enough to matter is unresolved at 18 drugs.

3.14 Which part of the prompt is read remains unresolved

§3.13 asked whether the information an unseen drug would need is present in the atlas at all. The complementary question is whether the model reads the drug information it is already handed, and that one is this section’s. The chapter has established that the output responds to that information as a block; it has not established *which* part of it is read, and what follows says why, and what would settle it.

The natural next step writes itself: append the protein targets to the prompt, retrain, and see whether unseen-drug performance rises. That arm assumes the bottleneck is information availability. Everything in this chapter points the other way, but the direct evidence is weaker than earlier drafts claimed, and it is worth saying exactly how weak it is. The scramble has until now replaced the drug name and the mechanism together, so a non-null gap establishes sensitivity to the combined prompt without attributing it to either span. A field decomposition costs one extra pair of generation arms and swaps exactly one span of the instruction line to the opposite-signature partner, leaving every other byte identical: the drug name with the mechanism kept, the mechanism with the name kept, and both together. One guard is essential rather than tidy: if the swap partner happens to share the original’s mechanism, then swapping only the mechanism produces a prompt identical to the model’s own, and the arm scores a gap of exactly zero by construction — read naively, that is indistinguishable from “the model ignores the mechanism”, the very claim under test, so such conditions are dropped rather than scored. This is why the mechanism arm carries fewer conditions than the others.

On the earlier target build, the decomposed arms suggested the response is carried almost entirely by the name span: swapping only the drug name gave Δ_{scr} of $+0.0809$ [$+0.0592, +0.1010$] ($n = 570$); swapping only the mechanism gave $+0.0091$ [$-0.0174, +0.0383$] ($n = 325$); and both together gave $+0.1024$ [$+0.0827, +0.1236$]. Those arms were scored against the **opposite**-signature comparator that §3.11 shows inflates every gap read against it. The arms *were* re-run under the corrected evaluation — Panel B of Table 16 carries the re-run figures — but the comparator defect survives the re-run, because it is a property of the swap partner and not of the target build. The claim that the model “reads the name and not the mechanism” is therefore withdrawn pending arms scored against the neutral **orth** partner; what the current run establishes is response to the combined drug-and-mechanism prompt, with the mechanism field’s contribution unresolved.

The practical consequence is narrower than the one earlier drafts drew, and more honest. The target-prompt arm is not retired by measurement — it is untested. What argues for building it carefully rather than eagerly is the shape of the failure documented above: the model’s deficit is not confined to missing information about unseen drugs, because on seen drugs it covers 13.9% of the distance from chance to the within-well

reference, against a training-only lookup that covers all of it and more (§3.12). And what argues for building it at all is the one thing a lookup structurally cannot do: handle a drug it has never seen — the regime where §3.13 finds measurable transfer through the protein-target channel, on the axis and the 18 drugs that verdict rests on. The model’s opportunity is not the part of the residual that fails to transfer across contexts: §3.9 puts $1 - \mathcal{T}$ at 44.7% [40.7%, 48.6%], but that is a shortfall in cross-context similarity, mixing cell line, dose, well and estimator — not a variance share and not an interaction — and whether any of it is learnable structure rather than noise is separated nowhere in this chapter. The opportunity is the unseen drug, and whether the model can be made to take it is not answered here.

3.15 The investigation in one table

The results are distributed across two measurement frames, and a reader assembling them has to keep track of which numbers may be compared with which. The two frames are not interchangeable and are separated deliberately. The expression frame scores a predicted cell profile against other drugs on the same plate; the residual frame scores a predicted drug-specific signature against other drugs in the same cell line, after the generic programme has been subtracted. Both have chance at 0.50 by construction, and that coincidence is the trap: the residual frame’s within-well precision reference is 0.854 where the expression frame’s is 0.576, so the 0.537 in one frame and the 0.498 in the other are not two attempts at one task, and the difference between them is neither an improvement nor a decline. Each has to be read against its own frame’s reference. Read that way, the expression-frame model does not clear chance at all, and the residual-frame model clears it by a margin §3.12 sizes at 13.9% of the distance from chance to the reference, on the support where the lookup comparison is made.

Five results in this chapter are not arms and cannot appear in the table. The metric audit: of the five candidate metrics graded, NIR is the only one whose DRF is positive — +0.635 [+0.604, +0.663] across plates, +0.545 [+0.516, +0.569] within plate, +0.704 [+0.667, +0.736] in the residual frame — and the other four are negative and exclude zero, which is to say they score an uninformative baseline above the noise ceiling (§3.2). Representation and readout: drug identity is decodable from the activations, and the generation readout consults it at one to two per cent of the preference gap, a fifth or less of the pre-registered material margin, in one training arm and at one rung of the displacement ladder — the chapter’s weakest load-bearing result (§3.6). Prompt sensitivity: re-encoding the target as the drug-specific residual produces a gap of +0.0898 [+0.0291, +0.1506] on trained conditions, clustered on the drug, and a budget-matched re-run of the optimal-transport control clears zero as well (+0.0292 [+0.0144, +0.0441]), so the re-encoding is first in the sequence rather than uniquely responsible (§3.10). Transfer:

$\mathcal{T} = 0.553$ [0.514, 0.593], a cross-context similarity coefficient and not a variance share (§3.9). And the channels: of the three routes to a drug never seen in training, only protein target clears the relevance margin, on the cell-line axis and not on the drug axis, on 18 annotated drugs of which two are held out (§3.13).

The distance between the model and its neutral scramble is the drug effect this work recovered; the distance between the model and the lookup is what it did not.

Task characterisation (streamed sample)		
conditions characterised		6,628
of which statistically active against DMSO	3,647 / 5,014 (72.7%)	
identifiable in principle (own within-plate split-half NIR $\geq$ 0.8)		3,064 (46.2%)
with a plate-mate closer than their own split-half		78.7%
median cells per condition		44
median #DEGs per condition (BH $q < 0.05$)		0
median SNR (effect $\div$ within-well split-half noise)		0.75
drugs with a real dose series	218 / 351 (62.1%)	
drugs seen in ≥ 3 cell lines		227
Residual-frame cache (rebuilt targets, §3.8)		
cells (treated / control)	620,564 (600,000 / 20,564)	
(cell line, plate) groups		198
drugs / cell lines		106 / 50
conditions with ≥ 40 treated and ≥ 20 control cells		6,617
by split after inventory filters (train / held-out wells / held-out drugs)	4,999 / 900 / 711	
retained by the data-resolved reliability line	2,523 of 4,999 train (50.5%)	
training examples / well-held-out validation examples		147,710 / 3,600
Evaluation sets		
tier 1 (seen conditions) / tier 2 (unseen drugs) / tier 3 (unseen combinations)	4,325 / 2,188 / 99	
conditions scored per tier in the generation evaluations		400
tier-2 drugs in the expression-frame NIR benchmark (within plate)		606
conditions scored in the residual-frame evaluation		1,394
by split (train / <code>unseen_combo</code> / <code>unseen_drug</code>)	300 / 895 / 199	
treated wells behind them (no well spans two splits)	136 / 36 / 28	
drugs per split (90 distinct in total)	76 / 34 / 10	
cell lines per split (42 distinct in total)	39 / 42 / 40	

Table 1: The populations behind every n quoted in this chapter. The 72.7% active figure is over the 5,014 conditions whose (cell line, plate) group carries plate-matched DMSO cells — the only ones a control-referenced test can reach — and not over all 6,628. The identifiability bar is a per-condition threshold — a condition’s own split-half NIR against its plate-mates must clear 0.8 — and is distinct from the aggregate ceilings of Table 2.

Ceiling	Frame and metric	Comparison set and budget	Where used
0.76	expression, DE- Δr	two real cells	metric audit (§3.2)
0.625 / 0.575	expression, NIR	cross-plate / within plate	plate leak (§3.3)
0.67–0.83	expression, forced choice	within plate, by cell budget	§3.3
0.576	expression, NIR	within plate, tier 2	the blindness verdict
0.61 / 0.76 / 0.85	expression, NIR	within plate; 10/40/120 cells	identifiability curve
0.956/0.824/0.836	residual, NIR	within cell line; train/combo/drug	§3.12
0.742 / 0.7432	token retrieval	within condition, split-half	§3.8

Table 2: One estimator, seven settings: every “ceiling” in this chapter is a within-well split-half precision reference — the score of a disjoint subset of the same treated well’s cells — and they differ in frame, comparison set, cell budget and the conditions the split contains. None is a biological replicate: this cache places 286 of 287 treatments in a single well, so for those there is no second well to replicate against. In the residual frame the reference is not constant across splits (0.956 / 0.824 / 0.836 on `train` / `unseen_combo` / `unseen_drug`, a spread of 0.132), so a distance to it is comparable within a split and not between splits.

Predictor	Information	DE- Δr (K50, tier 2)	partial-DE- Δr (tier 1)
<code>revert_center</code> (genes $\rightarrow G/2$)	none	0.9999	undefined
<code>revert_mean</code> (predict mean control)	none, no fit	0.961	0.25
ridge linear (control $\rightarrow$ shift)	control	0.947	0.28
noise ceiling (two real cells)	—	0.76	—
model (1B LLM)	control + drug	0.73	not measured

Table 3: DE- Δr is saturated by a control artefact. Regressing the control out leaves ≈ 0.26 of real but drug-agnostic skill in the two baselines with a defined partial value, and a baseline that performs no fitting at all reaches the same figure. The partial column comes from the tier-1 run that computed it; the rest of the table is tier 2.

Predictor (drug-agnostic)	cross-plate	within plate
mean baseline	0.399	0.161
control-copy	0.551	0.510
noise ceiling	0.625	0.575
fraction of drugs at chance	43%	49%

Table 4: Expression-frame NIR. A drug-agnostic mean baseline more than doubles its score when the plate is free to vary, which is what plate leakage looks like. The ceiling falls once the plate is held constant because the batch signature was inflating it too.

Instrument	Quantity	Value	Reading
Forced-choice grading	own-truth win rate	0.48 ($\approx$ scramble; ceiling 0.67–0.83)	chance
Output invariance	invariance gap	0.000 [−0.019, +0.016]	bounded, ≤ 0.016
Scramble, tier 1	DE- Δr	0.740 real vs 0.739 scrambled	indistinguishable
Scramble, tier 2	DE- Δr	0.733 real vs 0.729 scrambled	indistinguishable
Scramble (identifiable subset)	Δ_{scr} (NIR)	+0.014 [−0.016, +0.042]	bounded, $\leq +0.042$

Table 5: Four independent instruments agree, and the two that return intervals bound the drug’s effect rather than merely covering zero; both bounds are read out below. Near-ceiling DE- Δr is achieved identically with a wrong-mechanism drug, so that score is drug-independent.

Layer	Decode (raw)	Decode (context removed)	Ablate pure drug (KL)	Random matched dim	Ratio
2	0.408	0.354	0.155	0.012	12.7×
4	0.585	0.429	1.563	0.084	18.6×
6	0.692	0.529	1.435	0.251	5.7×
8	0.779	0.569	0.596	0.042	14.1×
9	0.821	0.602	0.418	0.100	4.2×
12	0.754	0.873	0.133	0.035	3.8×
16	0.758	0.883	0.015	0.004	3.6×

Table 6: All seven layers pass the removal gate (88–95% of above-chance drug signal removed); chance is 0.083. The raw probe peaks at layer 9 and declines; with context removed, decodability rises monotonically to layer 16.

Arm / seed	effect	95% CI	p
residual, seed 42	+0.0197	[+0.0063, +0.0335]	0.0010
residual, seed 43	+0.0118	[+0.0060, +0.0173]	0.0003
residual, seed 45	+0.0085	[+0.0029, +0.0143]	0.0015
single-cell arm, six depths	nothing detectable at any depth (median −0.0012)		

Table 7: The identity test, as a fraction of that model’s own clean preference gap between the two responses being compared (mean effect +0.0133 across the three seeds that measured it). Seed 44 produced no layer-12 measurement, it lost five of seven layers to the headroom gate, so the pre-registered criterion of two replications in three new seeds is met with the third a non-measurement rather than a failure.

Target	ε	Δ_{scr}	Verdict
single cell (real treated cell)	≈ 0	+0.014 [−0.016, +0.042]	null
OT barycentric (state-matched)	0.5	−0.001 [−0.018, +0.016]	null
consensus (global mean)	$\rightarrow \infty$	−0.010 [−0.022, +0.003]	null

Table 8: The regularisation parameter interpolates between the two target constructions tested independently, and no rung separates from its random-partner scramble on the identifiable stratum ($n = 204$). Two things bound how far that negative reaches. The consensus endpoint is the ladder’s weakest-evidenced point — that arm was stopped at roughly 60% of an epoch and its tier-1 NIR was unscorable — so within this table the negative is carried by the single-cell and optimal-transport endpoints, and the optimal-transport endpoint is not null in every frame. These rows score it on the identifiable stratum against a random partner at a 1200-token generation budget; a separate evaluation of the same checkpoint, scored against the neutral `orth` comparator over all its conditions rather than this stratum, moves from +0.0144 [−0.0051, +0.0338] at 600 generation tokens ($n = 250$) to +0.0292 [+0.0144, +0.0441] at the residual arm’s 1400 ($n = 248$), reported in §3.10. This table’s 1200 tokens are a third budget, lying between those two, and §3.1 shows the budget is a claim-level variable in this project rather than a runtime setting — so the two frames differ in comparator, population and budget at once. That rung has not been re-scored on this table’s terms, so the two frames are set side by side and not put into ratio. The table is a null under this scoring, not a null over the family of target constructions: outside this scoring the single-cell rung is the only one whose negative survives a matched budget and a neutral partner.

Representation	Top-200 tokens differing between drugs	Within-well split-half retrieval reference
full profile (the standard target)	34.6 / 200 (17%)	0.742
shift (treated − control)	111.4 (56%)	0.742
residual (drug-specific)	118.2 / 200 (59%)	0.7432

Table 9: 83% of the standard target’s tokens are identical between any two drugs in the same context. The residual row is quoted at cell-line scope; at plate scope it reads 120.3 and 0.7467. The third column is a split-half retrieval reference, not a biological replicate ceiling: it scores disjoint cell subsets of one drug in one cell line on one plate against each other. It is near-constant across the three rows; what the target encodings differ in is the first column — how much drug-discriminating content lands in the top-200 token positions. That is a fact about the target’s tokens; it does not on its own identify tokenisation as the binding constraint.

Check	What it bounds
disjoint control halves	shared-control inflation: +0.242 at cell-line scope, -0.000 at plate
structure-matched negative control	the estimand itself; returns -0.008
simulated $\kappa=0$ null	estimator bias: returns 1.001 where the truth is 1.000
dose strata on molar doses	a reference scale: the same drug and line at another dose
leave-one-drug-out generic	train/held-out target asymmetry (§3.8)

Table 10: The five checks accompanying the transfer-coefficient estimate. The simulated null matters most: at $\kappa = 0$ the *uncorrected* cross-line cosine is only 0.734, which read naively reports 27% interaction where the truth is zero, while the disattenuated estimator recovers the truth to within 0.001. That recovery is what licenses the reference column below. The dose strata carry a molar-resolved dose — an earlier cache carried the well identifier in the dose field, so what was labelled dose transfer compared physical samples rather than concentrations.

Quantity	Estimate	95% interval	Reference
$\mathcal{T}$ (cross-context, plate-scoped generic)	0.553	[0.514, 0.593]	simulated $\kappa=0$ null = 1.001
residual NOT shared across contexts, $1 - \mathcal{T}$	44.7%	[40.7%, 48.6%]	mixes cell line, dose, well and estimator
$\mathcal{T}$ across dose (same drug, same line)	0.707	[0.632, 0.782]	reference scale
matched negative controls	-0.017, -0.008	spanning zero	must be ≈ 0

Table 11: The transfer coefficient, estimated on real molar doses with dyadic cluster-robust intervals. The matched negative control sits on zero; the same-plate control that produced the earlier false alarm reads +0.485 under a cell-line-scoped generic, for the reasons given in the text. $\mathcal{T}$ is a cross-context similarity coefficient, not a variance share, and $1 - \mathcal{T}$ must not be read as a drug $\times$ cell-line interaction share.

Configuration	same cell line	different cell lines	excess within-line
plate-scoped generic	-0.0149 [-0.0251, -0.0050]	+0.0003 [-0.0071, +0.0081]	-0.0153

Table 12: No cell line bends different drugs in a consistent direction. Both entries are disattenuated by κ 's own split-half reliability, and their intervals come from a cell-line-clustered bootstrap (38 and 42 clusters). 4,000 pairs in the same-line arm and 3,884 across lines, over 1,282 conditions (`vardecomp_matched.json`). The same-line interval lies entirely below zero, so the arm that would have to be positive for a per-cell-line direction to exist is not merely null. The excess is the difference of the two point estimates; no interval was computed for it. A cell-line-scoped counterpart was not computed in the canonical run and the row an earlier draft carried for it is withdrawn.

Split	n (lines; drugs; wells)	model NIR	Δ_{scr} , clustered on <i>cell line</i>	the same, clustered on <i>drug</i>
<code>train</code>	300 (39; 76; 136)	0.598 [0.549, 0.647]	+0.0898 [+0.0382, +0.1415]	+0.0898 [+0.0291, +0.1506]
<code>unseen_combo</code>	895 (42; 34; 36)	0.533 [0.492, 0.575]	+0.0332 [+0.0069, +0.0594]	+0.0332 [−0.0031, +0.0695]
<code>unseen_drug</code>	199 (40; 10; 28)	0.459 [0.404, 0.514]	−0.0315 [−0.0836, +0.0206]	−0.0315 [−0.1049, +0.0419]

Table 13: Whole-well holdout on the corrected evaluation — the truth verified against the build’s own fit digest, and the full held-out population rather than the reliability-surviving subset. The two right-hand columns carry the SAME point estimate and differ only in what they treat as the unit of generalisation. On `train` the gap excludes zero either way. On `unseen_combo` it excludes zero clustered on cell line and does not clustered on drug — which, for the reason the text below gives, is why this row is reported as not established. Every gap printed here is measured against the neutral `orth` partner, the only comparator that can carry a magnitude claim; the `opposite` figures are quoted in prose only, as the diagnostic, never as a column. The `unseen_drug` row corroborates rather than proves: those drugs were held out of Tahoe fine-tuning, so a large gap there would indict the comparator — but the base model was pretrained on natural language and the prompt still supplies a mechanism, so a strictly zero effect was never guaranteed. Its neutral gap is negative (−0.0315) while its `opposite` gap is +0.0379: the same comparator asymmetry the other two rows show. What establishes comparator non-neutrality is the direct 0.550/0.501/0.423 ordering above. Model NIR intervals test against the chance value 0.5 with no comparator. Intervals are two-way cluster-robust over cell lines and wells, with a Student t reference on $\min(G_{\text{line}}, G_{\text{well}}) - 1$ degrees of freedom — which is why the well counts are printed: the `unseen_drug` arm spans 28 wells, and $t_{27} = 2.05$ against 1.96 widens its interval by five per cent (§3.1).

Arm	What it knows	NIR	Note
within-well split-half reference	disjoint cell halves, one well	0.854	precision, not a replicate
drug_lookup (training-only)	drug’s mean residual, other lines	0.913	the bar to beat
drug_lookup_1	the same, from one other cell line	0.822	still beats the model
moa_lookup ($n = 218$)	drugs sharing this drug’s MoA	0.552	narrower support; not separated
model	control cell + drug	0.537	
control_copy	nothing (copies the control)	0.506	drug-agnostic
generic	nothing (predicts the average)	0.500	the zero vector
random	nothing	0.501	floor

Table 14: The baseline ladder on the rebuilt targets, $n = 1,394$ scored conditions; the lookups are fitted from training conditions only. Reference and model rows by split: 0.956/0.824/0.836 and 0.598/0.533/0.459 (train / unseen-combination / unseen-drug) — the three precision references differ by 0.132, so the splits are not equally hard and their model rows are not directly comparable. **drug_lookup** is the pre-registered win condition: a lookup can reproduce a memorised average, so only a model that tailors the response to the control cell’s state can beat it. **moa_lookup** is scored on the 218 conditions that have a training-drug MoA partner, so that row does not sit on the $n = 1,394$ support the rest of the ladder does; paired on that support model – **moa_lookup** = +0.0436 [−0.0480, +0.1352] two-way clustered, an interval that spans zero, so the two are not separated in either direction; and the two pooled means are averages over different condition sets, so their ordering in the table is not a comparison. The paired comparison on the common support ($n = 1,192$) is model – **drug_lookup** = −0.3631 [−0.3982, −0.3280] two-way clustered, and of the distance from chance to the within-well split-half precision reference the model covers 13.9% against the lookup’s 115.7%. That pooled figure should not be the one quoted, because the lookup is not equally clean on both sides of the split. It averages the same drug from *other cell lines*, and a Tahoe well holds roughly 49 lines, so for a trained target those other lines can sit in the target’s own physical well and carry its batch. On held-out wells they cannot. The lookup scores 0.980 on **train** and 0.890 on **unseen_combo**, and that 0.09 is the size of the well advantage. The clean comparison is the held-out one, where the target’s well is absent from training by construction: there the model trails the lookup by −0.357 [−0.403, −0.311] over 893 conditions, against −0.381 [−0.430, −0.332] over 299 trained ones. The ordering is the same and the conclusion does not turn on the choice, but the held-out number is the one that means what the sentence says.

Channel	Drugs (annot./scored)	n	vs plate-matched null	(vs count-matched)	Verdict
protein target	264 / 18	712	+0.1351 [+0.0745, +0.1958]	+0.1498	live, cell-line
mechanism (MoA)	180 / 26	921	+0.0805 [+0.0279, +0.1332]	+0.0623	inconclusive
chemical structure	377 / 94	4131	+0.0261 [+0.0035, +0.0487]	+0.0122	inconclusive

Table 15: Coverage is stated first, as the number of Tahoe-annotated drugs behind each channel, for the reason given above: a null result on a channel that was never really tested is not a closure. The plate-matched null is the estimand-correct control; the count-matched column shows how much of each raw gap was co-plating. The verdicts are the gate’s own, and its rule is that the plate-matched interval’s lower bound must clear a pre-set margin of 0.03 when clustered on cell line and well: on that axis, and only on that axis, protein target clears it (+0.0745) while mechanism (+0.0279) and chemistry (+0.0035) do not. All three intervals exclude zero there, so the two that fail the margin are inconclusive rather than equivalent to their nulls, and neither may be written as closed. The two fail it for different reasons, and the distinction matters. Mechanism’s lower bound falls 0.0021 short — a knife-edge that should not be read as a negative result in either direction. Chemistry’s interval is the *narrowest* of the three (0.045 wide against 0.121 for target), so it is not underpowered: the effect is small and well estimated, and the margin simply falls inside it. [†]The axis is written into the target verdict because the word “live” is not usable without it: clustered on the drug, which is what a claim about unseen drugs has to generalise across, the target interval is +0.1351 [+0.0287, +0.2416] and falls 0.0013 short of the same margin (§3.13). Two further caveats bound every row. The counts are annotations in Tahoe’s drug metadata rather than testable coverage here: within this cache only 18, 26 and 94 of the 94 scored drugs received a channel arm at all. And the three arms average very different numbers of partners — 1.19 drugs for target, 2.46 for mechanism, 5.00 for chemistry — so the channel with the largest gap is also the one predicting from the fewest neighbours.

Panel A: the channel is available — metadata-neighbour retrieval, plate-matched				
Channel	Drugs (annot./scored)	n	Gap against a plate-matched random null	
protein target	264 / 18	712	+0.1351	[+0.0745, +0.1958]
mechanism (MoA)	180 / 26	921	+0.0805	[+0.0279, +0.1332]
chemical structure	377 / 94	4131	+0.0261	[+0.0035, +0.0487]
the same three, against the stricter DIFFERENT-PLATE null				
protein target		599	+0.1322	[+0.0648, +0.1996]
mechanism (MoA)		781	+0.0746	[+0.0243, +0.1249]
chemical structure		4090	+0.0118	[−0.0097, +0.0333]
Panel B: is it read? — one span swapped, every other byte identical				
Span swapped		Δ_{scr}	n	Comparator direction (partner NIR)
drug name, mechanism kept	+0.1026	[+0.0625, +0.1427]	1394	inflated (0.434)
mechanism, drug name kept	+0.0303	[−0.0078, +0.0685]	738	suppressed (0.520)
both spans (the standard arm)	+0.1139	[+0.0781, +0.1496]	1394	inflated (0.423)

Table 16: Availability and use are different questions, and only the first is currently answered. Panel A shows that a drug’s protein targets and mechanism carry response information recoverable by retrieval alone, so an unseen drug is not information-starved. Its first column gives drugs annotated in Tahoe’s metadata over drugs actually carrying that arm in this cache: chemistry is the only channel defined for all 94 scored drugs, and the target arm rests on 18. Read the two halves together. Under a plate-matched null all three channels clear zero; under the stricter different-plate null, which removes the last shared batch structure, target and mechanism continue to clear it and *chemical structure does not* — +0.0118 [−0.0097, +0.0333]. Much of chemistry’s apparent signal is co-plating. Clearing zero is not the same as clearing the gate’s relevance margin of 0.03, and only target does both: mechanism’s different-plate lower bound of +0.0243 leaves it inconclusive rather than live. Intervals here are two-way cluster-robust over cell line and treatment well; under the one-way cell-line interval this table previously carried, every row looked significant. Panel B asks whether the model uses the field it is already handed. Swapping the drug name moves the output by +0.1026; swapping the mechanism, with the name left in place, moves it by +0.0303 with an interval containing zero. The attribution that reading invites — that the name carries the effect and the mechanism field does not — is *not* supported, because the comparator is not neutral and is non-neutral in opposite directions across the two rows. The name arm’s scrambled partner scores 0.434, below chance, so about two thirds of its +0.1026 is the comparator falling rather than the model rising; the mechanism arm’s partner scores 0.520, above chance, which suppresses its gap. Measured instead as the model’s own distance from chance on each row — the decomposition of §3.11 — the ordering inverts: +0.0367 on the name row against +0.0506 on the mechanism subset. The panel therefore establishes that the model responds to the combined prompt, and leaves the split between the two spans unresolved until the arms are re-run against the neutral comparator. Both intervals are clustered on the drug and the treatment well. The mechanism row carries fewer conditions by construction: where a swap partner happens to share the target’s mechanism, swapping only the mechanism reproduces the model’s own prompt byte for byte and the gap is zero by definition, so those conditions are dropped rather than scored.

Frame	Arm	NIR	What it knows
expression (within plate, tier 2, $n = 606$)	within-well reference	0.576	two disjoint halves of one treated well; precision, not a replicate
	control-copy	0.504	nothing — it reproduces the untreated cell
	linear map	0.500	the control cell, no drug
	model	0.498	control cell + drug
	global mean	0.180	nothing
	model , identifiable stratum	0.768	the same model, on the 204 conditions whose own reference clears 0.80
	control-copy, same stratum	0.766	nothing — and it matches the model there
residual (cell-line sets, $n = 1394$)	within-well reference	0.854	two disjoint halves of the same treated well; not a biological replicate, and the lookup exceeds it
	drug lookup (training-only)	0.913	this drug’s signature from other cell lines
	drug lookup, one line	0.822	the same, from a single other cell line
	model	0.537	control cell + drug
	scramble, orth partner	0.501	the neutral comparator: the partner is chosen by measured residual cosine to the target ($\cos \approx 0$) before any scoring, and its own score is then <i>empirically</i> at chance
	scramble, opposite partner	0.423	an active treatment, not a null
	generic response	0.500	nothing (the zero vector, by construction)

Table 17: Every arm on the calibrated metric, by frame. Chance is 0.50 in both. The reference row in each frame is a within-well split-half figure and not a biological replicate: this cache places 286 of 287 treatments in exactly one well, so no second well exists to replicate against, which is also why a training-only lookup can score *above* it. The residual-frame figures are the corrected evaluation of §3.11 and §3.12 — 1,394 conditions, truth verified against the build’s own fit digest, and the full held-out population rather than the reliability-surviving subset. The last two expression-frame rows are the identifiable stratum of §3.3, the 204 of the 606 conditions whose own reference clears 0.80: where the task is provably winnable is also where a predictor holding no drug information matches the model. The **orth** comparator is one of three swap partners fixed before any scoring by residual cosine to the target (+0.22, 0.00, −0.20); that the neutral one also scores at chance is what its construction predicted and the audit then confirmed, and it is why only that comparator can carry a magnitude claim (§3.11).

4 Limitations and Future Research Directions

4.1 Limitations of scope

Drug coverage. The residual analysis is built on a cache of 620,564 cells covering 106 drugs across 50 cell lines, sampled from Tahoe’s shards — roughly a tenth of the atlas’s compound set. This is ample to test whether the model uses the drug, and to measure transfer, but it is not a population estimate. Any claim about the *distribution* of drug behaviours should be read against the broader characterisation, which used 6,628 conditions.

Only reproducible conditions. The residual arm trains on the 50.5% of conditions whose residual reproduces across a half-split — 2,523 of 4,999 training conditions, kept at $\cos > 0.109$, the 95th percentile of a null that pairs half A of one condition with half B of another. This is the right choice — training on irreproducible residuals is training on noise — but it means the model was taught the winnable fraction of the task, not all of it. A reliability filter might be expected to select *toward* context-independent conditions and so bias the transfer coefficient upward. Measured, it does the opposite: the unfiltered estimate is 0.597 [0.563, 0.632] against the filtered 0.553 [0.514, 0.593]. The filtered value is the one quoted, because disattenuation is unstable at the very low reliabilities the filter removes. Two mismatches belong with it. The filter inside that calculation is the inherited $\cos > 0.2$ line rather than the 0.109 line the training set actually uses, so the pair brackets the trained-on population without reproducing it. And the generation evaluation is selected differently again — it applies the 0.109 line to the `train` split and leaves the held-out splits unfiltered. That second mismatch is load-bearing rather than cosmetic, and it is taken up on its own below.

The intervals omit generation variance. The confidence intervals in this thesis are cluster-robust with each condition’s score treated as fixed. The headline gaps carry two-way intervals over *cell line* and *treatment well*, with a Student t reference on $\min(G_{\text{line}}, G_{\text{well}}) - 1$ degrees of freedom, and are reported a second time clustered on the *drug*; the transfer coefficient carries a multiway dyadic interval; a few tables retain a one-way clustered bootstrap and name its unit. They therefore capture between-cluster variance and not the variance introduced by sampling the model’s output. With four sampled generations per condition at temperature 0.8 in the residual-frame arms and eight in the expression-frame ones (§3.1), that second source is not negligible, and no across-seed replication of the generation arms was run — not deferred, not planned. Every interval quoted from a generation arm is therefore a within-seed interval and should be read as conditional on the seed: it covers the sampling of conditions, cell lines and wells, and it does not cover the sampling of the model’s output.

The one measurement here that was repeated across seeds gives a sense of what is being left out, though it cannot calibrate it. The activation probe of Section 3.6 returned +0.0197, +0.0118 and +0.0085 on three seeds, each excluding zero. That is a different instrument on a different quantity and its spread does not transfer to NIR; what it does establish is that a factor of two between seeds is realistic in this project, which is reason enough not to read a generation interval as though it covered seed variation as well.

A swing that looked like generation variance and was not. Two superseded evaluations of the held-out-drug control returned values of Δ_{scr} far enough apart to be recorded as evidence of seed instability. Only one of the pair survives in the results archive, so the size of the swing is no longer quotable here; the diagnosis, however, was wrong — or rather, it named the trigger and missed the cause. Both figures are gaps against the **opposite** swap, which Section 3.11 shows is not a null: it contains a term measuring how far the *partner* drug was pushed the wrong way, a quantity that has nothing to do with the target and that changes with which partner happens to be drawn. An arm whose value is partly a nuisance quantity will move when the draw moves. On the canonical evaluation the same control reads +0.0379 against **opposite** and -0.0315 $[-0.0836, +0.0206]$ against the neutral stratum, the interval clustered two-way on cell line and well. Two mechanisms are involved and both are now addressed: generation was genuinely unseeded — `do_sample` advanced from global torch state, so the two arms of a contrast were drawn from different points of one stream — and each call now seeds its own generator from (condition, arm, batch); and the comparator has been corrected. The general caution above stands regardless; this particular swing is no longer evidence for it.

Pseudobulk resolution. Residuals are condition-level pseudobulks over ≥ 40 treated cells. The single-cell heterogeneity of the response is not modelled in this frame, and the project’s own earlier results show why: at single-cell resolution the drug moves a cell less than noise does ($\text{SNR} \approx 0.75$, median 0 differentially expressed genes at FDR $q < 0.05$).

Cross-plate absolutes. The residual-space scores are computed with comparison sets grouped by cell line rather than by plate. The *difference* Δ_{scr} is leak-immune, because both arms share the same control cell, so the grouping choice does not touch Δ_{scr} itself. The absolute NIR values do not travel. They are not comparable to the within-plate tables elsewhere in the chapter, and the rescore that would make them comparable — regrouping the generation-free arms within plate, which needs no GPU time — was not run, so the size of the correction is unknown for this frame. Table 4 is the only guide to how large it might be, and it suggests not small: measured in the expression frame, holding the plate constant moves a drug-agnostic mean baseline from 0.399 to 0.161 and the precision reference from 0.625 to 0.575. Residual-frame absolutes should therefore be

read within their own frame and never set beside a within-plate number.

Frame coverage. The drug-specific residual is $\|r\|/\|s\| = 0.786$ of the response, against 0.620 for the generic program and batch, and the two are essentially orthogonal: the shift’s energy decomposes into residual 0.621, generic 0.390 and a cross term of -0.011 , which sum to 1.000. So this frame covers 62.1% of the response *energy* and every result here should be read within that bound — a real but not total share. The generic 39.0% is precisely the component the field’s absolute-accuracy metrics are dominated by, which is why a drug-agnostic predictor scores so well on them.

The splits are scored on different populations. The within-well split-half precision reference differs sharply across the three splits as they are actually scored — 0.956 on `train`, 0.824 on `unseen_combo`, 0.836 on `unseen_drug` — and the obvious reading of that spread is the wrong one. It is not that held-out conditions carry less recoverable signal. It is the selection asymmetry noted above: `train` is filtered at the 0.109 reliability line, which costs it one of its 300 conditions, and the held-out splits are not filtered at all, so their scored populations still contain conditions whose own two halves disagree outright. Impose the training line on all three and the references agree — 0.955, 0.961 and 0.960, a spread of 0.006 where there had been 0.132. Matched on reproducibility, a held-out condition is as recoverable as a trained one.

The consequence runs against this thesis rather than for it. *Those conditions carry less signal* is not a reason to discount a held-out shortfall, because once the same line is applied to both they carry the same signal; where the raw spread has been used to soften a held-out result, that softening is withdrawn. What survives is narrower and still binding: the three splits are scored on populations selected by different rules, so a raw NIR or a Δ_{scr} read across them is not answering the same question on each, and the training-exposure contrast of Section 3.11 inherits that directly. The matched-population rescore was carried out for the precision reference only and was not run through the model’s own arms, so what the held-out shortfall becomes once the populations agree is not reported here.

4.2 Limitations of the unseen-drug evaluation

Two defects, both of which must be stated rather than papered over. A third — the comparator — was found and corrected rather than merely stated, and Section 3.11 reports the arm against the neutral stratum throughout; the two below survive that correction.

It is not a clean held-out-drug design. The prompt contains `Mechanism: {moa}`, so holding out a drug still hands the model the held-out drug’s class label. The cor-

rect design is leave-one-*mechanism*-out. The size of the leak is measured rather than assumed. Scrambling the **Mechanism** field alone moves the model’s NIR by +0.030 ($n = 738$; $[-0.008, +0.069]$ clustered on drug and well), a channel small enough that it is not distinguishable from zero. A mechanism-only lookup is the same measurement from the other side, and is equally inconclusive: pooled over the $n = 218$ conditions on which both are scored, $\text{model} - \text{moa_lookup} = +0.044$ ($[-0.048, +0.135]$ clustered two-way on cell line and well), an interval that spans zero, so the two are not separated and the comparison settles nothing in either direction.

The exception is the split where it would matter. On **unseen_drug** the mechanism lookup is ahead of the model by 0.099 on 44 conditions, with an interval spanning zero ($[-0.222, +0.008]$, clustered on cell line) and far too wide to settle anything. That is not evidence the leak is doing work; it is a measurement with no power at exactly the point where the design defect bites. So the position is: the design is wrong, the arm cannot say what the wrongness cost, and its proper use is as an instrument control rather than as the basis for a positive unseen-drug claim.

It is underpowered by construction. The confidence interval half-width on the **unseen_drug** gap is ± 0.052 clustered two-way on cell line and well, and ± 0.073 clustered on the drug; the smallest effect the arm would detect at 80% power is ≈ 0.079 . The largest effect the mechanism channel could deliver is the distance from chance to **moa_lookup**, which scores 0.552 pooled across splits against a chance value of 0.500, so ≈ 0.052 . The arm therefore *cannot* detect the largest effect it could plausibly be measuring.

That cuts both ways, and the two directions are not symmetric. As a *control* the arm is not falsified: the question there is whether it returns zero when the drug name carries nothing the model was fine-tuned on, and the estimate -0.0315 $[-0.1049, +0.0419]$ (clustered on drug) has an interval covering zero. It does not sit *at* zero — it is negative and of the same order as the largest effect the arm could plausibly be measuring — so what it supplies is the absence of a contradiction, not a confirmed null. As *evidence of absence* it carries almost nothing, because an effect of the size the mechanism channel could deliver would sit comfortably inside the interval. That distinction is what licenses using the arm to validate the comparator in Section 3.11 while declining to use it to argue that unseen-drug transfer is zero. And “unseen” is narrower than it sounds in any case: it means held out from Tahoe fine-tuning only. Pythia-1b’s own pretraining is natural-language text in which compound names occur, and the prompt still supplies **Mechanism**.

4.3 Directions

The arm to build carefully rather than eagerly. The obvious continuation of Section 3.13 is a channel-conditioned prompt: append protein targets to the instruction, retrain, and evaluate on a leave-one-mechanism-out split. We do not recommend building it eagerly, and the reason is a shape of failure rather than a settled measurement. Section 3.14 does not resolve *which* span of the prompt is read. On the canonical run, swapping the drug *name* alone moves the output by +0.1026 [+0.0625, +0.1427] ($n = 1394$) while swapping the *mechanism* field alone moves it by +0.0303 [−0.0078, +0.0685] ($n = 738$), both clustered on drug and well. All of those arms swap to an opposite-signature partner, which Section 3.11 shows is not a null, so each gap carries a comparator term — and one that runs in opposite directions across the two rows, inflating the name arm, whose partner scores 0.434, and suppressing the mechanism arm, whose partner scores 0.520. Read instead as each row’s own distance from chance the ordering inverts (+0.0367 on the name row against +0.0506 on the mechanism subset), so the split between the two spans is unresolved rather than a direction. Pending a re-run against the neutral partner, the case against a target-conditioned prompt rests not on demonstrated mechanism-blindness but on the size of the readout deficit: on the common support the model covers 13.9% of the distance from chance to the within-well bar that a context-blind lookup covers entirely.

What would be worth building instead is an intervention that makes the drug *causal* rather than merely present — a learned per-drug modulation of the transformation from control to treated, so that the drug parameterises the computation instead of sitting in the residual stream as an input the readout can round away. That is a research architecture rather than a prompt change, and it is outside the scope of what remained here.

Two channel questions that remain open. Target and mechanism overlap, so whether target adds anything *over* mechanism is unresolved — and it matters, because the prompt already exposes mechanism and whether the model uses that field is itself unresolved. Separately, the channel gate measures *similarity transfer*: an unseen drug resembling a seen drug with the same target. It does not test a *compositional rule* — that a named target’s own pathway moves, which would apply to any drug whose target is named, including one with no similar drug in training. The second is the stronger claim, it does not reduce to a lookup, and it is untested. Both questions sat behind a defect in the gate itself — the partner pool was not restricted to training drugs, so a held-out drug’s response could be predicted from a partner that was also held out — and the gate has since been re-run with the restriction in place. The repair narrowed what may be claimed, and the clustering axis narrows it again: protein target clears the relevance margin on the cell-line axis, misses it on the drug axis, and rests on 18 annotated drugs; mechanism misses its lower bound by 0.0021 on the cell-line axis and spans zero on the

drug axis; chemistry’s interval straddles the margin on both. Only the target channel supports a statement that unseen-drug generalisation is reachable through retrieval, and only as a direction worth testing rather than a measured effect of useful size. The binding limitation is coverage: 18 drugs cannot settle it, and the experiment that would is a target-annotated panel broad enough to cluster on.

Unseen cell lines. A cell line never seen with any drug is a different generalisation axis from the held-out pairings tested here, and one not blocked by the structure gate: it asks whether the control cell’s state is used *compositionally* with the drug. The transfer coefficient scopes it only loosely: 44.7% [40.7%, 48.6%] of the drug-specific residual is not shared across contexts, and no structure was recovered in that unshared part. That figure is not a drug×cell-line interaction share — it mixes cell line, dose, well and estimator (Section 3.9) — so it bounds what a per-drug constant leaves on the table without establishing that a new cell line is the axis the remainder lies along. What can be said is that no consistent per-line direction was found, so the achievable target on a new cell line plausibly reduces to recalling the drug’s per-drug constant — which a training-only lookup already does at 0.913. The axis is therefore less promising than it appeared before Section 3.9 was measured, though the argument is weaker than a coefficient-based one.

Objectives that were considered and rejected. Recorded so they are not re-litigated:

- **Reinforcement learning with a contrastive reward.** Its mechanism — a per-sequence reward, so the token-dilution ratio never enters the gradient — attacks a bottleneck that the residual encoding has already lifted ($34.6 \rightarrow 118.2$ differing tokens); and its reward optimises similarity to the true residual, which is precisely what `drug_lookup` already achieves at 0.913. Its best case is convergence to the lookup.
- **A differentiable profile-level contrastive loss.** The natural construction, $\text{score}(g) = \sum_i P(\text{token}_i = g) w(i)$, presumes $\text{gene} \approx \text{token}$, but 86.9% of the 946 panel genes share a first sub-token under this tokeniser (mean 3.254 BPE tokens per gene, only 8 single-token). Under teacher forcing the difficulty compounds: the residual target *is* the drug signature, so the true prefix identifies the drug within a handful of symbols and the gradient vanishes exactly at the drug-blind solution the loss exists to penalise.
- **Loss re-weighting toward differentially expressed genes.** Re-weighted cross-entropy against a single target contains no negatives, so nothing in it forces the output for one drug to differ from another’s. It therefore attacks token dilution and leaves the discrimination failure exactly where it was — and the discrimination

failure is what the measurements establish, whether or not the readout also turns out to be blind to direction specifically, which Section 3.6 declines to establish.

One model size, and the smallest one. Every result here comes from a 1B-parameter backbone. C2S-Scale releases checkpoints up to Gemma-2 at 27B, and we tested none of them, so the most predictable question an examiner can ask — does it work if the model is bigger — has no answer in this thesis.

Two things can be said about what we would expect, and they point in different directions. The constraint this work points to is not capacity. It is that the target tokens are overwhelmingly shared across drugs — only 17% of them differ between two drugs in the same cell line under the standard target, against 59% under the residual encoding (Section 3.10) — so very little of the training signal can be carrying drug identity. That last step is a hypothesis about the optimisation rather than a measurement: no token-level loss or gradient attribution was computed, so no fraction of the gradient itself is claimed. A larger model receives the same diluted signal from the same objective, and re-encoding the target lifted the effect while leaving capacity untouched — which points away from capacity without establishing what it points towards. The re-encoding changed aggregation, control referencing, sign coding and length at the same time, and a budget-matched optimal-transport target lifts the effect too, so Section 3.10 declines to identify the tokenisation as the constraint that binds.

The genuinely different argument for scale is not capacity but *prior knowledge*. A 27B model pretrained on far more text plausibly knows, as a fact about the world, that a named compound inhibits a named kinase — knowledge a 1B model may simply lack. That matters here because Section 3.14 leaves unresolved whether the model uses the *mechanism* field it is already handed — swapping that field alone moves the output by $+0.0303$ $[-0.0078, +0.0685]$ ($n = 738$), an interval containing zero in an arm whose comparator sits above chance and so suppresses that gap — and a model with real pharmacological knowledge is exactly the case where that might change. This is a sharper hypothesis than “bigger is better” and it is testable cheaply with a low-rank adaptation of a mid-sized checkpoint rather than a full fine-tune. We did not run it, and it is the first thing we would run with more time.

The starting checkpoint was never prepared for this task. We fine-tune from C2S-Scale-Pythia-1b-pt, whose cell-sentence pretraining corpus is atlas data alone — CELLxGENE and the Human Cell Atlas, no perturbation corpus of any kind — and whose released capabilities do not include perturbation prediction. This cuts two ways and both should be stated. It is why the held-out splits are as clean as they are: a drug withheld here entered neither the fine-tuning data nor the generic that defines its target, the split being assigned from metadata before any expression was aggregated,

so no Tahoe perturbation response for it can have been memorised. But it also means the failure we report is a failure of *this* starting point, and a reader may reasonably ask whether a checkpoint already tuned for perturbation prediction would behave differently. Within the C2S-Scale line that question cannot be answered from what is public: the perturbation-specialised model there is a Gemma-3 4B checkpoint that was not released, and the available C2S-Scale checkpoints are the pretrained ones. Outside that line the question is answerable, and we did not answer it. Tahoe-x1 [13] is a perturbation-trained foundation model built on this same atlas and public before this thesis was submitted. It is not a cell-sentence model, so it does not isolate the representation variable this work is about, but it is the obvious external check on whether a readout failure of this kind survives perturbation-specific pretraining, and it was not run here. That is a gap in this work, not a gap in the public record. What we can say is narrower and worth separating from the broader claim: every arm in this thesis starts from the same base and differs only in the variable under test — with one material exception, the consensus arm, which was stopped at roughly 60% of an epoch with its tier-1 NIR unscorable (Table 20) — so the *comparisons* between arms are unaffected by this choice even where the absolute level is not, and that arm’s null is quoted with the qualification wherever it is used.

Generality. Everything here is one backbone on one atlas in one representation. Whether the readout-specificity failure is a property of cell-sentence models specifically, or of autoregressive conditional generation over transcriptomes generally, is open. The mechanistic protocol of Section 3.4 — purified subspace, removal gate, positive control, matched-norm swap — transfers unchanged to any model with an accessible residual stream, and is probably the most reusable artifact of this work.

5 Conclusions

This thesis asked whether a cell-sentence language model fine-tuned on the largest available single-cell drug-perturbation atlas learns drug-specific biology. The model measured here does not. That subject is deliberately narrow — one 1B C2S-Scale Pythia checkpoint, five fine-tuning arms cold-started from it with identical hyperparameters at a single learning rate, none trained for longer than one epoch, and no larger checkpoint tested (Section 4) — and nothing below widens it to conditional generative models in general. What travels past this checkpoint is the chain of results establishing why it fails, how far the failure can be repaired, and what bounds the repair — and, before any of them, the evaluation result that had to come first.

The evaluation had to be fixed first. The metric the field standardises on is exploitable: a zero-information baseline placing every gene at the middle rank scores $\text{DE-}\Delta r = 1.000$, above both the noise ceiling and the model. Under a Dynamic Range Fraction analysis computed across plates, and replicated within plate, only NIR — a discrimination metric, in which the generic response cancels — has a positive dynamic range fraction: $+0.635$ across plates and $+0.545$ within plate. All four prediction metrics tested come out negative — -0.163 to -0.686 across plates, -0.104 to -0.515 within plate — inverted, not merely uninformative. The released calibration artifacts clustered the drug rows as though each row were its own cell line, so their intervals were far too narrow. Both arms have been rerun with the bootstrap clustered on cell line and the null imposed: $\text{DRF}(\text{NIR}) = +0.635$, $[+0.604, +0.663]$ over 25 cell lines across plates, and $+0.545$, $[+0.516, +0.569]$ over 27 within plate, each at Holm $p = 0.002$ across the family of five. That p is still at the resampling floor, but now because the estimate sits some forty cluster-robust standard errors from zero rather than because each row was counted as its own cluster. The intervals are percentile cluster bootstraps without a small-sample correction, and plate is crossed with cell line while only the cell-line channel is clustered, so they are robust on that axis alone. The residual frame is audited separately rather than assumed to inherit that result, and reads $\text{DRF}(\text{NIR}) = +0.704$ $[+0.667, +0.736]$ over 42 cell lines against a zero-drug-information control, so both frames this thesis scores in are calibrated by measurement rather than by assertion. The metric audit is the most portable *result* here, since it applies to any work scored on $\text{DE-}\Delta r$ -family metrics. It is not the most reusable *object*: that is the mechanistic protocol of Section 3.4 — purified subspace, removal gate, positive control, matched-norm swap — which needs only an accessible residual stream and which Section 4 nominates on exactly that ground.

This also settles how the thesis enters the dispute of Section 2.4. Miller et al. [20] argue that reports of deep models failing to beat uninformative baselines reflect miscalibrated evaluation rather than model failure, and on genetic perturbation they support that with

calibrated-metric benchmarks. We port their framework, confirm its methodological claim emphatically — metric choice dominates, and only the rank-based metric survives — and then find that on drug perturbation the model is at chance *on the metric their own analysis endorses*. That is the strongest form the sceptical result can take: it cannot be answered with the objection their paper raises. Their conclusion is not overturned so much as bounded — it does not extend from genetic perturbation to this task, and the per-drug-constant finding below suggests why.

The model is drug-blind, and the representation is not the reason. On the calibrated metric, within plate, the model scores 0.498 against a 0.576 within-well split-half precision reference — and a control-copy that reproduces the untreated cell, carrying no drug information whatever, scores 0.504. The obvious objection is that this is a noise floor rather than a model failure, and the difficulty stratification answers it: on the 204 conditions whose own split-half ceiling clears 0.8 and averages 0.938, where the task is provably winnable, the model reaches 0.768 and the same zero-drug-information control-copy reaches 0.766 — a paired contrast of +0.002 $[-0.026, +0.028]$ clustered over cell lines. The model is blind precisely where it could see, which is what makes this a result about the model rather than about the data.

Three further instruments agree, and what they deliver is a bound rather than a zero. Replacing the drug name and mechanism in the prompt with a different, different-mechanism drug’s leaves the output unchanged to within +0.014 $[-0.016, +0.042]$ on that same identifiable subset; the truth-free invariance gap is 0.000 $[-0.019, +0.016]$, so two generations for different drugs resemble each other as closely as two generations for the same one; and forced-choice grading against the condition’s own truth sits at chance, far below the value the same test returns for a real sample of that condition’s cells. None of these shows the drug effect to be zero. Each bounds it inside an interval containing zero, which is the strongest statement measurements of this size support — and a bound is what the argument needs, since an effect four independent instruments cannot resolve is not one a prediction task can use.

The representation is not where the failure lives. With cell line, plate and dose projected out, drug identity becomes *more* linearly decodable with depth, reaching 0.88 at layer 16 against an 8% chance floor, while the purified drug subspace’s variance share stays flat at ≈ 0.03 against the cell line’s ≈ 0.6 . The network does consult that region — ablating it costs $3.6\text{--}18.6\times$ a matched-dimension random subspace — but at a bit player’s weight, 9–37% of what ablating cell line costs at layers 2–9. So the dissociation is not between something encoded and something never touched. It is between how explicitly the drug is written and how little the process that writes the output does with it.

How far that can be pushed toward a mechanism is limited, and the limit should be stated plainly. Three earlier designs for localising the failure by activation swapping

were retracted once their controls were found confounded — mismatched geometry, an ablate-then-add that made the treatment a restoration, and an off-manifold direction that inflated the normaliser — and what replaces them (Section 3.6) is a small, replicated, single-layer effect present in one training arm and absent in another, whose survival of multiplicity correction depends on which rung of the displacement ladder the correction is applied to. That supports the modest reading — the generation readout consults the drug representation at very low gain — and does not support a claim about privileged subspaces or global workspaces, which is offered here as an interpretive frame rather than a result. It is the weakest load-bearing element in the thesis, and no argument here rests on it alone: the one architectural intervention this thesis proposes takes its warrant from the size of the readout deficit, not from this.

The target’s tokens barely distinguish one drug from another, and the construction that changed that produced the first non-null effect. Because a cell sentence ranks genes by expression, the targets for two different drugs in the same context are nearly the same string: 82.7% of the top-200 genes are shared, so only about a sixth of the target distinguishes one drug from another. Since the training loss weights every token position equally, the natural inference is that very little of the gradient can be carrying drug identity — but that inference is a *hypothesis* about the optimisation, not a measurement of it. What is measured here is the overlap of the target gene sets; no token-level loss or gradient attribution was computed, and the claim should be read at that strength. Its warrant is that it predicts correctly, twice. It predicts that changing *what the target is* while leaving *what the tokens encode* alone will not help — and an ε -ladder of target constructions, from a single state-matched treated cell to a full pseudobulk collapse, stays in the near-null regime under the random-partner test at the budget and comparator that test was run at. That qualification is load-bearing, and the paragraph below pays it: re-scored at a matched budget against a neutral partner one of the three rungs is not null, and a second arm was stopped mid-epoch and so evidences little, which leaves the ladder’s negative resting outside that scoring on the single-cell rung alone. And it predicts that changing what the tokens encode will help: re-encoding the target as the drug-specific residual raises the differing fraction to 59% and produces the first non-null drug effect in the project, with a gap that grows monotonically with swap dissimilarity.

Two qualifications stop that from being a demonstration that the tokenisation is the binding constraint, and both are load-bearing. The residual construction changes aggregation, control referencing, generic subtraction, sign coding and sequence length at the same time as it changes token divergence, and no matched-budget factorial here separates them. And *first* means first in sequence, not uniquely responsible. Re-run at the residual arm’s 1400-token budget, the optimal-transport control that had been read as inert clears zero against its neutral comparator, +0.0292 [+0.0144, +0.0441]; at the 600-token cap

it was originally scored at, the same gap read $+0.0144$ $[-0.0051, +0.0338]$ and spanned zero. The single-cell control stays flat at both budgets. That correction is also a transferable measurement lesson: a generation token cap can manufacture a null, and the same cap had already halved the residual arm’s own effect before anyone noticed it was doing anything.

What that drug use amounts to, and what it does not. One correction here generalises past this project, and it belongs in the record as a finding rather than as an erratum. A label-swap ablation — rename the condition, leave every other byte identical, read the drop — is the standard way to ask whether a conditional model uses its condition, and the swapped arm is normally read as a null. It is not one. Across the three swap strata, which are defined in advance by residual cosine and not by their scores, the scrambled model’s own NIR measures 0.550, 0.501 and 0.423 as the named partner moves from aligned through orthogonal to anti-aligned with the target. A comparator whose score is a monotone function of which drug it was handed is an active treatment tracking that drug, and a gap read against a deliberately dissimilar partner therefore adds “how much naming a lie hurts” — $+0.077$ against the anti-aligned partner — to “how much naming the truth helps”, which is the only term anyone wanted. Any label-swap ablation scored by a discrimination metric inherits this, and the remedy is cheap: report the comparator’s own score, and read magnitudes only off a stratum that measures at chance.

It cost this thesis a claim, which is how it was found. An earlier form of this paragraph claimed the drug use generalises, on a holdout keyed at the (drug, cell line) level where pairings never seen in training scored $+0.1002$ $[+0.0661, +0.1368]$ and held-out *drugs* came out null. Two corrections removed that claim: the targets were rebuilt so the generic is fitted after the split rather than over the holdout, and the non-neutral comparator was replaced by the neutral stratum. The held-out-drug arm is what exposed the comparator, by returning a positive result on a split where a positive result is impossible.

Measured against the neutral comparator on the rebuilt targets, trained pairings score $+0.0898$, excluding zero clustered on cell line and well $[+0.0382, +0.1415]$ and on drug and well $[+0.0291, +0.1506]$; held-out pairings score $+0.0332$ and held-out *drugs* -0.0315 $[-0.0836, +0.0206]$. Transfer to held-out combinations is *not* established. Its interval excludes zero when the bootstrap is clustered on cell line and well $[+0.0069, +0.0594]$, and does not when it is clustered on drug and well $[-0.0031, +0.0695]$ — and drug is the axis along which a claim about transferring a *drug* signature has to generalise. The split is also mostly recombination rather than novelty: 30 of its 34 drugs and 39 of its 42 cell lines also appear in the trained arm, and 82% of its conditions pair a training drug with a training cell line. The memorisation premium that earlier carried this argument goes with it. Pooled it is $+0.0567$ $[-0.0014, +0.1148]$, cluster-robust $p = 0.056$; restricted to the 30 drugs common to both splits it changes sign, to -0.0606 with $p = 0.18$. It is confounded

in any case, because the splits are not equally hard: the within-well split-half precision reference is 0.956 on the trained split against 0.824 and 0.836 on the two held-out ones, a spread of 0.132, so gaps measured on them are not directly comparable.

What survives is narrower than what was claimed for it. The comparator-free test — the slope of NIR on the cosine similarity between the real response of the drug named in the prompt and that of the true one — is positive on every split: +0.373 on trained pairings, +0.336 on held-out pairings, +0.342 on held-out drugs. This does *not* show that the model knows held-out compounds. The slope is fitted over the three scrambled *partner* prompts alone (**near**, **orth**, **opposite**); the prompt naming the true target drug is never a point on the line, and on the held-out-drug split 82% of the 597 points name a drug that also appears in the trained arm. What it shows is that the model responds to the identity of a named trained drug even when it is scored in a context whose true drug it has never seen — a real effect, and a much weaker one than zero-shot drug knowledge. Alongside it, on both held-out splits the model scores at chance when its output is ranked against other drugs in the same cell line (NIR = 0.533 [0.492, 0.575] and 0.459 [0.404, 0.514]), and its output is no closer to the true drug’s response than to other drugs’ (own-minus-other cosine +0.011 and −0.014, neither interval excluding zero). The map exists and is coarse: what the model has captured is close to a per-drug average response — which is also why the lookup below beats it, since the lookup *is* that average, measured directly.

But a lookup table wins — and what it beats is not a ceiling. A per-drug mean residual measured in other cell lines scores 0.913; the model reaches 0.537 overall, and 0.549 on the $n = 1192$ conditions on which both are scored. The within-well split-half precision reference is 0.854 over all conditions and 0.857 on that common support, so the lookup does not approach that reference — it *exceeds* it, covering 115.7% of the distance from chance to that bar against the model’s 13.9%. That excess is a property of the bar rather than a headroom estimate, for the three reasons Section 3.12 gives. The natural reading of those numbers is that the drug’s response is essentially a per-drug constant, that a conditional model therefore has almost nothing to condition on, and that the correct formulation of the task is retrieval rather than generation. We initially drew exactly that conclusion, and it is wrong.

It is wrong because NIR compresses near the top of its range — the same class of failure this thesis opens by documenting, arriving a second time in our own preferred instrument, at the top of the scale instead of the bottom. Measured in cosine space with measurement noise divided out, the transfer coefficient is $\mathcal{T} = 0.553$ [0.514, 0.593], against a simulated null in which the residual is entirely shared across contexts — a null this estimator recovers at 1.001 — and with a structure-matched negative control at −0.008 [−0.036, +0.020]. So roughly 45% [41%, 49%] of the drug-specific residual is *not shared across contexts*. That is deliberately weaker than the reading in an earlier draft,

which took $1 - \mathcal{T}$ to be a drug×cell-line interaction share. It is not one. The estimator pools pairs that differ in cell line, in dose and in well, and it was run with `same_dose_only` off, `loo_generic` off and a reproducibility threshold of 0.2; on pairs that vary dose within a single cell line it returns $\mathcal{T} = 0.707$ [0.632, 0.782], so nearly 0.3 of the residual still fails to transfer with the cell line held fixed. The unshared remainder therefore mixes interaction with dose, well and estimator slack, and cannot be attributed to a drug×cell-line term without a design that separates them. What can be said is that the per-drug constant is what the lookup captures and what the model gropes toward, and that the unshared remainder is reached by neither.

So the honest statement of what the model does is sharper, and less flattering, than “it misses the part that is not shared across contexts”. It fails to recover the *shared* part that a numpy dictionary obtains for free, and the unshared remainder — which mixes cell line, dose, well and estimator, and in which no learnable structure was found — is reached by no method here, ours least of all. Two different failures, and collapsing them undersells both.

That conclusion connects this work to an assumption the field has carried since the Connectivity Map: that a compound induces a characteristic signature stable enough across cellular contexts for signature-similarity querying to be meaningful. This thesis measures that assumption directly, with a discrimination metric and a within-well split-half precision reference, at single-cell resolution — and the modern finding that a per-perturbation mean is hard to beat may be less a warning about weak baselines than a restatement of it.

It also places this work as the single-cell, generative instance of DrEval’s critique [19]. DrEval audits drug *sensitivity* prediction and finds that deep models barely outperform a naive additive predictor of mean drug and cell-line effects. Our `drug_lookup` is that predictor’s drug term, and the finding replicates on a 946-dimensional transcriptomic response. On the 1,192 conditions where both are defined the lookup scores 0.913 against the model’s 0.549, with a within-well split-half precision reference of 0.857 that the lookup exceeds. Three of DrEval’s six pathologies — pseudoreplication, biased evaluation metrics, and missing ablation studies — turn out to structure the entire methodology of this thesis, which arrived at clustered bootstraps over cell lines, a discrimination metric immune to between-drug variance, and a token-level scramble ablation independently. That convergence is worth recording: the controls a critical benchmark demands are the same controls a single negative result needs in order to mean anything.

Where the opportunity actually is. Two further measurements bear on where it lies — one unresolved, the other now settled on a corrected run — and between them they redirect what follows from this work.

The part of the drug-specific residual that fails to transfer across contexts is large —

roughly 45% of the similarity scale $\mathcal{T}$ is measured on, which is not a variance share — and we could not find structure in it. No cell line modulates different drugs in a consistent direction. Within-line consistency is -0.0149 $[-0.0251, -0.0050]$ against $+0.0003$ $[-0.0071, +0.0081]$ across lines, an excess of -0.0153 ; the within-line interval lies entirely below zero, so there is no positive consistency to find. A direct test of whether mechanism-matched drugs share a context term with a given cell line returned nothing detectable.

Those are two negative results and they do not license the conclusion that the unshared part is irreducible. The second test bounds its own power severely: the mean split-half reliability of κ in this atlas is 0.326, so the quantity being correlated is barely estimable, and mechanism is partially confounded with plate, by an amount that was never measured for this test. The honest statement is that the structure of the unshared part is *unresolved here*, and that resolving it needs more cells per condition or a plating design that crosses mechanism with plate — neither of which is an analysis choice. What can be said without qualification is that a lookup scoring 0.913, above the 0.857 within-well split-half precision reference on that same common support, already outscores anything conditional built here, so on seen drugs nothing we built beats retrieval.

But a lookup requires the drug to have been seen. For drugs that have not been, the gate has now been re-run with its partner pool restricted to training conditions — the construction the question actually asks about, since a partner the model never saw cannot stand for what training supplied. On that run one channel clears the pre-set relevance margin of 0.03 and two do not — and the one that does, does so on one clustering and not the other. Protein target clears it on all three constructions clustered on cell line and well: count-matched $+0.1498$ $[+0.0958, +0.2038]$, plate-matched $+0.1351$ $[+0.0745, +0.1958]$, and the stricter different-plate requirement $+0.1322$ $[+0.0648, +0.1996]$. Clustered on the drug — the unit a claim about unseen drugs generalises over, and one carrying only 18 annotated drugs — the same gaps read $[+0.0536, +0.2460]$, $[+0.0287, +0.2416]$ and $[+0.0199, +0.2445]$: still clear of zero, but clearing the margin on the count-matched null only. Mechanism excludes zero on all three (plate-matched $+0.0805$ $[+0.0279, +0.1332]$) yet its lower bound falls 0.0021 short of the margin, so it is neither live nor equivalent to its null but inconclusive on a knife-edge. Chemical structure carries the narrowest interval of the three and still cannot be placed: it excludes zero against a plate-matched null ($+0.0261$ $[+0.0035, +0.0487]$) and spans zero once the plate is held apart ($+0.0118$ $[-0.0097, +0.0333]$), so the effect is small and well estimated rather than unmeasured, with the margin falling inside it. Table 16 carries all three constructions. What these gaps establish is that protein-target annotation carries transcriptional signal of a useful size in this atlas; that mechanism carries signal whose size cannot be placed on either side of usefulness; and that chemistry carries a small one. None of the three may be written as closed, and none may be written as live without naming the axis: target is live clustered

on cell line and well, and short of the margin clustered on the drug.

This overturns a pre-registration made earlier in this project. Two design documents concluded that unseen drugs were unreachable, and both inferred it from a gate that tested chemical *structure* — the channel our own drug-knowledge specification ranked lowest, and the only one of the three whose interval spans zero once the plate is held apart. The inference generalised from the weakest channel to the strongest, and it was wrong. The protein-target annotations that carry the signal were sitting unread in the atlas’s own metadata throughout.

So the thesis does not end in closure. It ends with the opportunity located. On seen drugs a generative model has nothing here to add over retrieval — not because the unshared part is irreducible, which is exactly what the two negative results above decline to establish, but because nothing built here reached it and a lookup already takes everything else. On unseen drugs, where retrieval is impossible by construction, protein target is the only one of the three channels that clears the pre-set relevance margin against a different-plate null with the partner pool restricted to training drugs; mechanism excludes zero there too and falls short of the margin. That makes target annotation the one place a model might do something a dictionary cannot — with the size of the opening unresolved, since on the drug axis its interval clears zero but not the margin, on 18 annotated drugs of which two are held out. That is also the setting the benchmark literature calls Leave-Drugs-Out, and the one that matters for drug design.

The controls are most of the result. In a literature whose central problem is that uninformative baselines score well, a negative result is worth precisely what its controls are worth, and that — not any single number above — is what this work would defend first. The substantive results here rest on controls that were added in response to specific objections: plate-matched comparison sets after a zero-information baseline more than doubled its score when the plate was free to vary; a swap-distance sweep after it was pointed out that scrambling to a similar drug is a weak test; a removal gate before trusting any ablation null; matched-norm injection before reading a swap; clustered bootstraps over cell lines rather than conditions; and a pre-registered interpretation written before the decisive result, whose middle branch is the one that fired. Several claims were retracted along the way and are struck rather than deleted; that list is the reason the surviving claims can be read at the strength they are stated, and no stronger.

References

- [1] Jesse Zhang et al. “Tahoe-100M: A Giga-Scale Single-Cell Perturbation Atlas for Context-Dependent Gene Function and Cellular Modeling”. In: *bioRxiv* (2025). Preprint, posted February 2025 (v3, 10 May 2025); not peer reviewed. Tahoe Therapeutics, formerly Vevo Therapeutics. Dataset: <https://huggingface.co/datasets/tahoebio/Tahoe-100M>. DOI: 10.1101/2025.02.20.639398.
- [2] Daniel Levine et al. “Cell2Sentence: Teaching Large Language Models the Language of Biology”. In: *Proceedings of the 41st International Conference on Machine Learning (ICML)*. Vol. 235. Proceedings of Machine Learning Research. Preprint: bioRxiv 2023.09.11.557287. PMLR, 2024, pp. 27299–27325.
- [3] Syed Asad Rizvi et al. “Scaling Large Language Models for Next-Generation Single-Cell Analysis”. In: *bioRxiv* (2025). Preprint, v1 posted 17 April 2025; author list above is v4, posted 1 January 2026. Not peer reviewed. DOI: 10.1101/2025.04.14.648850.
- [4] Stella Biderman et al. “Pythia: A Suite for Analyzing Large Language Models Across Training and Scaling”. In: *Proceedings of the 40th International Conference on Machine Learning*. Vol. 202. Proceedings of Machine Learning Research. PMLR, 2023, pp. 2397–2430. URL: <https://proceedings.mlr.press/v202/biderman23a.html>.
- [5] Haotian Cui et al. “scGPT: toward building a foundation model for single-cell multi-omics using generative AI”. In: *Nature Methods* 21.8 (2024), pp. 1470–1480. DOI: 10.1038/s41592-024-02201-0.
- [6] Christina V. Theodoris et al. “Transfer learning enables predictions in network biology”. In: *Nature* 618.7965 (2023), pp. 616–624. DOI: 10.1038/s41586-023-06139-9.
- [7] Charlotte Bunne et al. “How to build the virtual cell with artificial intelligence: Priorities and opportunities”. In: *Cell* 187.25 (2024), pp. 7045–7063. DOI: 10.1016/j.cell.2024.11.015.
- [8] Mohammad Lotfollahi, F. Alexander Wolf, and Fabian J. Theis. “scGen predicts single-cell perturbation responses”. In: *Nature Methods* 16.8 (2019), pp. 715–721. DOI: 10.1038/s41592-019-0494-8.
- [9] Mohammad Lotfollahi et al. “Predicting cellular responses to complex perturbations in high-throughput screens”. In: *Molecular Systems Biology* 19.6 (2023), e11517. DOI: 10.15252/msb.202211517.
- [10] Yusuf Roohani, Kexin Huang, and Jure Leskovec. “Predicting transcriptional outcomes of novel multigene perturbations with GEARS”. In: *Nature Biotechnology* 42.6 (2024), pp. 927–935. DOI: 10.1038/s41587-023-01905-6.

- [11] Stefan Peidli et al. “scPerturb: harmonized single-cell perturbation data”. In: *Nature Methods* 21.3 (2024), pp. 531–540. DOI: 10.1038/s41592-023-02144-y.
- [12] Yan Wu et al. “PerturBench: Benchmarking Machine Learning Models for Cellular Perturbation Analysis”. In: *arXiv preprint arXiv:2408.10609* (2024).
- [13] Gandhi, Javadi, et al. “Tahoe-x1: Scaling Perturbation-Trained Single-Cell Foundation Models to 3 Billion Parameters”. In: *bioRxiv* (2025). Preprint, posted 23 October 2025; not peer reviewed. scGPT-style encoder pretrained on Tahoe-100M among other corpora; code at <https://github.com/tahoebio/tahoe-x1>. DOI: 10.1101/2025.10.23.683759.
- [14] Constantin Ahlmann-Eltze, Wolfgang Huber, and Simon Anders. “Deep-learning-based gene perturbation effect prediction does not yet outperform simple linear baselines”. In: *Nature Methods* 22.8 (2025), pp. 1657–1661. DOI: 10.1038/s41592-025-02772-6.
- [15] Gerold Csendes et al. “Benchmarking foundation cell models for post-perturbation RNA-seq prediction”. In: *BMC Genomics* 26.1 (2025), p. 393. DOI: 10.1186/s12864-025-11600-2.
- [16] Ihab Bendidi et al. “Benchmarking transcriptomics foundation models for perturbation analysis: one PCA still rules them all”. In: *arXiv preprint arXiv:2410.13956* (2024).
- [17] Eric Kernfeld et al. “A comparison of computational methods for expression forecasting”. In: *Genome Biology* 26.1 (2025). Preprint: bioRxiv 2023.07.28.551039, titled “A systematic comparison of computational methods for expression forecasting”, p. 388. DOI: 10.1186/s13059-025-03840-y.
- [18] Tony Hauptmann and Stefan Kramer. “A fair experimental comparison of neural network architectures for latent representations of multi-omics for drug response prediction”. In: *BMC Bioinformatics* 24.1 (2023), p. 45. DOI: 10.1186/s12859-023-05166-7.
- [19] Judith Bernett et al. “From Hype to Health Check: Critical Evaluation of Drug Response Prediction Models with DrEval”. In: *bioRxiv* (2025). Preprint, posted 29 May 2025. DOI: 10.1101/2025.05.26.655288.
- [20] Henry E. Miller et al. “Deep Learning-Based Genetic Perturbation Models Do Outperform Uninformative Baselines on Well-Calibrated Metrics”. In: *bioRxiv* (2025). Preprint, posted 21 October 2025. Shift Bioscience Ltd. DOI: 10.1101/2025.10.20.683304.
- [21] Gabriel M. Mejia et al. “Diversity by Design: Addressing Mode Collapse Improves scRNA-seq Perturbation Modeling on Well-Calibrated Metrics”. In: *arXiv preprint arXiv:2506.22641* (2025).

- [22] Stuart H. Hurlbert. “Pseudoreplication and the Design of Ecological Field Experiments”. In: *Ecological Monographs* 54.2 (1984), pp. 187–211. DOI: 10.2307/1942661.
- [23] Alexander Partin et al. “Deep learning methods for drug response prediction in cancer: Predominant and emerging trends”. In: *Frontiers in Medicine* 10 (2023). DOI: 10.3389/fmed.2023.1086097.
- [24] Justin Lamb et al. “The Connectivity Map: using gene-expression signatures to connect small molecules, genes, and disease”. In: *Science* 313.5795 (2006), pp. 1929–1935. DOI: 10.1126/science.1132939.
- [25] Aravind Subramanian et al. “A Next Generation Connectivity Map: L1000 Platform and the First 1,000,000 Profiles”. In: *Cell* 171.6 (2017), 1437–1452.e17. DOI: 10.1016/j.cell.2017.10.049.
- [26] Yanai Elazar et al. “Amnesic Probing: Behavioral Explanation with Amnesic Counterfactuals”. In: *Transactions of the Association for Computational Linguistics* 9 (2021), pp. 160–175. DOI: 10.1162/tac1_a_00359.
- [27] Wes Gurnee, Nicholas Sofroniew, Adam Pearce, et al. “Verbalizable Representations Form a Global Workspace in Language Models”. In: *Transformer Circuits Thread* (2026). Published 6 July 2026. <https://transformer-circuits.pub/2026/workspace/index.html>.
- [28] Nelson Elhage et al. “Toy Models of Superposition”. In: *Transformer Circuits Thread* (2022). arXiv:2209.10652.
- [29] Marco Cuturi. “Sinkhorn Distances: Lightspeed Computation of Optimal Transport”. In: *Advances in Neural Information Processing Systems 26 (NIPS 2013)*. Ed. by C.J. Burges et al. Vol. 26. Curran Associates, Inc., 2013, pp. 2292–2300. URL: https://proceedings.neurips.cc/paper_files/paper/2013/file/af21d0c97db2e27e13572cbf59eb343d-Paper.pdf.
- [30] Geoffrey Schiebinger et al. “Optimal-Transport Analysis of Single-Cell Gene Expression Identifies Developmental Trajectories in Reprogramming”. In: *Cell* 176.4 (2019), 928–943.e22. DOI: 10.1016/j.cell.2019.01.006.
- [31] Romain Lopez et al. “Deep generative modeling for single-cell transcriptomics”. In: *Nature Methods* 15.12 (2018), pp. 1053–1058. DOI: 10.1038/s41592-018-0229-2.

A Appendix

A.1 Estimator validation against planted ground truth

Three of the measurements in this thesis rest on estimators built for the purpose rather than taken off the shelf. Each was validated by planting a known answer and checking recovery — in both directions, since an estimator that only ever returns the answer one hopes for is not evidence.

Estimator	Planted truth	Recovered	
transfer coefficient $\mathcal{T}$	$\mathcal{T} = 1.00$ (fully shared)	0.999	
	$\mathcal{T} = 0.70$	0.703	
	$\mathcal{T} = 0.40$	0.404	
channel gate	a channel that genuinely predicts	+0.397 (sd 0.029)	over 24 worlds
	a channel carrying nothing	−0.016 (sd 0.091)	over 24 worlds
κ -structure	a consistent per-line modulation	excess +1.173	
	an idiosyncratic interaction	excess −0.002	

Table 18: All three recover planted answers in both directions. Every row is a synthetic planted-world check — no data, no network, fixed seed — and none of these numbers comes from the evaluation run or is stored in a result artifact. The $\mathcal{T}$ rows and the κ -structure rows reproduce from `variance.decomposition.py --selftest`, and the channel-gate rows from `channel_gate.py --selftest`. The κ -structure check plants two worlds and requires the estimator to separate them: one in which every cell line bends whatever it is given in a characteristic direction, where within-line consistency must exceed across-line (+1.17), and one in which κ is specific to each drug-and-line pairing with no line-level direction, where the excess must vanish (−0.00). An estimator that reported structure in the second would be the failure this table exists to rule out.

The $\mathcal{T}$ rows above are a data-free unit test: a purely synthetic world in which the only axis a drug’s response is not shared across is cell line, so recovery there shows that the estimator inverts its own generative model — and nothing more. In the canonical run $\mathcal{T}$ was computed with `same_dose_only=false`, `loo_generic=false` and `repro_thr=0.2`, so cell line, dose, well and estimator noise all enter the same number. Repeated on synthetic data calibrated to that run’s own gene dimension, per-condition cell counts and empirical noise, the estimator returns 1.001, 0.747 and 0.494 for planted $\mathcal{T}$ of 1.00, 0.75 and 0.50. Neither the planted rows above nor this calibrated repeat licenses reading $1 - \mathcal{T}$ as a drug $\times$ cell-line interaction share: it is the fraction not shared across the compared contexts.

Two of these validations changed the estimators rather than merely endorsing them. The channel gate’s count-matched null originally drew its random partners from a pool *excluding* the channel’s own selections, which depletes the null of exactly the partners the channel chose and anti-correlates the two arms. And the useless-channel gap has a standard deviation of roughly 0.09 on a single small world, so the original validation —

asserting that one draw fell within 0.10 — tested luck rather than correctness; it now averages 24 independent worlds. That variance is also why the reported channel results are gated on a clustered interval rather than on a point estimate.

A third defect was found after the canonical run, and the channel numbers in this thesis are the re-run that followed it. The gate’s *retrieval* pool — the partners whose mean residual forms the prediction — was every retained condition in the cell line other than the target’s own drug, rather than the training conditions only. Conditions held out of training could therefore contribute to the prediction made for a held-out drug, so the gate did not do what §3.13 says it was intended to do — draw partners only from drugs that were in training. The code now restricts the pool to the split’s training conditions, logs the restricted size, and warns rather than silently falling back to an unrestricted pool, and the gate has been re-run under it. Of the 4,131 retained conditions 2,523 are training, so the old pool was around 64% larger than it was entitled to be. Every channel-gate figure reported in the main text is from the corrected run; the planted worlds above are a separate synthetic path with no held-out split and are unaffected. The repair did not move the channels in a common direction — target fell, mechanism rose, chemistry roughly halved — which is why the pre-fix figures are not quoted alongside them.

A.2 Absent-gene convention

A predicted cell sentence lists only expressed genes, so some panel genes are unranked and a convention is required. Two were carried throughout: **worst** assigns absent genes the last rank G , **francesca** a fixed middle rank $G/2$. The choice does not matter — spike-in DE- Δr reads 0.952 under the first and 0.954 under the second, and no quantity scored under both conventions moves by more than 0.01 between them. The convention is an input only to the rank-based scores reported here, DE- Δr and panel- τ ; the canonical run’s headline metric, NIR, is scored in residual space by a separate instrument that never receives the absent-gene fill, so the choice cannot reach it. All figures in the body are quoted under **worst**.

A.3 Swap-distance sweep

The stratified scramble of Section 3.3 was preceded by a continuous version: Δ_{scr} measured as a function of the response distance between the real and the swapped-in drug, binned into quintiles. It is a separate, earlier run and none of its numbers belong to the canonical set: the two columns are the single-cell and optimal-transport checkpoints rather than the residual arm the canonical results describe, they were scored on the unseen-drug tier of a different evaluation set, and their swap partners were not restricted to the same plate, whereas the canonical run draws partners within plate. The artifacts are `scramble_sweep_single.json` and `scramble_sweep_T2.json`.

Bin (swap distance)	single-cell model	optimal-transport model
0 — nearest (2.75–4.03)	−0.046	−0.001
1 (4.03–4.60)	+0.060	+0.031
2 (4.60–4.93)	−0.006	+0.017
3 (4.93–5.32)	−0.004	+0.043
4 — farthest (5.32–6.25)	−0.019	−0.005
trend, corr(distance, Δ_{scr})	+0.051	+0.012

Table 19: No bin in either model has an interval excluding zero, and the trend correlations are ≈ 0 . 200 (drug, swap) pairs over 20 cell lines, $n = 40$ per bin.

Two features make this an argument rather than five nulls. The single borderline value sits in an *inner* bin — the second of the five, +0.060 [−0.002, +0.124] — whereas a genuine dissimilarity effect would be monotone and peak in the far bin. And the comparison is paired: the model’s own NIR moves across bins (0.497, 0.612, 0.523, 0.401, 0.326 from nearest to farthest — falling overall, but not monotonically), because drugs that are far from everything populate the far bins, but Δ_{scr} uses the same drug on both arms so that drift cancels.

The intervals are wide — $n = 40$ per bin gives half-widths of 0.06–0.10 in the single-cell column and 0.03–0.06 in the optimal-transport column — so this sweep excludes a large effect rather than establishing a null, which is why the stratified scramble replaced it in the body.

A.4 Training configuration

Every arm was cold-started from the same base checkpoint with identical hyperparameters, so that the target or objective is the only variable. One exception is material and is stated wherever the arm is quoted.

A.5 The identity test: failed designs and remaining guards

The substitution test of Section 3.6 is the fourth design. The three that failed are recorded here with their algebra, because each produced results this thesis previously quoted as findings, and the constraints they teach are the surviving design.

Failed design one: mismatched geometry. The first test overwrote drug A’s code with drug B’s, inside a ≤ 10 -dimensional purified subspace, and compared the resulting KL against a random vector of identical norm drawn in the *full* 2048-dimensional stream. Roughly 99.5% of the control’s energy then lay in directions the intervention never touched, including the cell-line directions measured at ≈ 0.6 variance share. “The swap moves the output less than noise” is exactly what that geometry predicts, whatever the readout does.

Arm	Target	Training	Note
single cell	the observed treated cell	1 epoch	
consensus	group pseudobulk	$\approx 60\%$ of an epoch	tier-1 NIR unscorable
optimal transport (T2)	barycentric, $\varepsilon=0.5$	1 epoch, 23,842 steps	
residual	signed residual signature	1 epoch, 15,340 steps	
residual (holdout)	as above, 738 withheld	1 epoch	

Table 20: Learning rate 10^{-5} , batch size 1, gradient accumulation 16, maximum length 8192, bf16, weight decay 0.01, warmup ratio 0.03 throughout. The consensus arm stopped early on a flat loss, so it is the ε -ladder’s weakest-evidenced endpoint; and at the residual arm’s generation budget against a neutral partner the optimal-transport endpoint is not null either (Section 3.10), so outside the scoring of Table 8 the ladder’s negative rests on the single-cell endpoint alone. These five arms — **single_cell**, **consensus**, **ot.T2**, **residual**, **residual_holdout** — are also the complete family over which the identity test of Section 3.6 takes its multiplicity correction: five arms by up to seven depths, 26 measured tests — the family size behind the q -values quoted there.

Failed design two: ablate-then-add. The second drew the control inside the subspace, which looked like a repair and was not. The hook ablates before adding, so what is delivered is (added – removed), and an uncentred class mean shares its global-mean component with the thing removed. The substitution arm was then close to a restoration and the control arm a real displacement. Simulated with no model present, the control delivers $5.3\times$ the perturbation in 100% of draws, with the predicted gap $2\langle\mu_B, u\rangle = 9.179$ against an observed 9.162, where μ_B is the injected class mean and u the cell’s own component in the slab. The retracted result was geometry, not biology.

Failed design three: the normaliser. The third injected the displacement and scored $\Delta \log P(g_B)$ against a norm-matched random direction — and was confounded the opposite way. A random direction inside the slab is off the class-mean manifold: it inflates the logsumexp normaliser (measured $+2.75$ nats) where the on-manifold displacement does not, so every target is depressed under the control arm whether or not it carries identity. With the scored target stripped of all identity information, 36% of the effect still stood. Norm-matched is not disruption-matched. The same argument retires the simplest test of all, “how far did the output move”: distance moved is exactly what norm-matching equalises, so that test has no power in precisely the world where the readout does read the drug.

Three further guards from the surviving design are engineering rather than inference and are recorded here. A *drug-proxy screen* refuses any context label that is a near-bijection with drug — it caught the metadata field `dose_float` holding a sample identifier in the residual builders, which, had it been orthogonalised against, would have deleted the drug axis by construction. An *injected-norm diagnostic* reports the displacement against the cell’s centred slab component (0.69–1.40: the injection is the size of the thing it displaces, so a null cannot be blamed on a whisper). And a *selftest with teeth*

runs three planted worlds — drug inert, drug read, and drug inert behind a readout sensitive to displacement magnitude rather than direction — through the actual headline code path; it separates them $(-0.042, +1.520, -0.010)$, and it failed four times during development, each failure a real defect that the guard list now covers. That third world is not a normaliser control: each planted response is a single token, so nothing diverges, and the state-dependent part of the normaliser is precisely what the selftest cannot exercise (Section 3.6).

A.6 Calibration self-test

The calibration framework of Section 3.2 is itself checked, since a metric audit whose instrument is untested proves nothing. Beyond deterministic unit checks — a perfect prediction must reach the metric’s maximum, a zero shift and a flat truth must return no score — the self-test asserts the exploit under a *second, independent* zero-information baseline: a leave-one-out mean scores $\text{DE-}\Delta r = 1.000$ while scoring $\text{NIR} = 0.000$. Two unrelated constructions both saturating the standard metric, and both scoring zero on the discrimination metric, establishes the exploit as a property of the metric rather than of any one baseline.